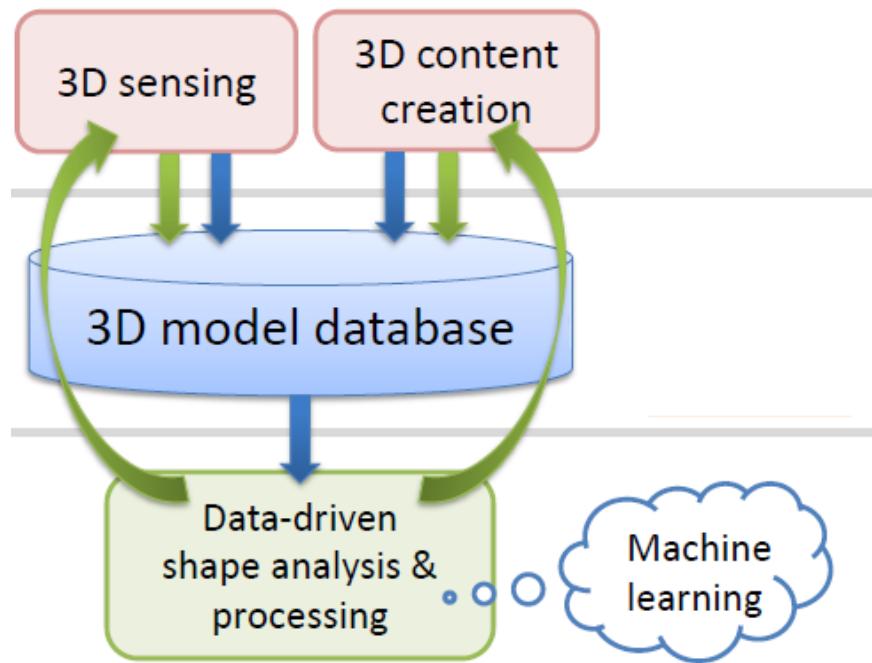


Data-Driven Shape Analysis and Processing



Presenter: Evangelos Kalogerakis, UMass Amherst

STAR report authors: Kai Xu, Vladimir Kim, Qixing Huang, Evangelos Kalogerakis

3D shapes for computer-aided design



Architecture



Interior design

3D shapes for information visualization



Geo-visualization



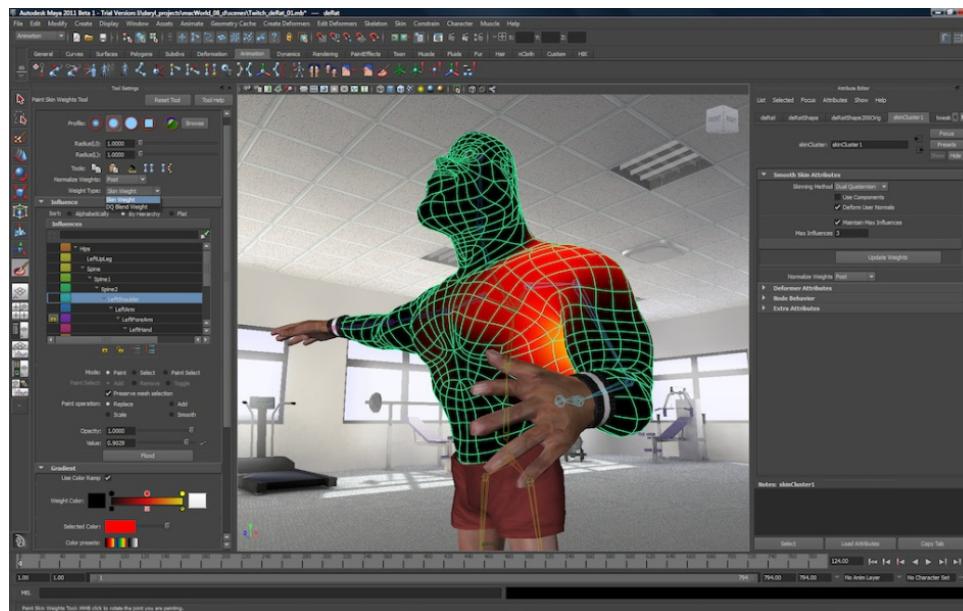
Scientific visualization

3D shapes for digital entertainment



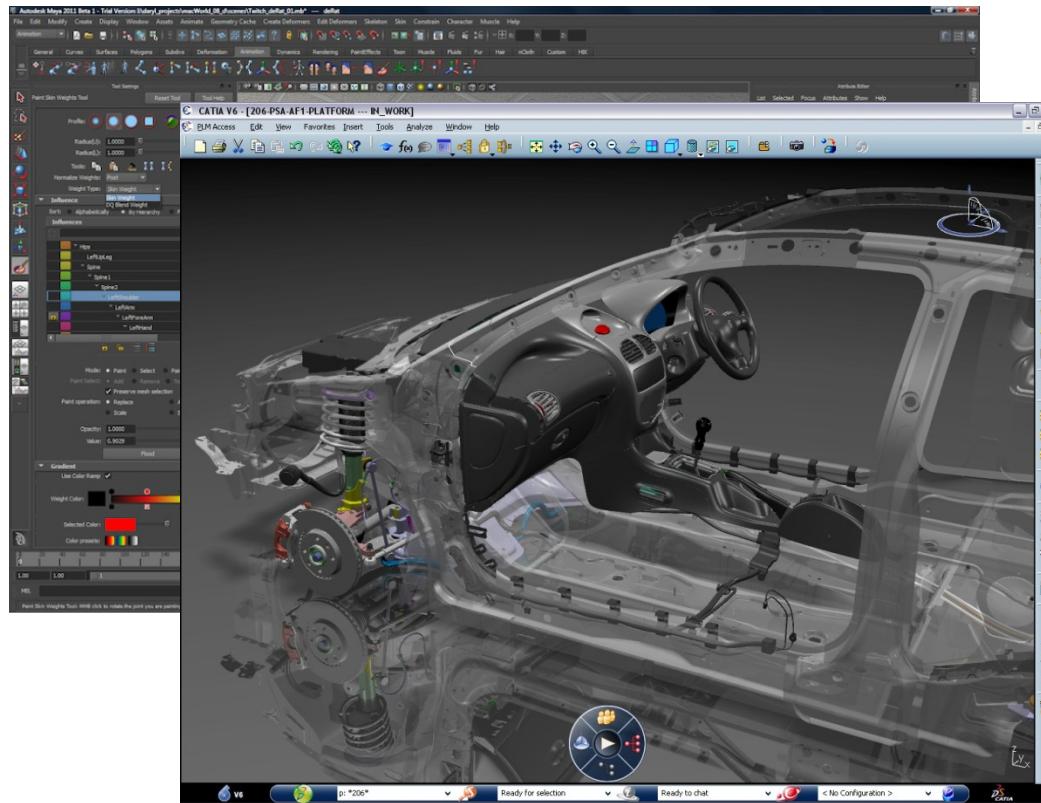
Video games

Digitizing our imagination



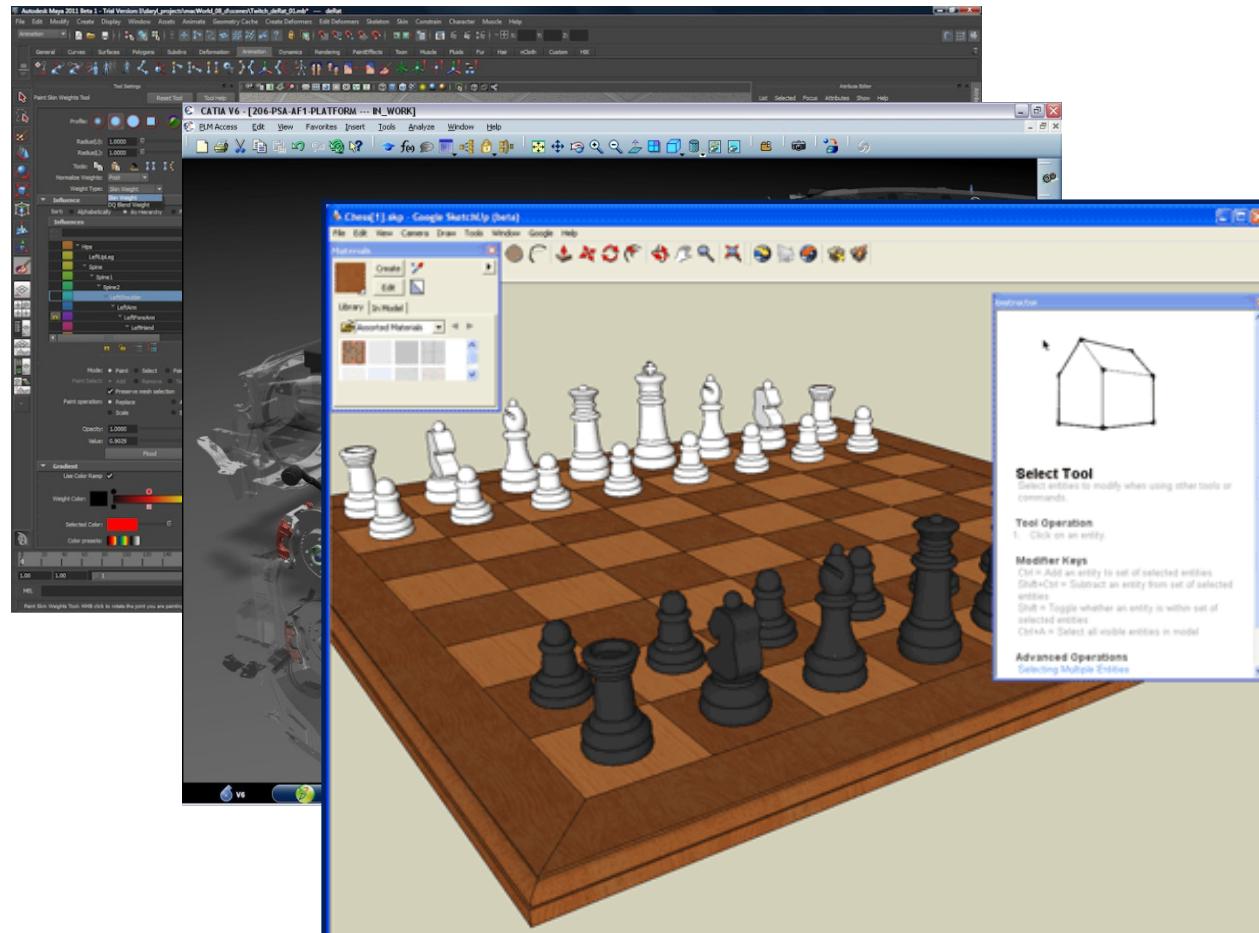
**Professional 3D modeling tools
[Autodesk Maya]**

Digitizing our imagination



**Computer-Aided Design tools
[Catia]**

Digitizing our imagination



General-Purpose Modeling tools
[Sketch-up]

3D shape repositories

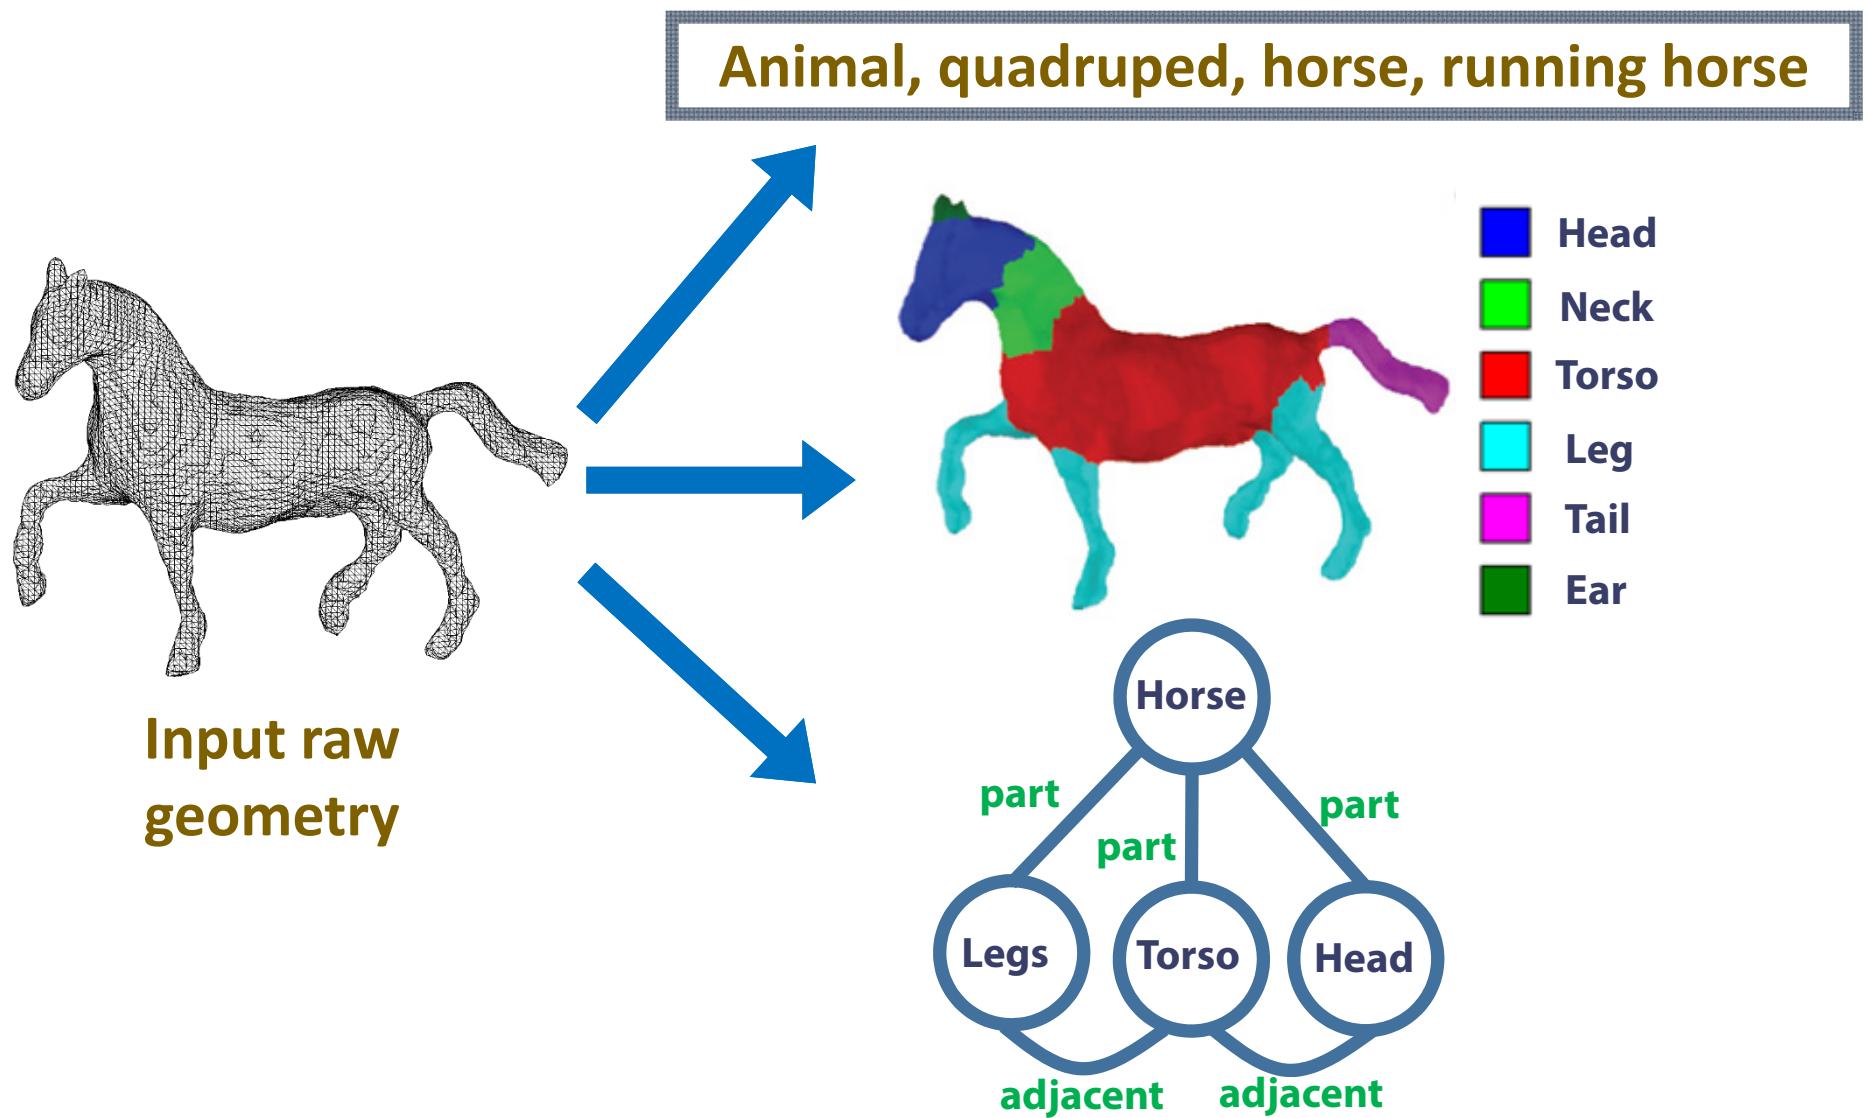
Google 3D warehouse [Advanced Search](#)

3D Warehouse Results Results 1 - 12 of many for cat (0.3 seconds) - [RSS](#)

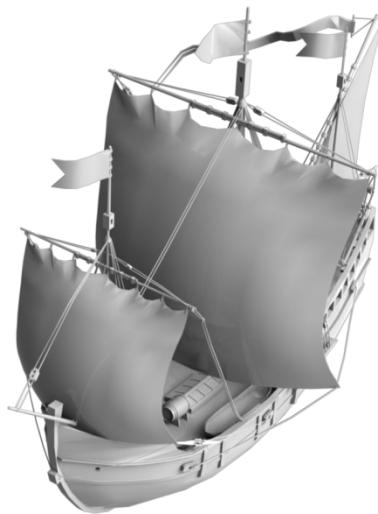
 CAT 797B Franco Peña by einstein El camión más grande del... Download to Google SketchUp 6 	 Cat by Stefy Gatto, Felino, Download to Google SketchUp 6 	 cats by rubicundo2 (2D) tres gatitos para... Download to Google SketchUp 6
 MM Tank-Bot V2 [For Phat... by Will I actually have a reason for... Download to Google SketchUp 6 	 Mad cat - Timber Wolf battle... by grenier.day This is a mad cat that I drew... Download to Google SketchUp 6 	 Cat Souvenir by Piper From 3D Collections Download to Google SketchUp 7
 MM Plasma Sniper [For Phat... by Will The Marble Men... Download to Google SketchUp 6 	 MM Assault Rifle [Entering it... by Will Fully automatic Marble Man... Download to Google SketchUp 6 	 Big solar powered Space... by Shogun(The rarely... This is a big solar powered... Download to Google SketchUp

[Trimble Warehouse]

Shape understanding



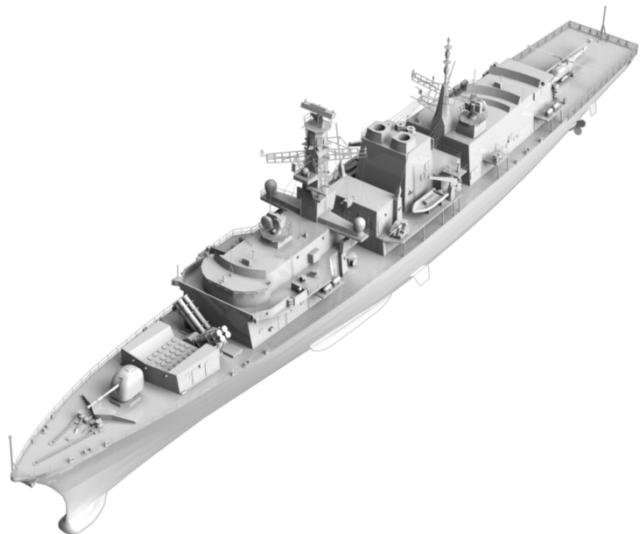
Why shape understanding? Shape categorization



**Sailing Ship,
Galleon**



**Sailing ship,
Yawl**



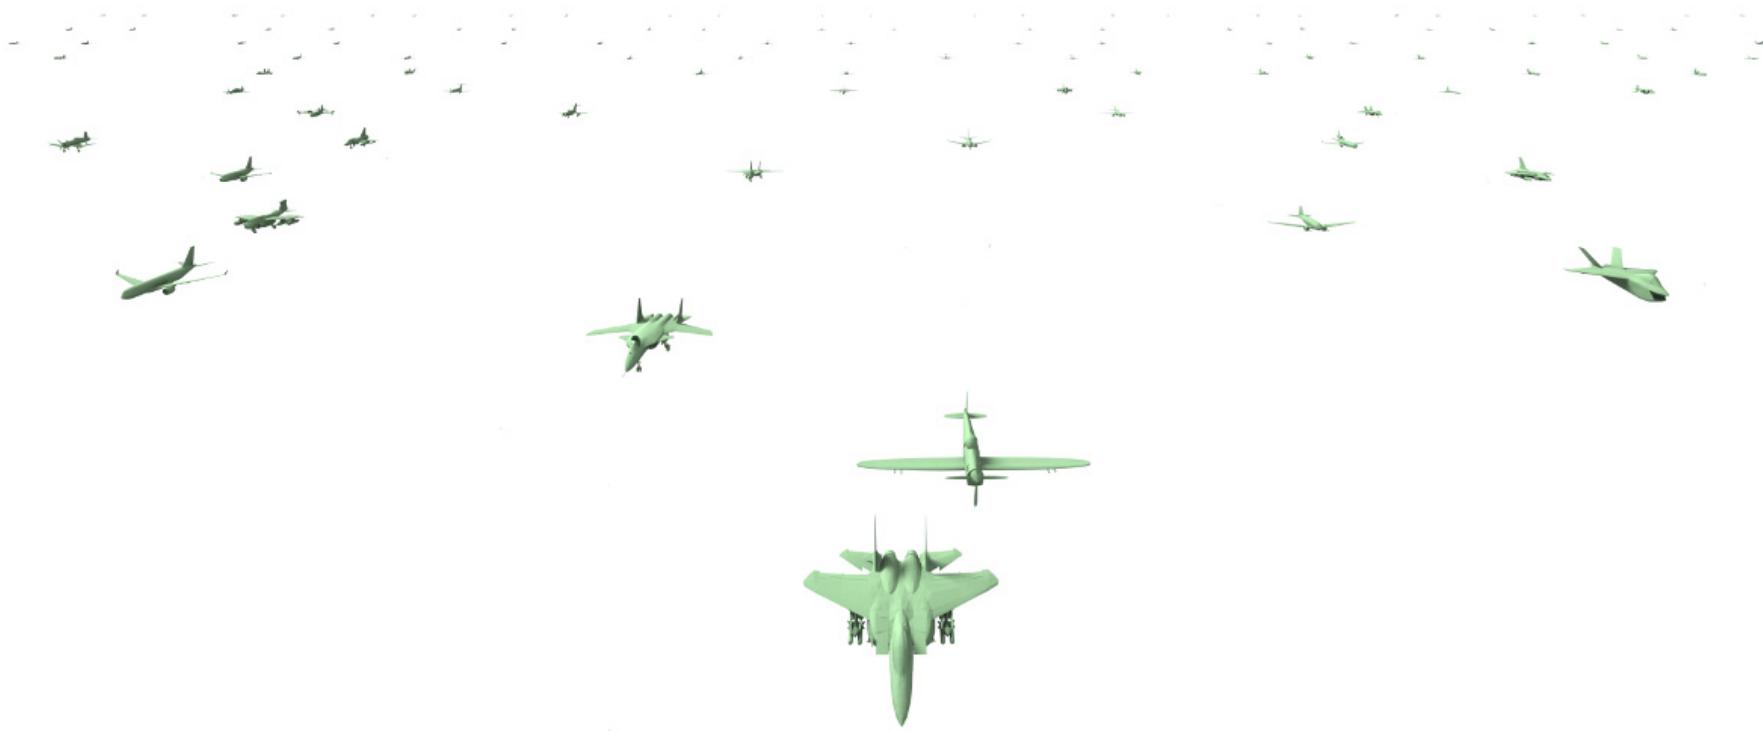
**Military ship,
Frigate**

Why shape understanding? 3D Modeling

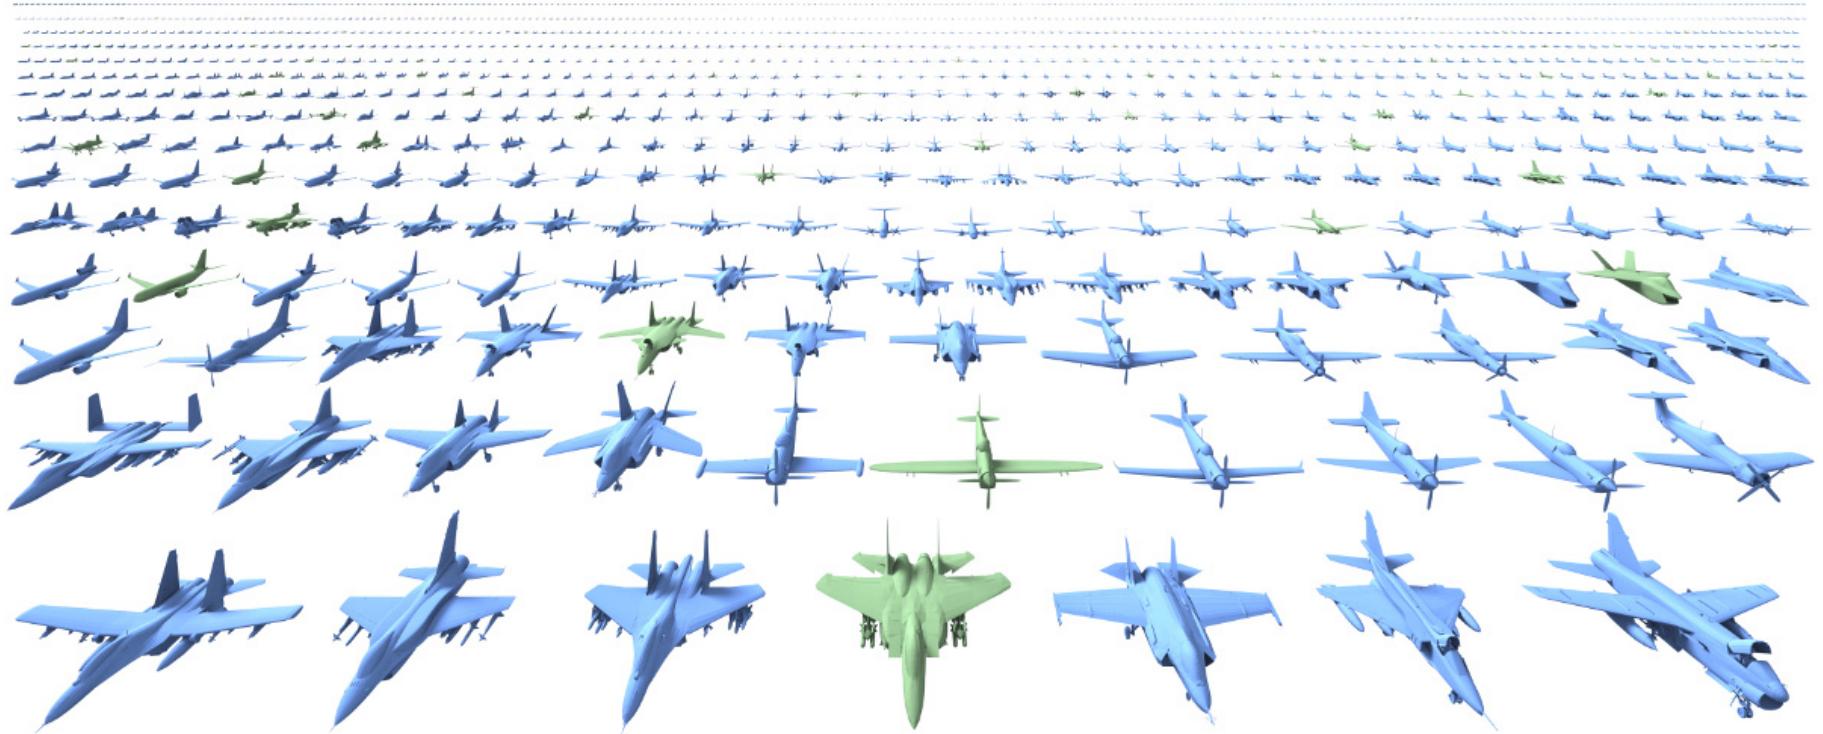


*Chaudhuri, Kalogerakis, Guibas, Koltun, SIGGRAPH 2011
access video: <https://www.youtube.com/watch?v=7Abki79WIOY>*

Why shape understanding? Shape synthesis



Why shape understanding? Shape synthesis



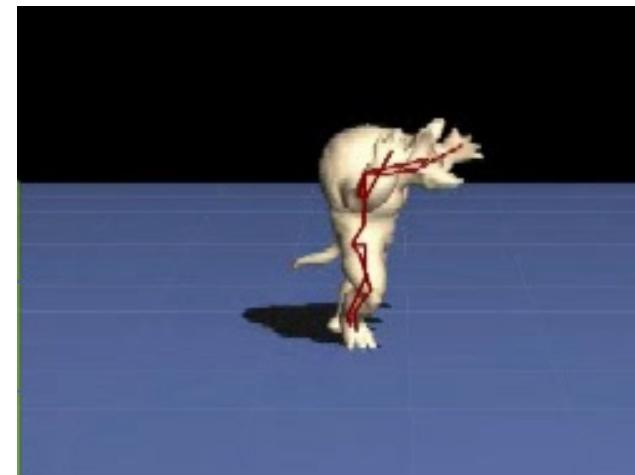
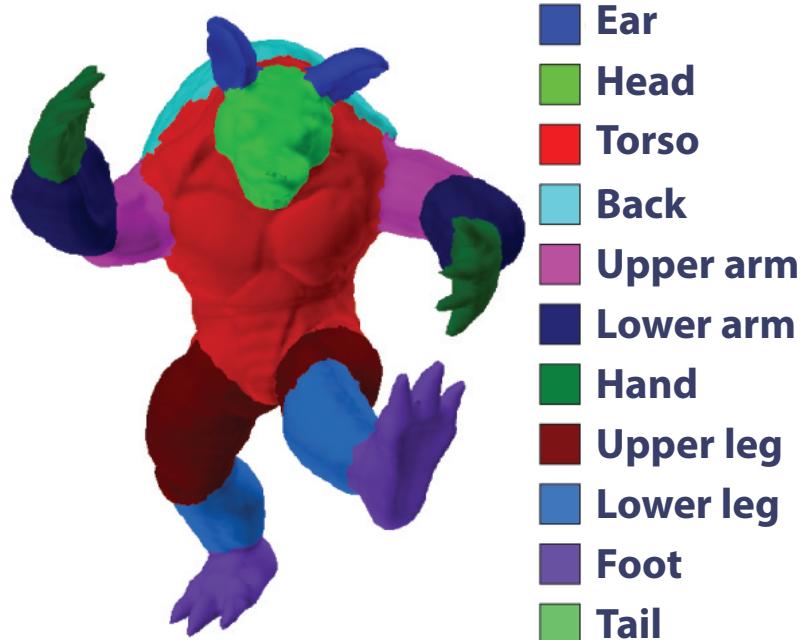
Kalogerakis, Chaudhuri, Koller, Koltun, SIGGRAPH 2012

Why shape understanding? Texturing



Kalogerakis, Hertzmann, Singh, SIGGRAPH 2010

Why shape understanding? Character Animation



Simari, Nowrouzezahrai, Kalogerakis, Singh, SGP 2009

Why data-driven methods for shape understanding?

It is extremely hard to perform shape understanding **with a set of deterministic, manually specified rules!**

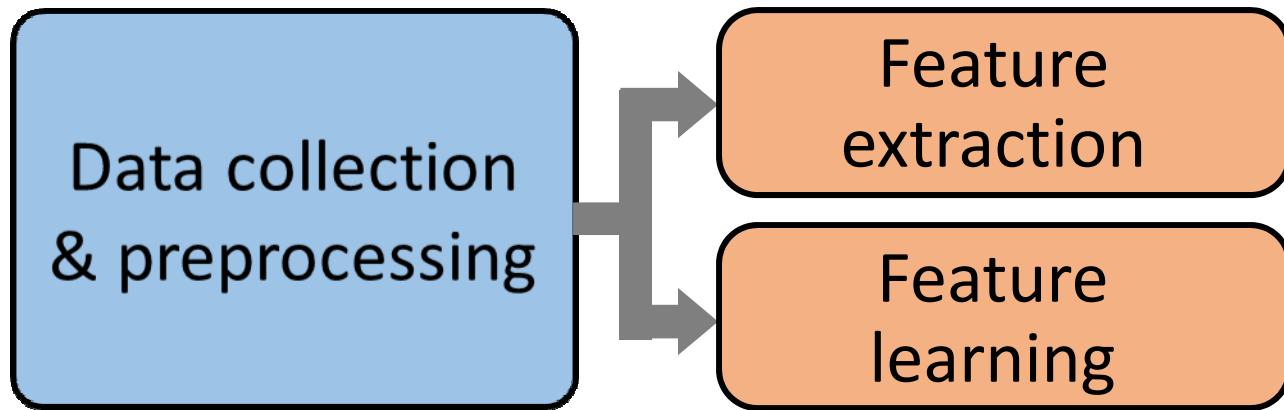
We should not treat shapes **in complete isolation of all others.**



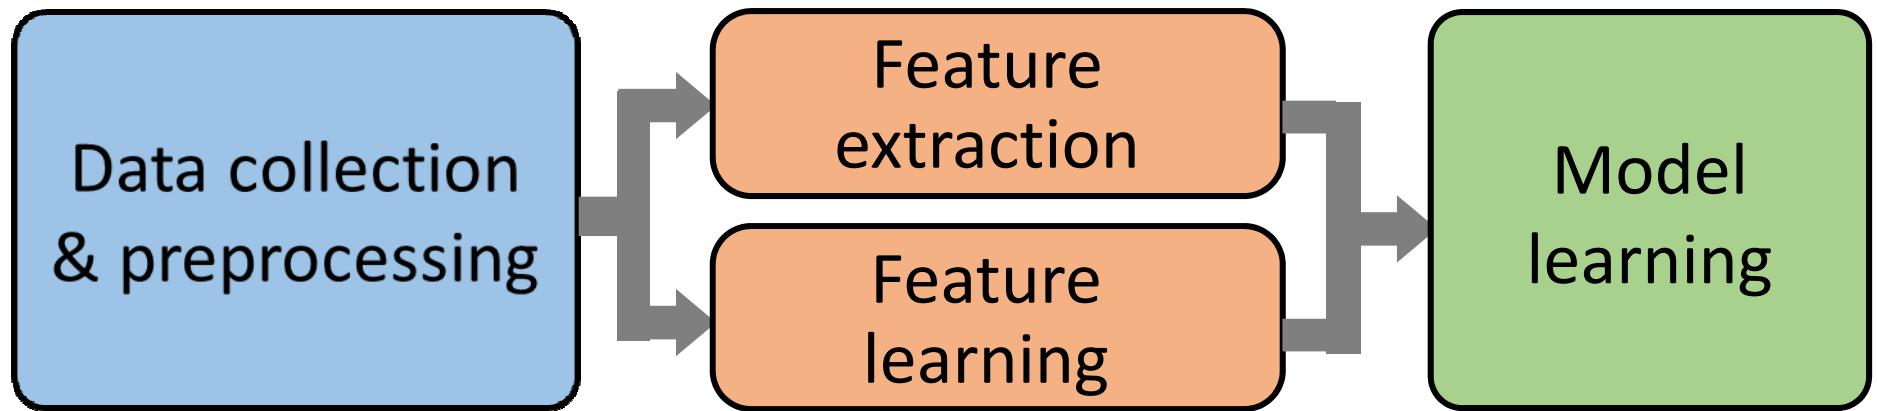
Pipeline of data-driven methods

Data collection
& preprocessing

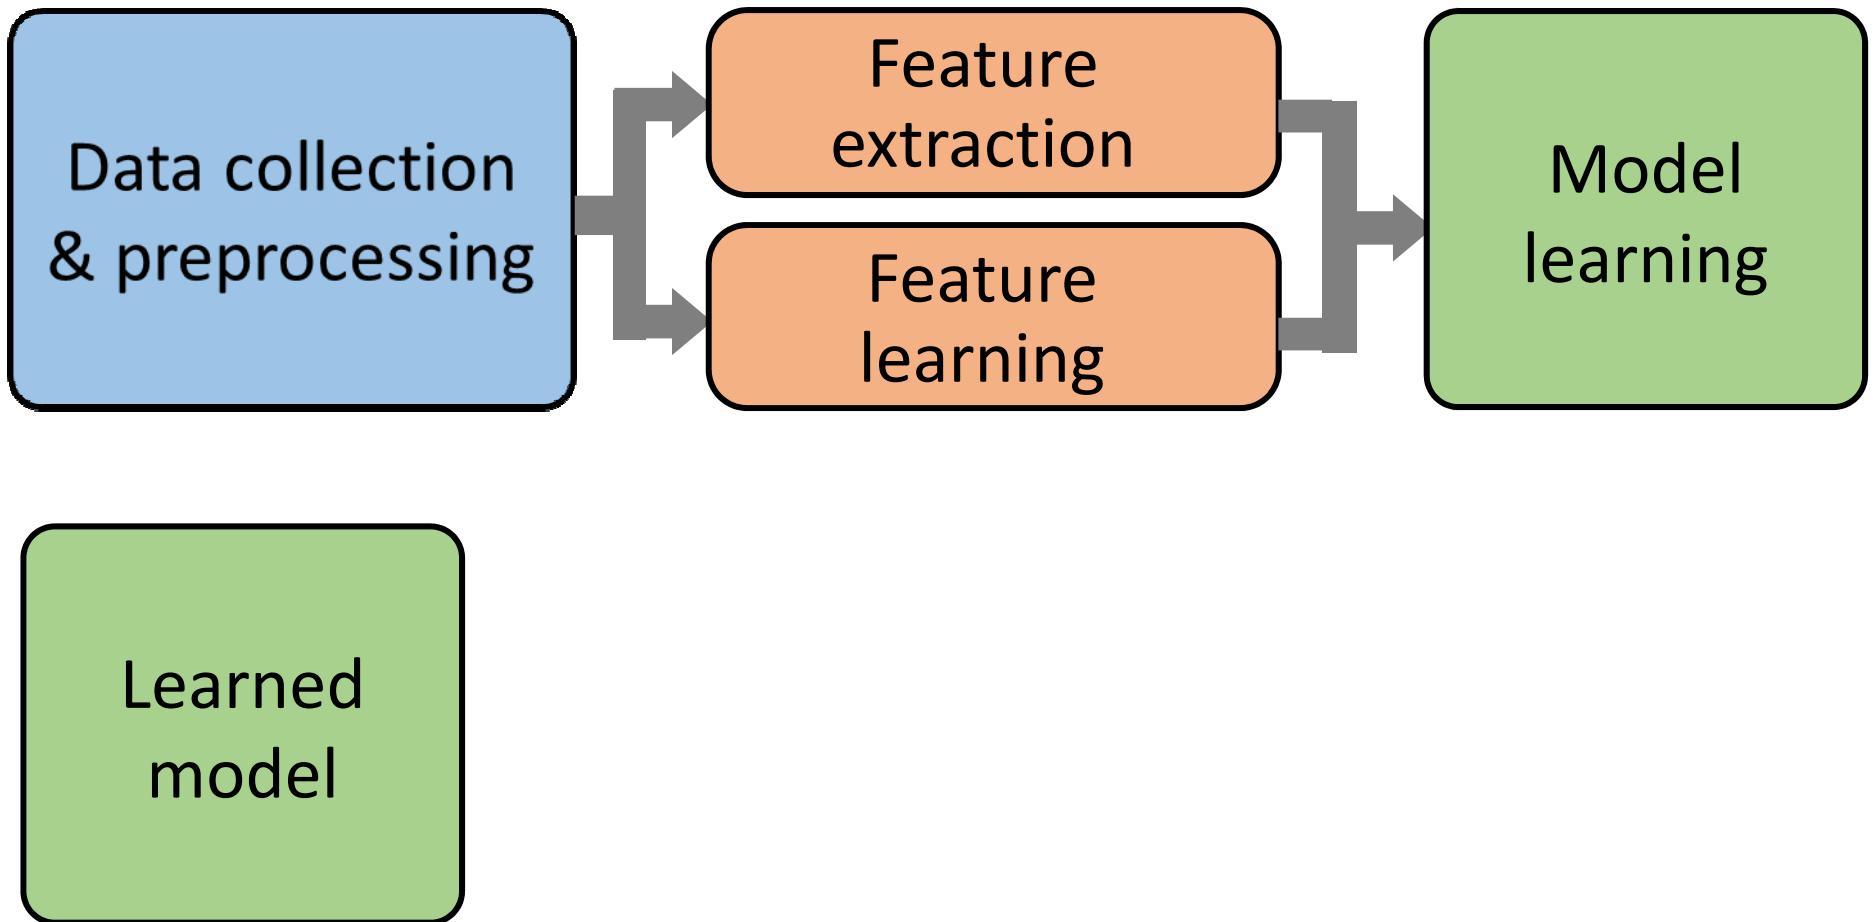
Pipeline of data-driven methods



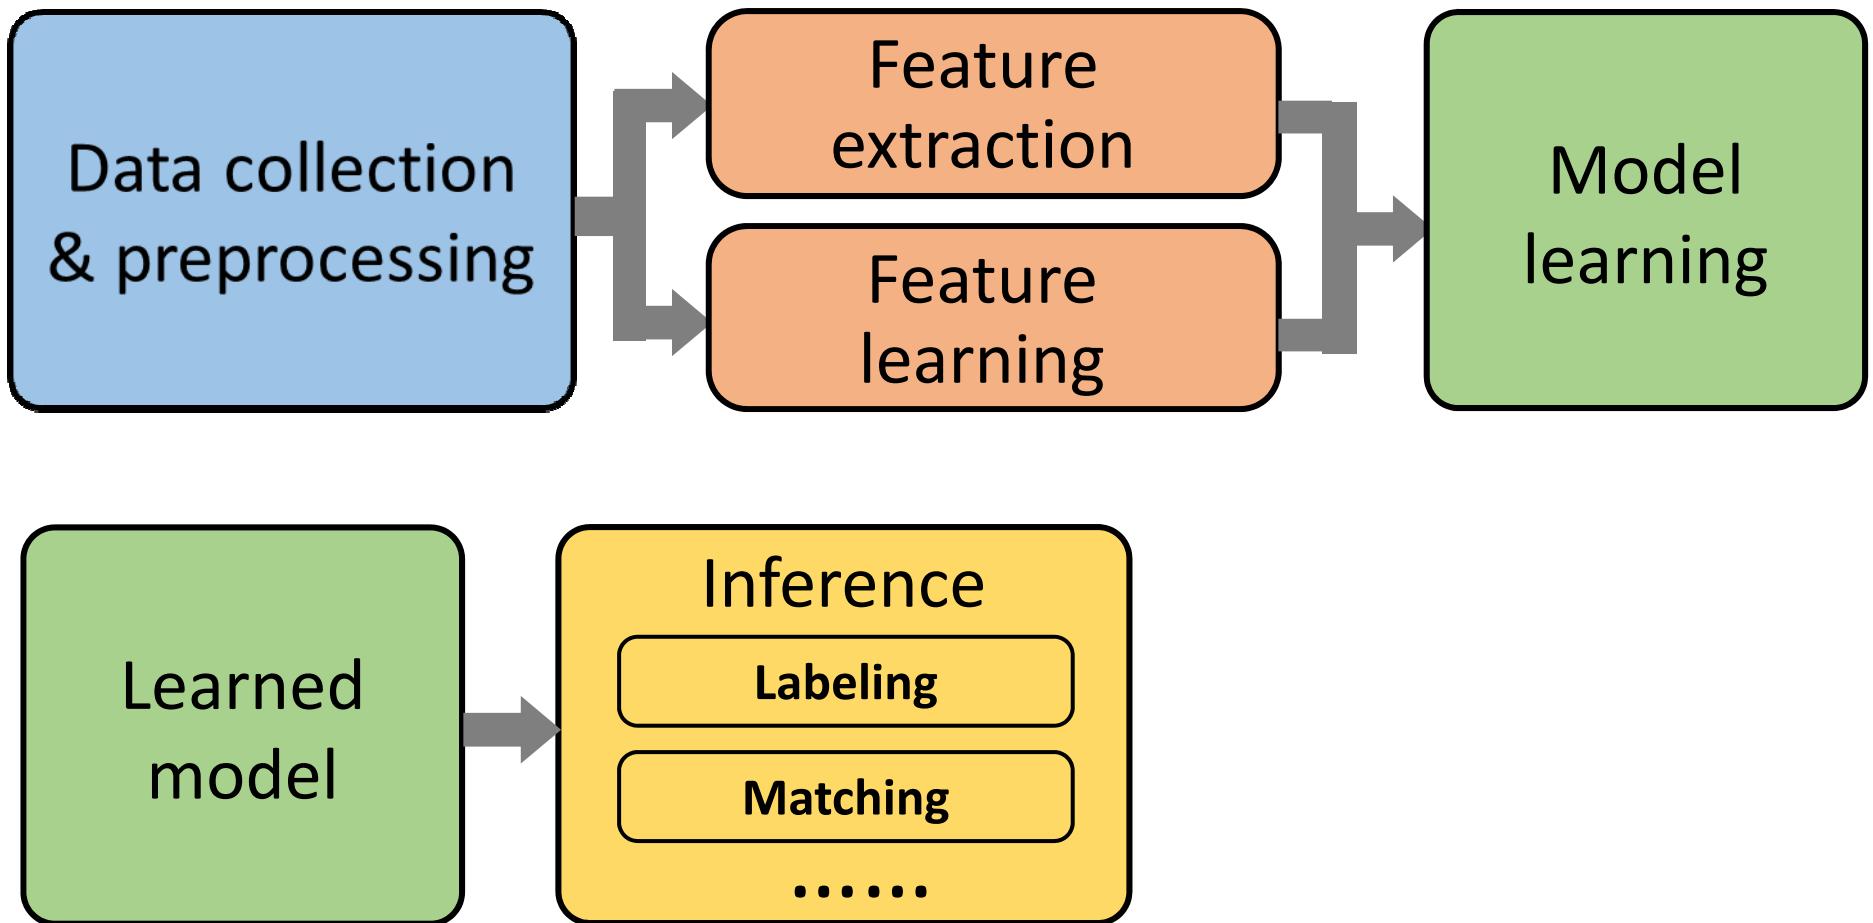
Pipeline of data-driven methods



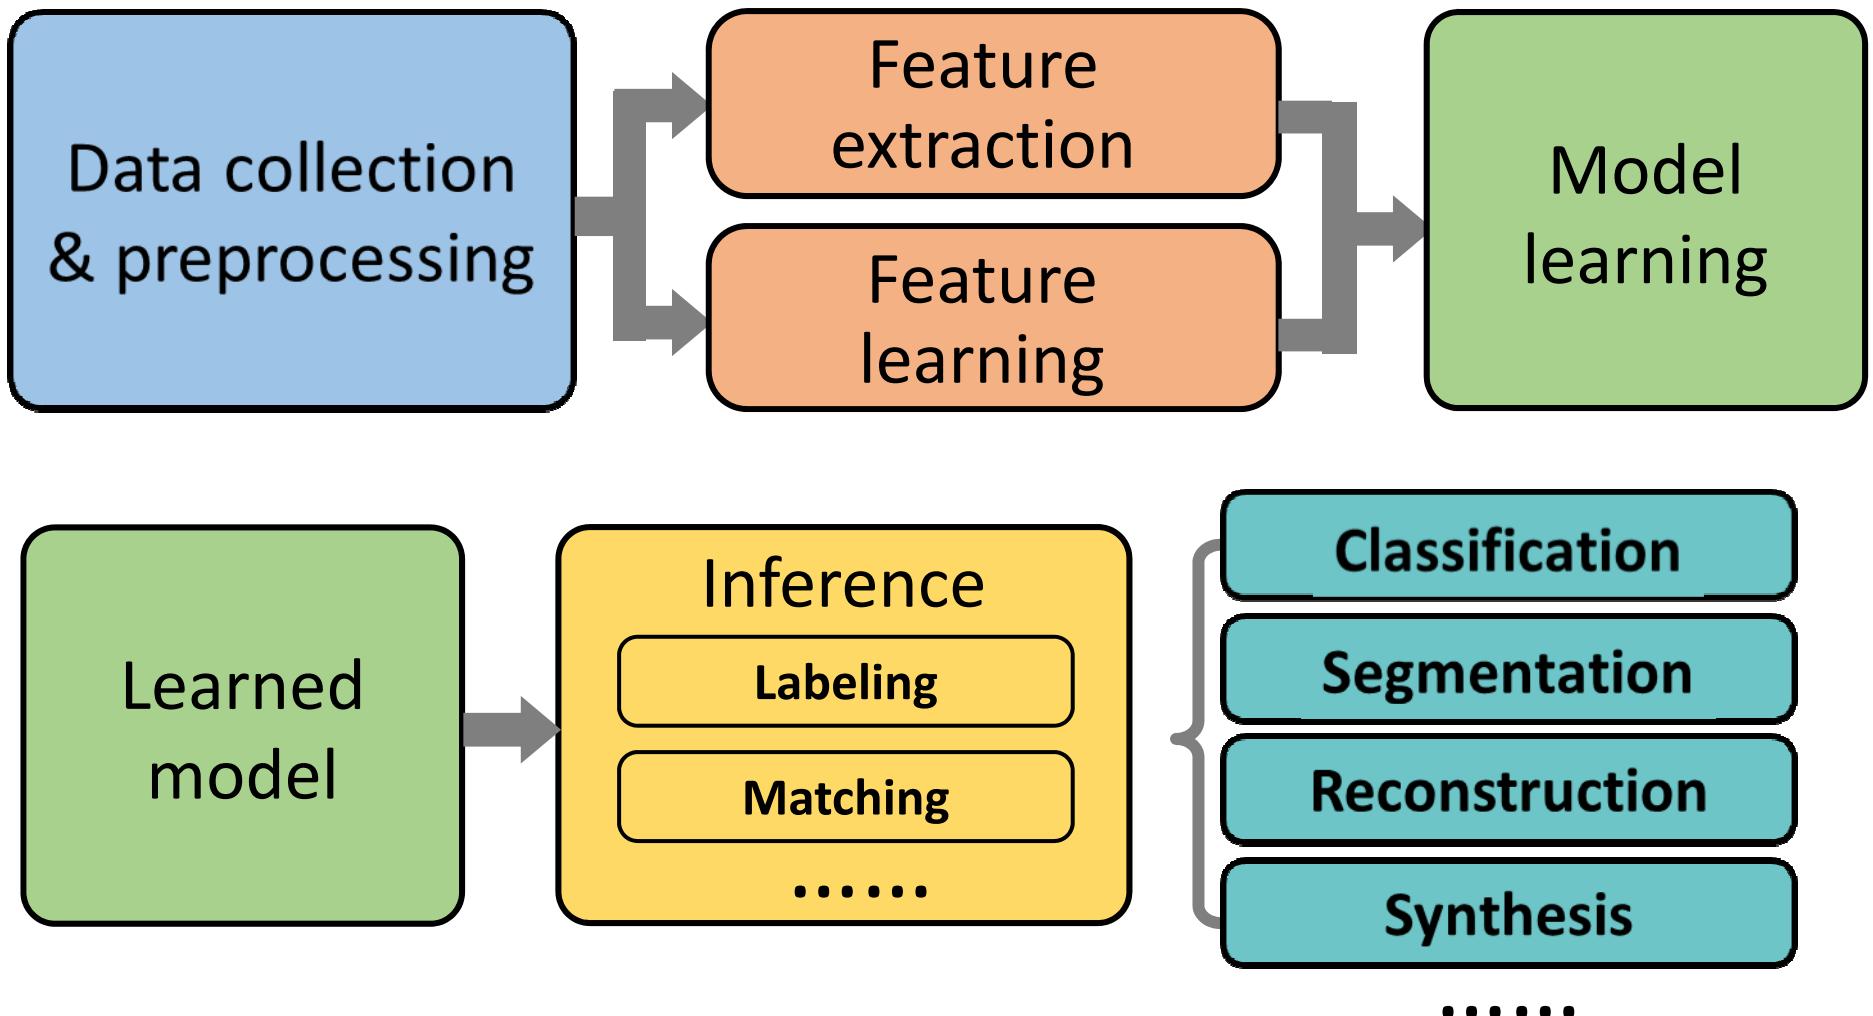
Pipeline of data-driven methods



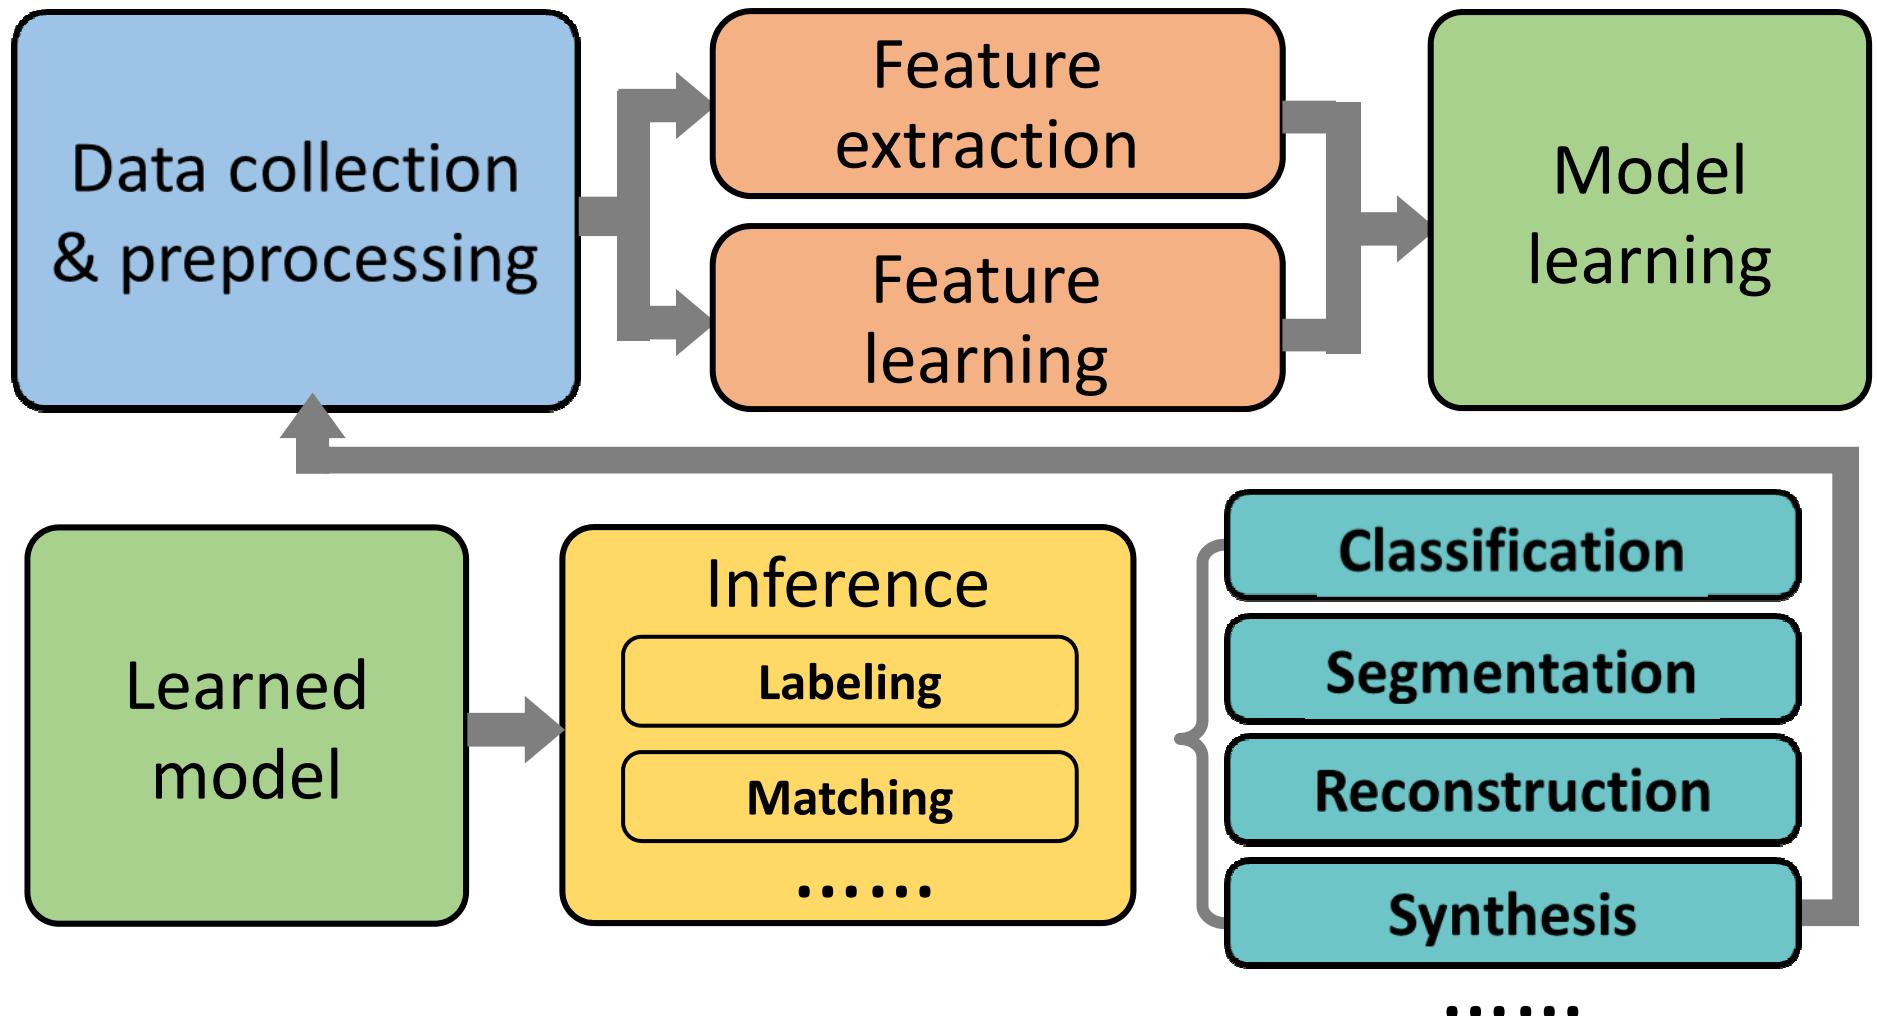
Pipeline of data-driven methods



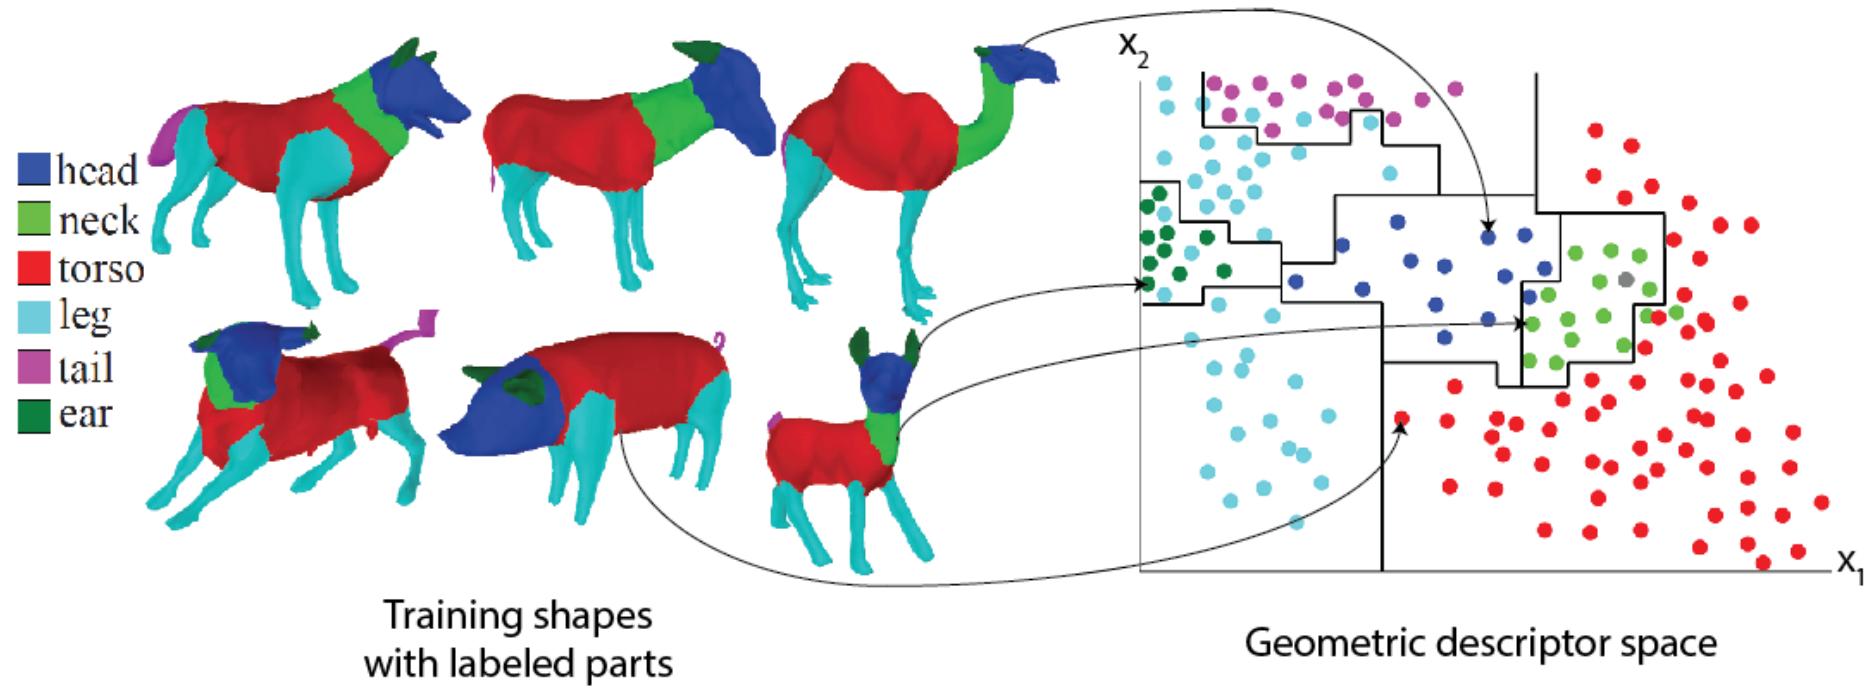
Pipeline of data-driven methods



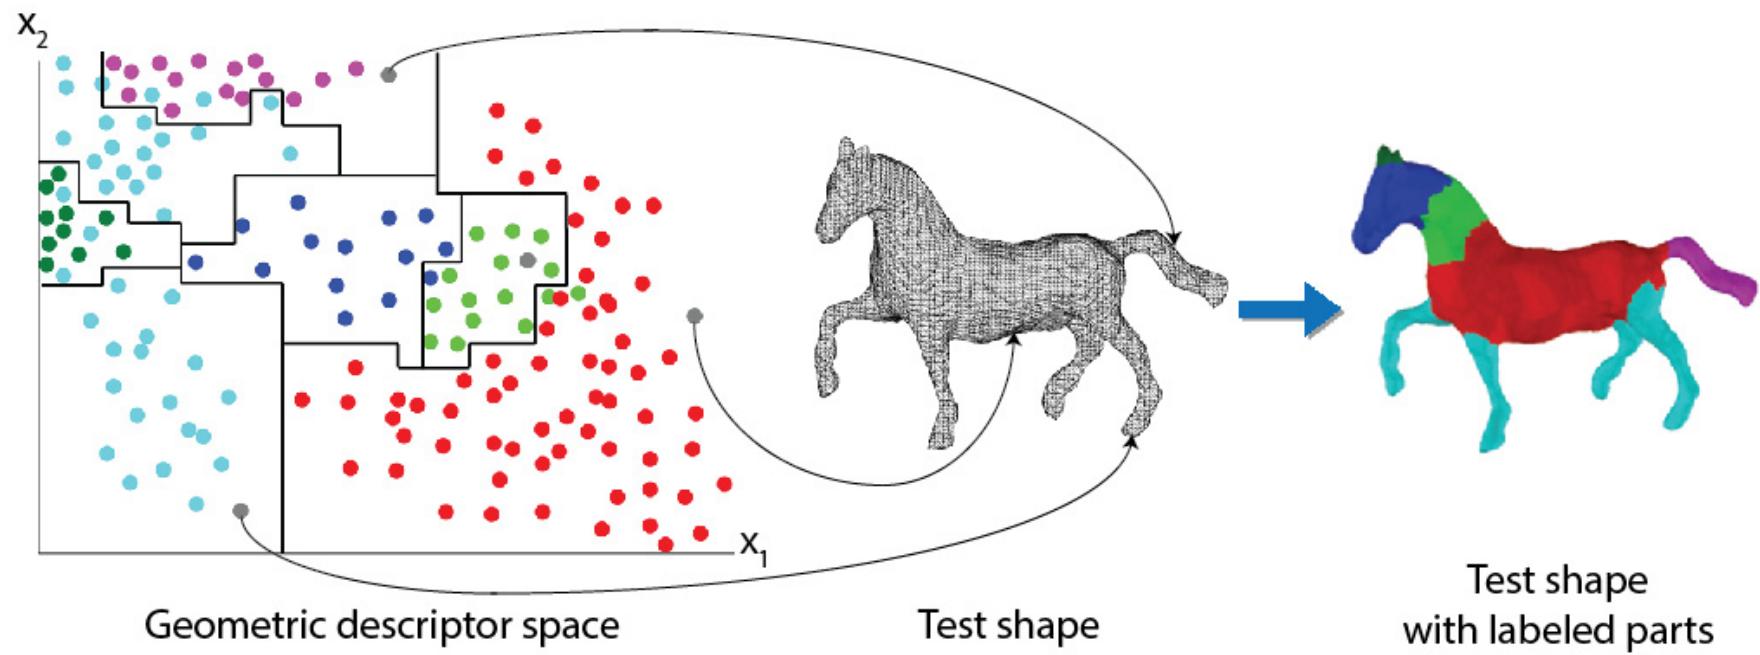
Pipeline of data-driven methods



Example: part labeling [training stage]



Example: part labeling [test stage]



STAR report

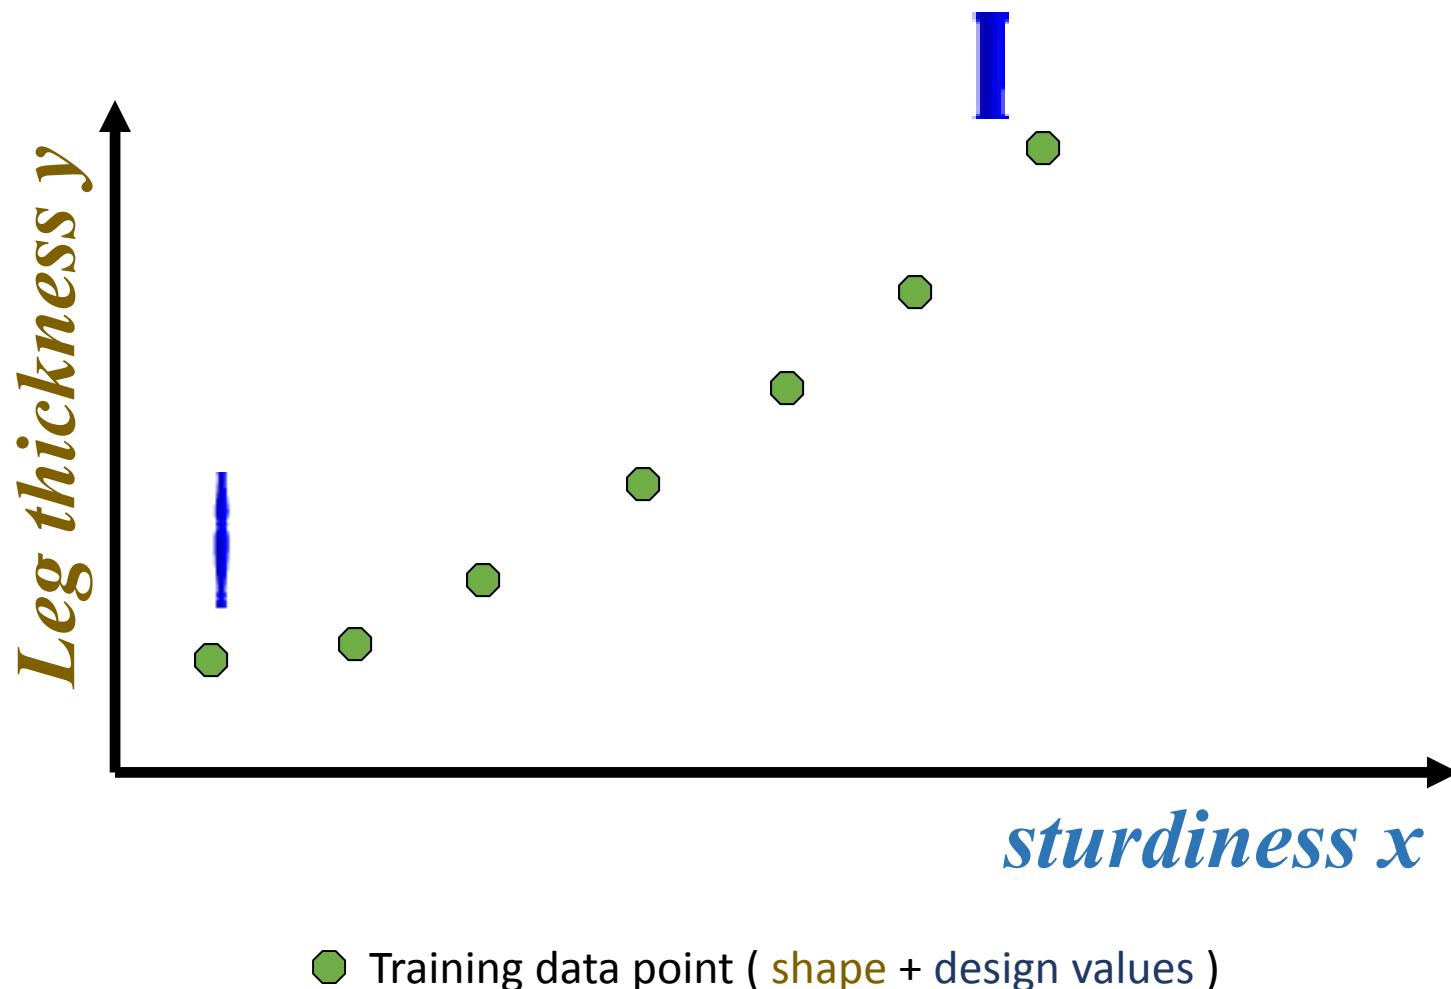
Overview of 200+ works published in the field, 30 pages. Enjoy!

1. Pipeline of data-driven shape processing techniques
2. Learning
3. Shape classification and retrieval
4. Shape segmentation
5. Shape correspondences
6. Shape reconstruction
7. Shape modeling and synthesis
8. Scene analysis and synthesis
9. Exploration and organization of 3D model collections

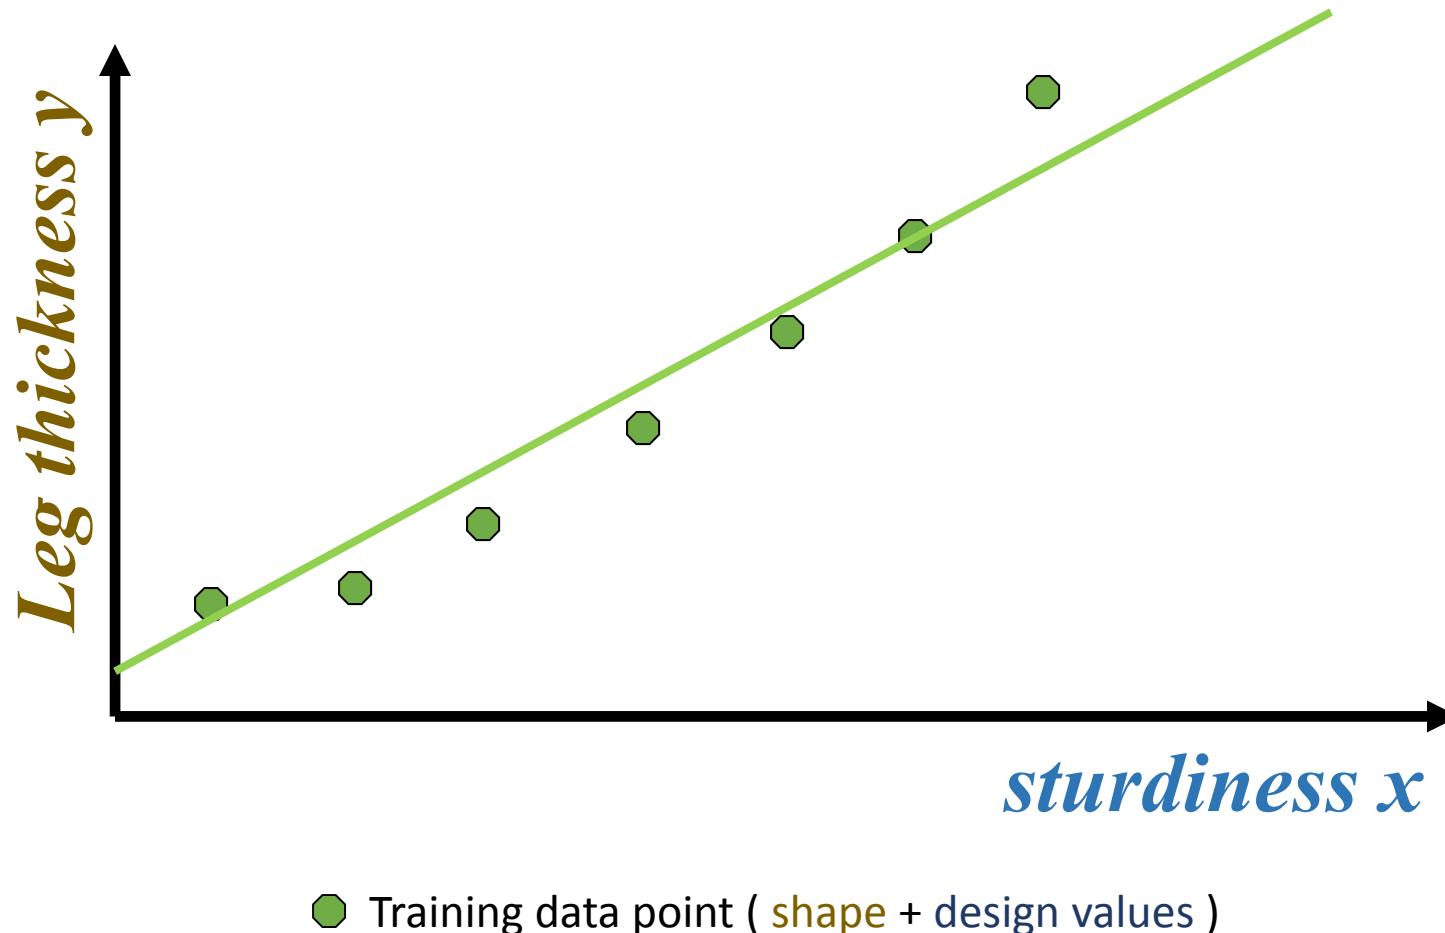
In this talk....

1. **Pipeline of data-driven shape processing techniques**
2. **Learning**
3. **Shape classification and retrieval**
4. Shape segmentation
5. Shape correspondences
6. Shape reconstruction
7. **Shape modeling and synthesis**
8. Scene analysis and synthesis
9. Exploration and organization of 3D model collections

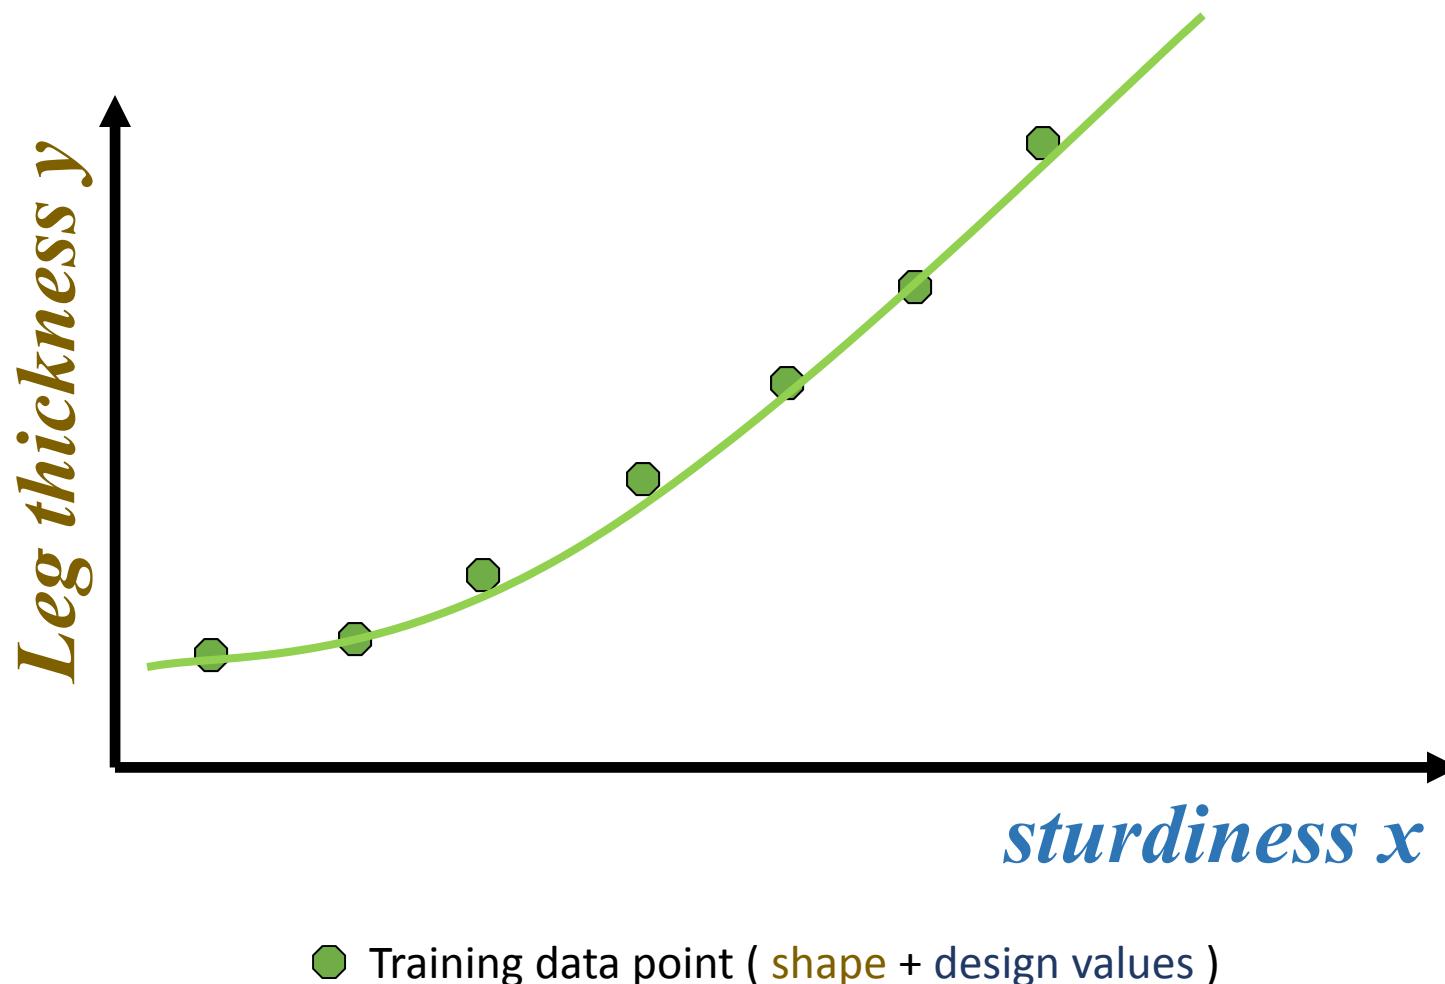
Learning basics: regression



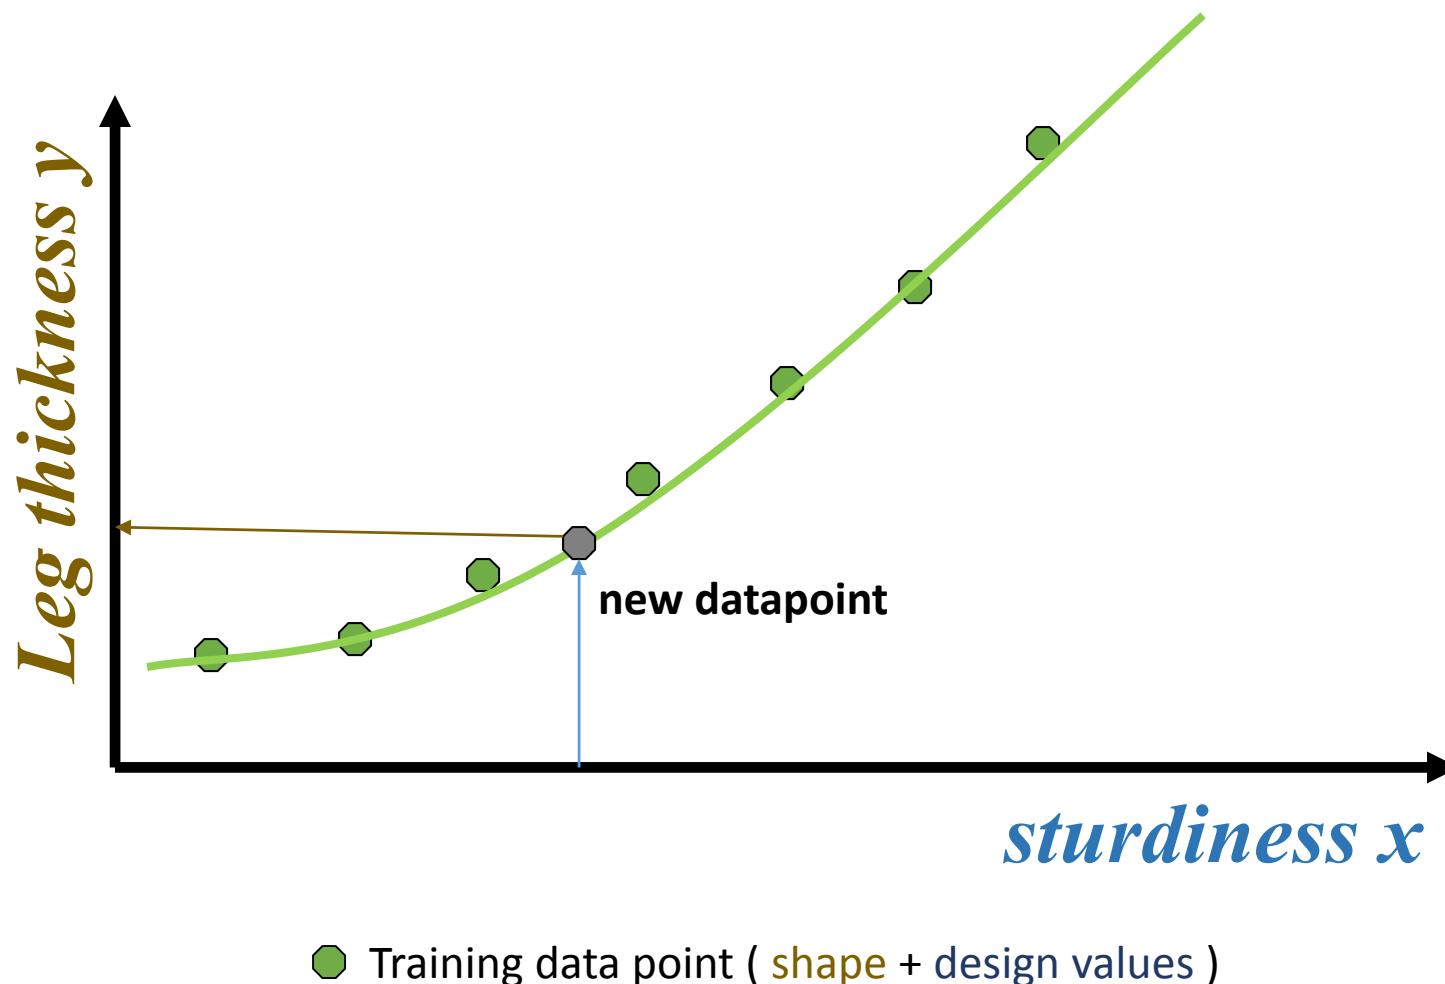
Learning basics: regression



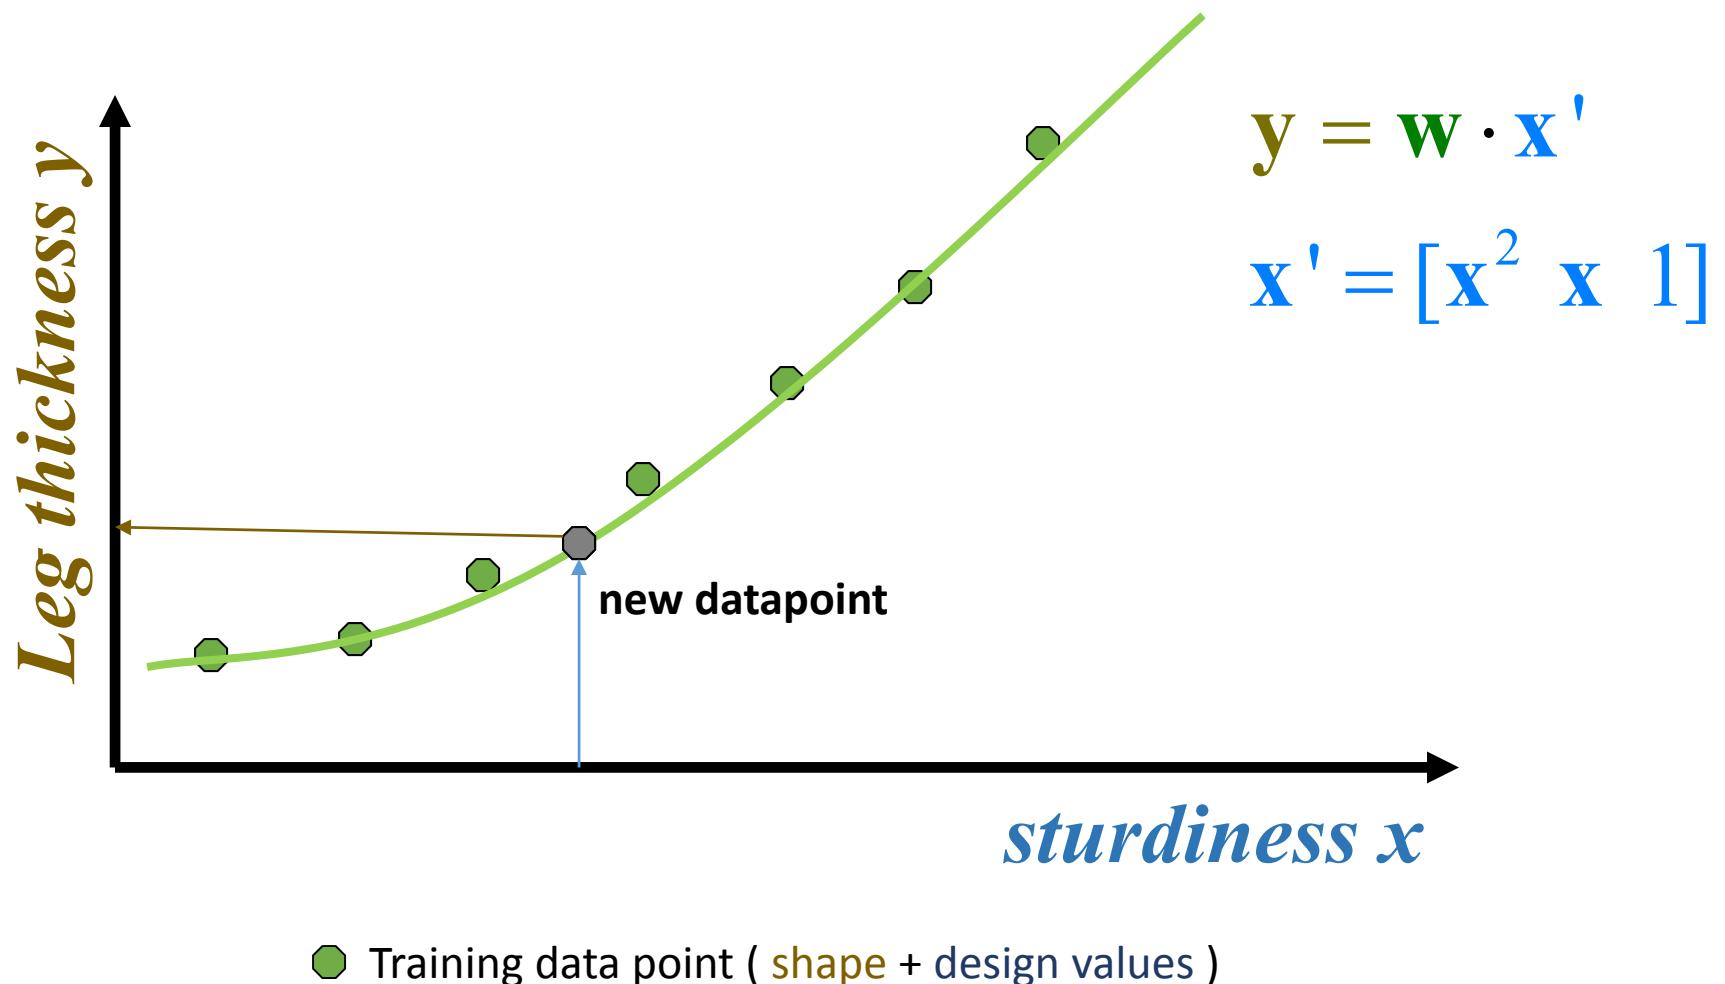
Learning basics: regression



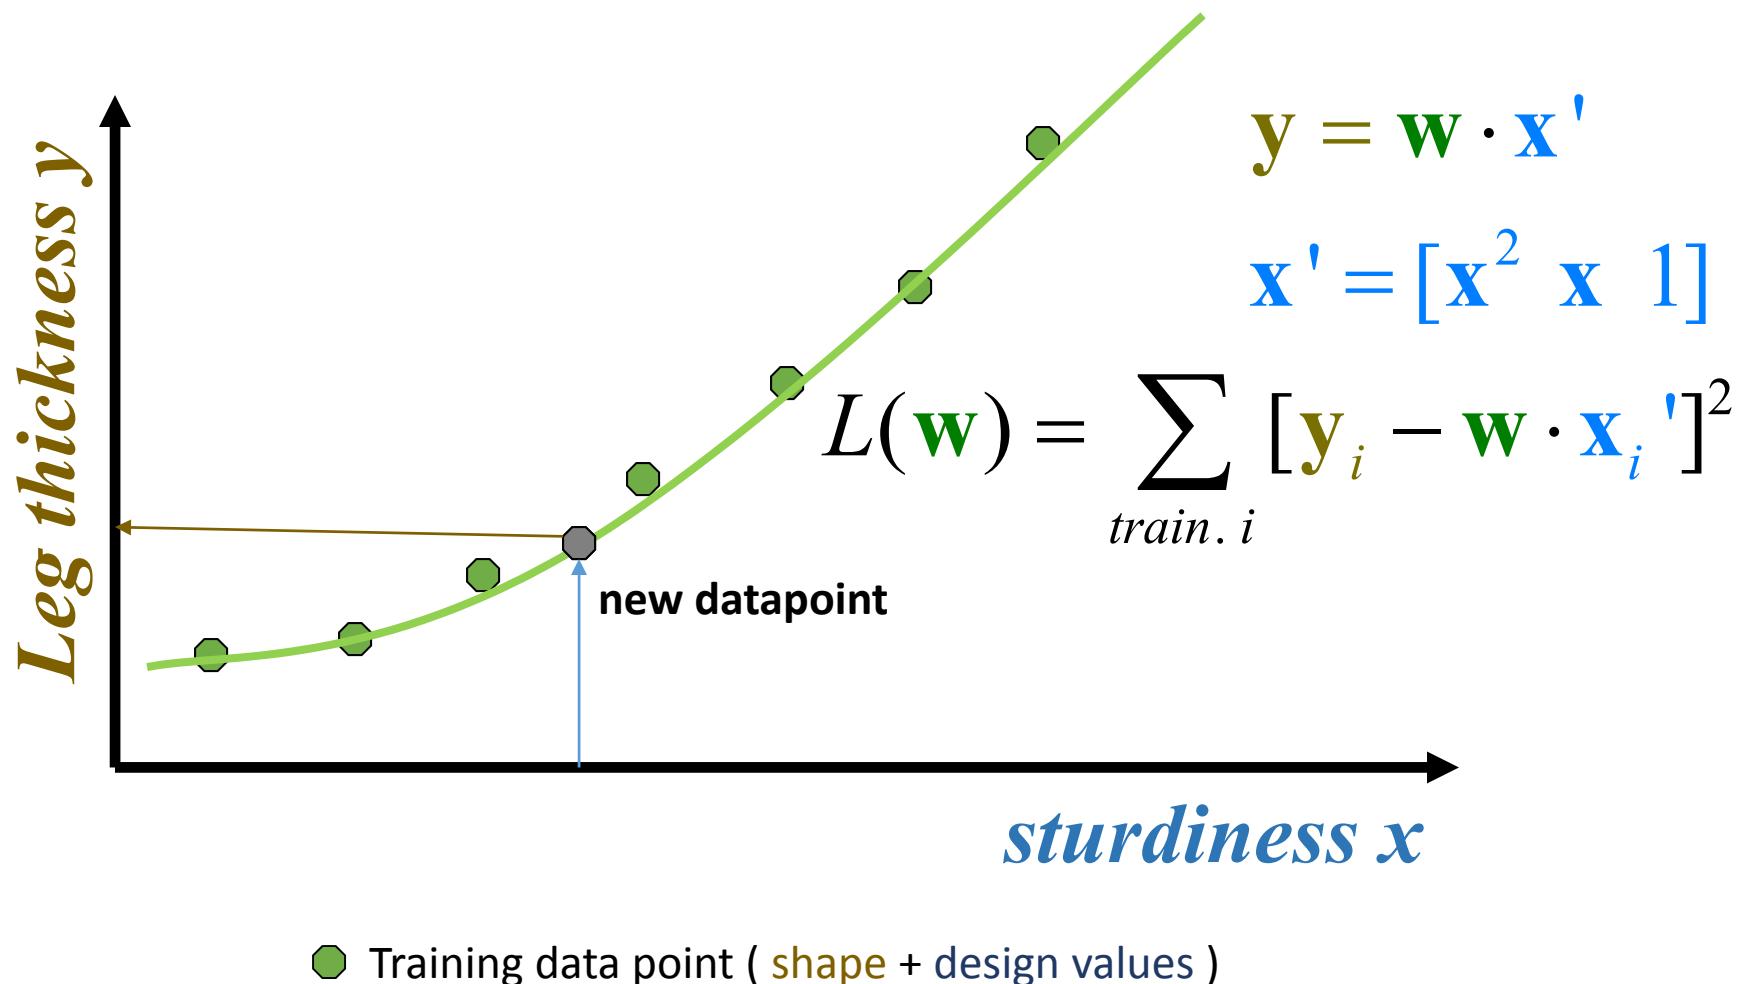
Learning basics: regression



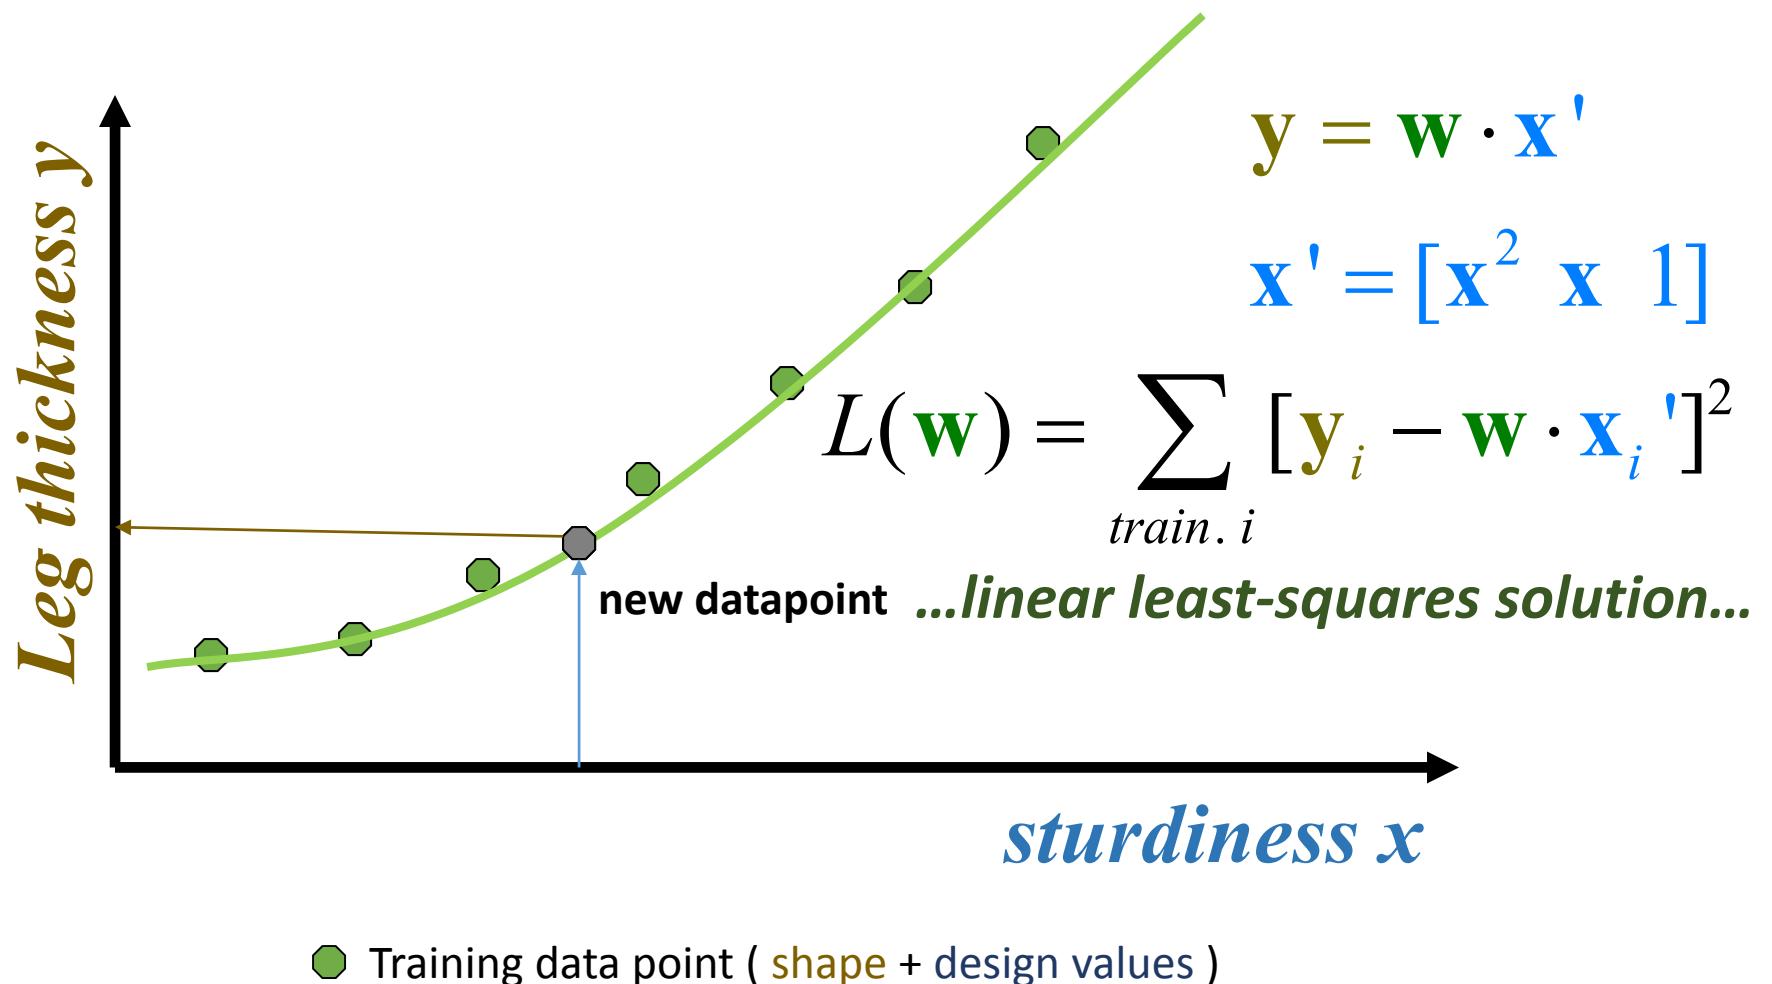
Learning basics: regression



Learning basics: regression



Learning basics: regression



Overfitting

Important to select a function that would **avoid overfitting & generalize** (produce reasonable outputs for inputs not encountered during training)

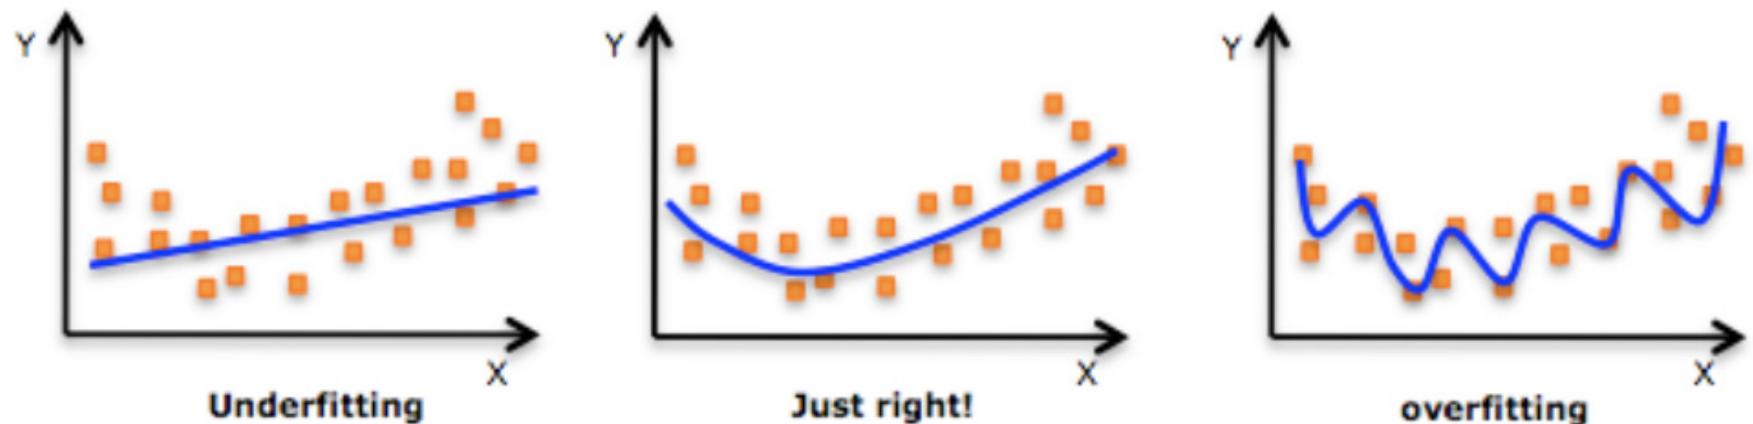


image from Andrew Ng's ML class

Learning basics: Logistic Regression

Suppose you want to predict **mug** or **no mug** for a shape.

Output: $y = 1$ [*coffee mug*], $y = 0$ [*no coffee mug*]

Input: $\mathbf{x} = \{x_1, x_2, \dots\}$ [curvature histograms, HKS etc]

Learning basics: Logistic Regression

Suppose you want to predict **mug** or **no mug** for a shape.

Output: $y = 1$ [*coffee mug*], $y = 0$ [*no coffee mug*]

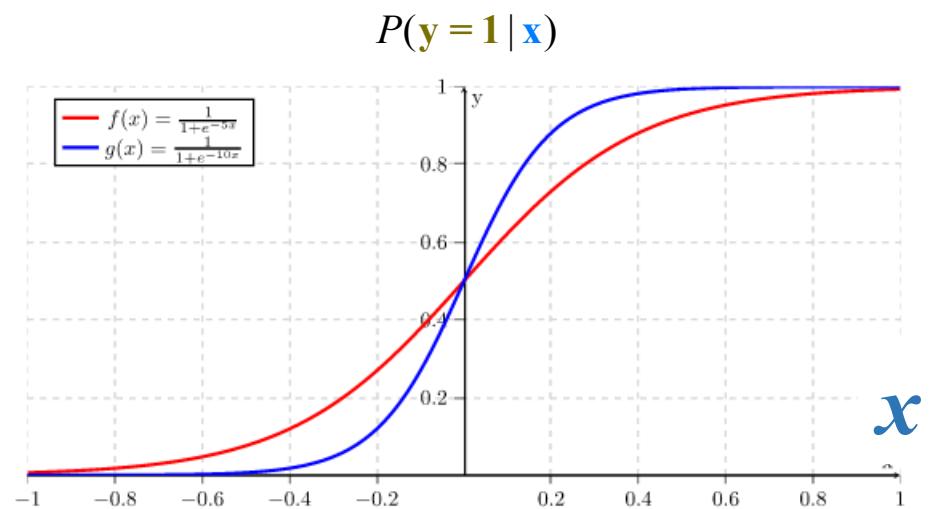
Input: $\mathbf{x} = \{x_1, x_2, \dots\}$ [curvature histograms, HKS etc]

Classification function:

$$P(y = 1 | \mathbf{x}) = \mathbf{f}(\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x})$$

where \mathbf{w} is a **weight vector**

$$\sigma(\mathbf{w} \cdot \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w} \cdot \mathbf{x})}$$



Logistic regression: training

Need to estimate parameters w from training data e.g., shapes of objects \mathbf{x}_i and given labels \mathbf{y}_i (mugs/no mugs) ($i=1\dots N$ training shapes)

Find parameters that **maximize probability of training data**

$$\max_w \prod_{i=1}^N P(\mathbf{y}_i = 1 | \mathbf{x}_i)^{[\mathbf{y}_i == 1]} [1 - P(\mathbf{y}_i = 1 | \mathbf{x}_i)]^{[\mathbf{y}_i == 0]}$$

Logistic regression: training

Need to estimate parameters w from training data e.g., shapes of objects x_i and given labels y_i (mugs/no mugs) ($i=1\dots N$ training shapes)

Find parameters that **maximize probability of training data**

$$\max_w \prod_{i=1}^N \sigma(w \cdot x_i)^{[y_i==1]} [1 - \sigma(w \cdot x_i)]^{[y_i==0]}$$

Logistic regression: training

Need to estimate parameters w from training data e.g., shapes of objects x_i and given labels y_i (mugs/no mugs) ($i=1\dots N$ training shapes)

Find parameters that **maximize the log prob. of training data**

$$\max_w \log \left\{ \prod_{i=1}^N \sigma(\mathbf{w} \cdot \mathbf{x}_i)^{[y_i==1]} [1 - \sigma(\mathbf{w} \cdot \mathbf{x}_i)]^{[y_i==0]} \right\}$$

Logistic regression: training

Need to estimate parameters \mathbf{w} from training data e.g., shapes of objects \mathbf{x}_i and given labels \mathbf{y}_i (mugs/no mugs) ($i=1\dots N$ training shapes)

Find parameters that **maximize the log prob. of training data**

$$\max_{\mathbf{w}} \sum_{i=1}^N [\mathbf{y}_i == 1] \log \sigma(\mathbf{w} \cdot \mathbf{x}_i) + [\mathbf{y}_i == 0] \log(1 - \sigma(\mathbf{w} \cdot \mathbf{x}_i))$$

Logistic regression: training

Need to estimate parameters w from training data e.g., shapes of objects x_i and given labels y_i (mugs/no mugs) ($i=1\dots N$ training shapes)

This is called **log-likelihood**

$$\max_w \sum_{i=1}^N [y_i = 1] \log \sigma(w \cdot x_i) + [y_i = 0] \log(1 - \sigma(w \cdot x_i))$$

Logistic regression: training

Need to estimate parameters w from training data e.g., shapes of objects x_i and given labels y_i (mugs/no mugs) ($i=1 \dots N$ training shapes)

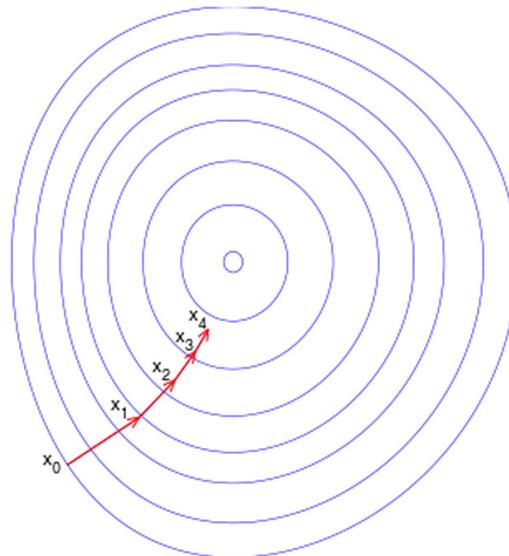
We have an **optimization problem**.

$$\max_w \sum_{i=1}^N [y_i = 1] \log \sigma(w \cdot x_i) + [y_i = 0] \log(1 - \sigma(w \cdot x_i))$$

$$\frac{\partial L(\mathbf{w})}{\partial w_d} = \sum_i x_{i,d} [y_i - \sigma(\mathbf{w} \cdot \mathbf{x}_i)]$$

(partial derivative for d^{th} parameter)

How can we minimize/maximize a function?



Gradient descent: Given a random initialization of parameters and a step rate η , update them according to:

$$\mathbf{w}_{new} = \mathbf{w}_{old} - \eta \nabla L(\mathbf{w})$$

See also **quasi-Newton** and **IRLS** methods

Regularization

Overfitting:

few training data vs large number of parameters!

Penalize large weights:

$$\min_{\mathbf{w}} -L(\mathbf{w}) + \lambda \sum_d \mathbf{w}_d^2$$

Called **ridge regression (or L2 regularization)**

Regularization

Overfitting:

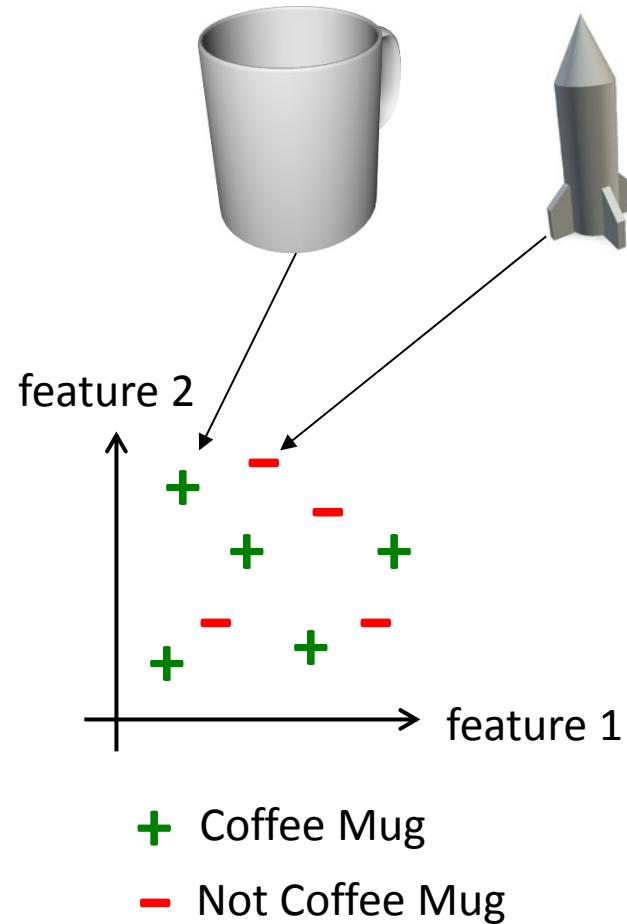
few training data vs large number of parameters!

Penalize non-zero weights - push as many as possible to **0**:

$$\min_{\mathbf{w}} -L(\mathbf{w}) + \lambda \sum_d |\mathbf{w}_d|$$

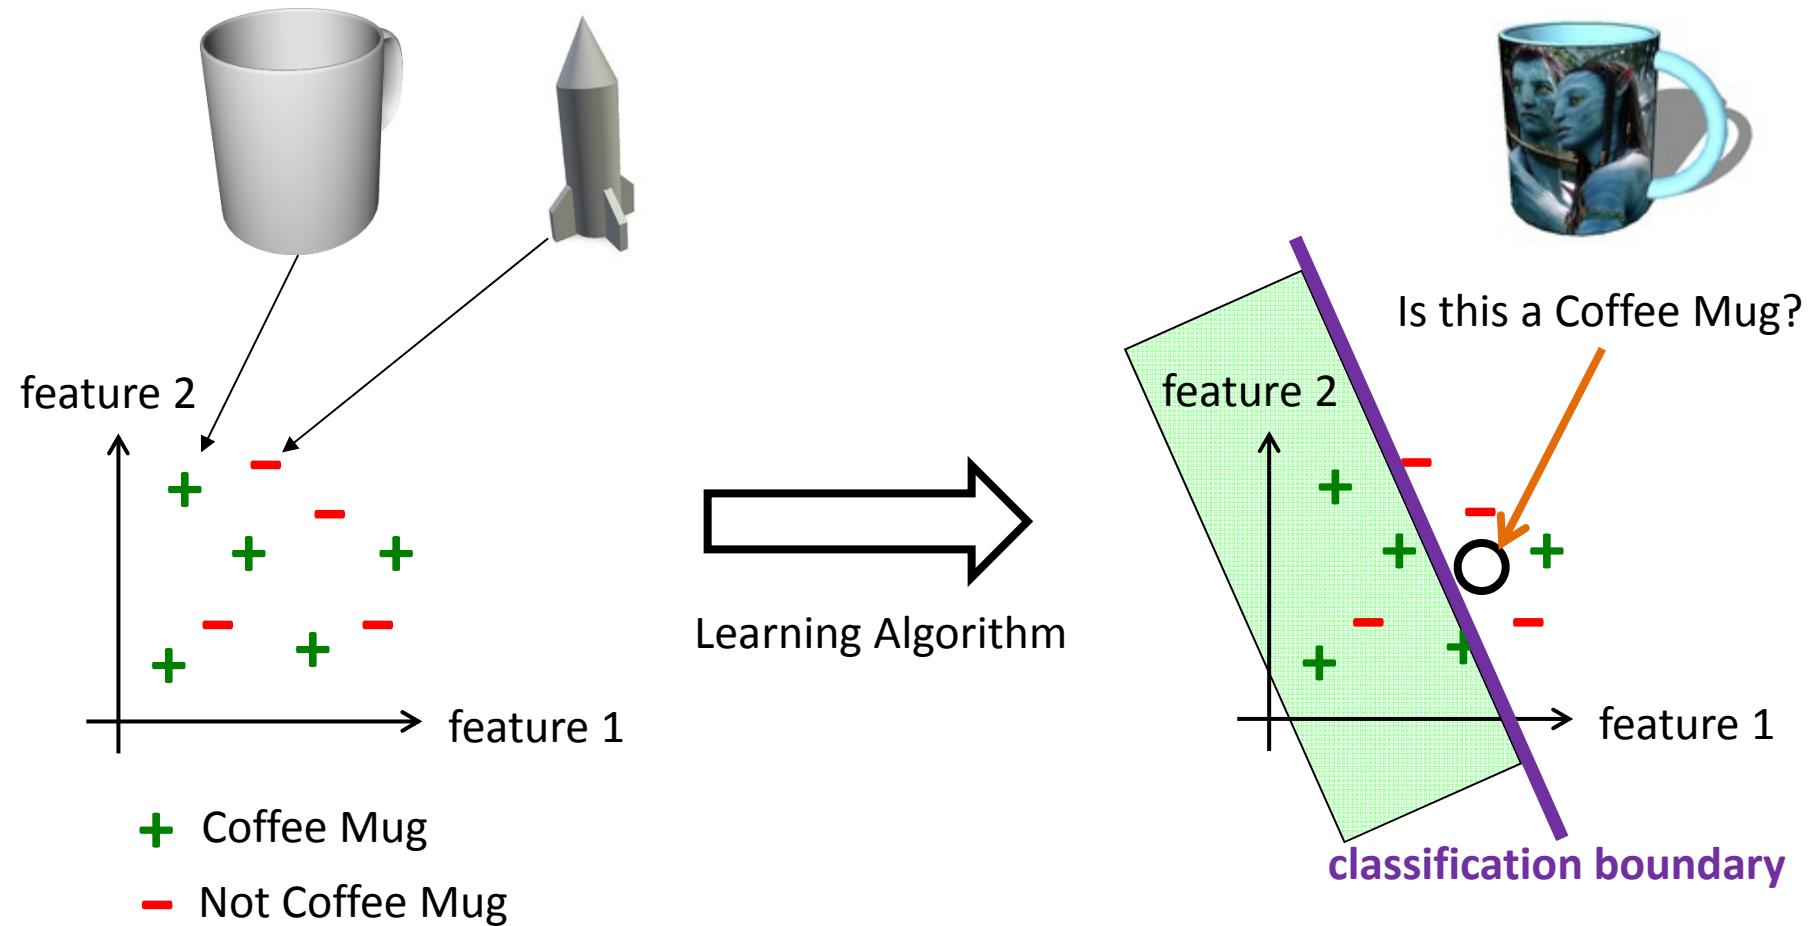
Called **Lasso** (or L1 regularization)

The importance of choosing good features...



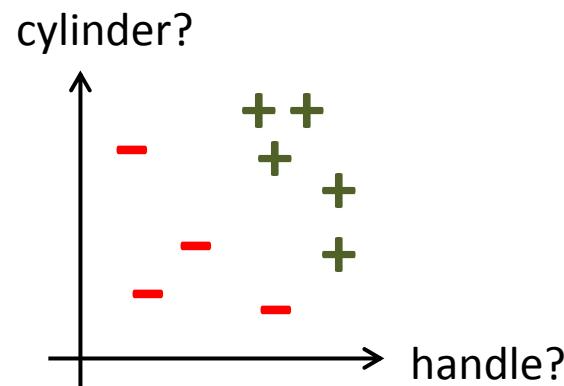
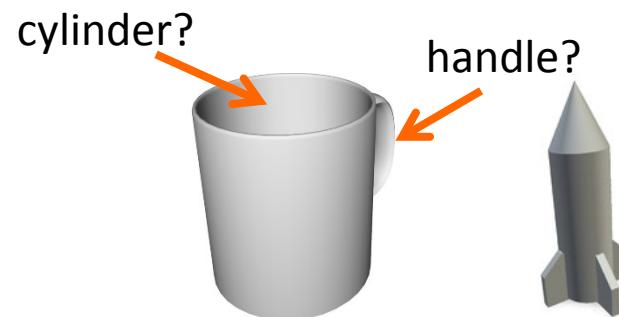
*modified slides originally
by Adam Coates*

The importance of choosing good features...



*modified slides originally
by Adam Coates*

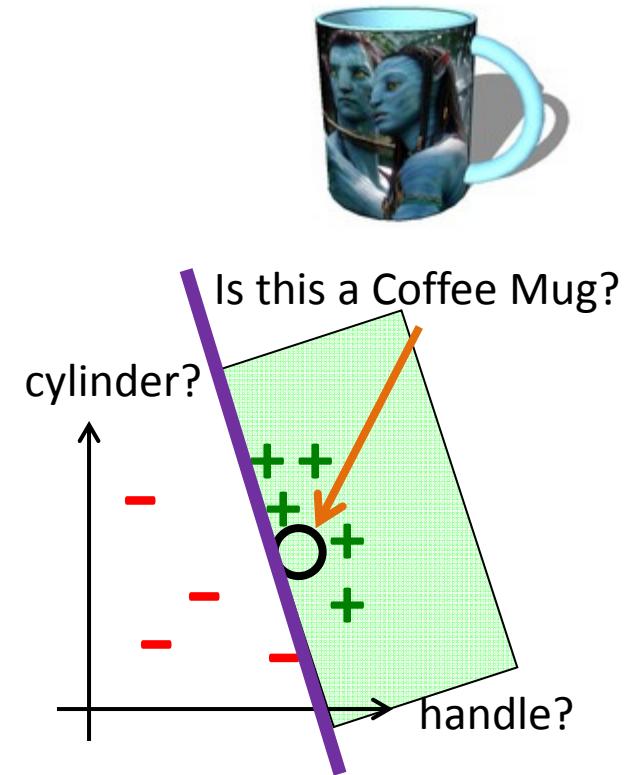
The importance of choosing good features...



+ Coffee Mug

- Not Coffee Mug

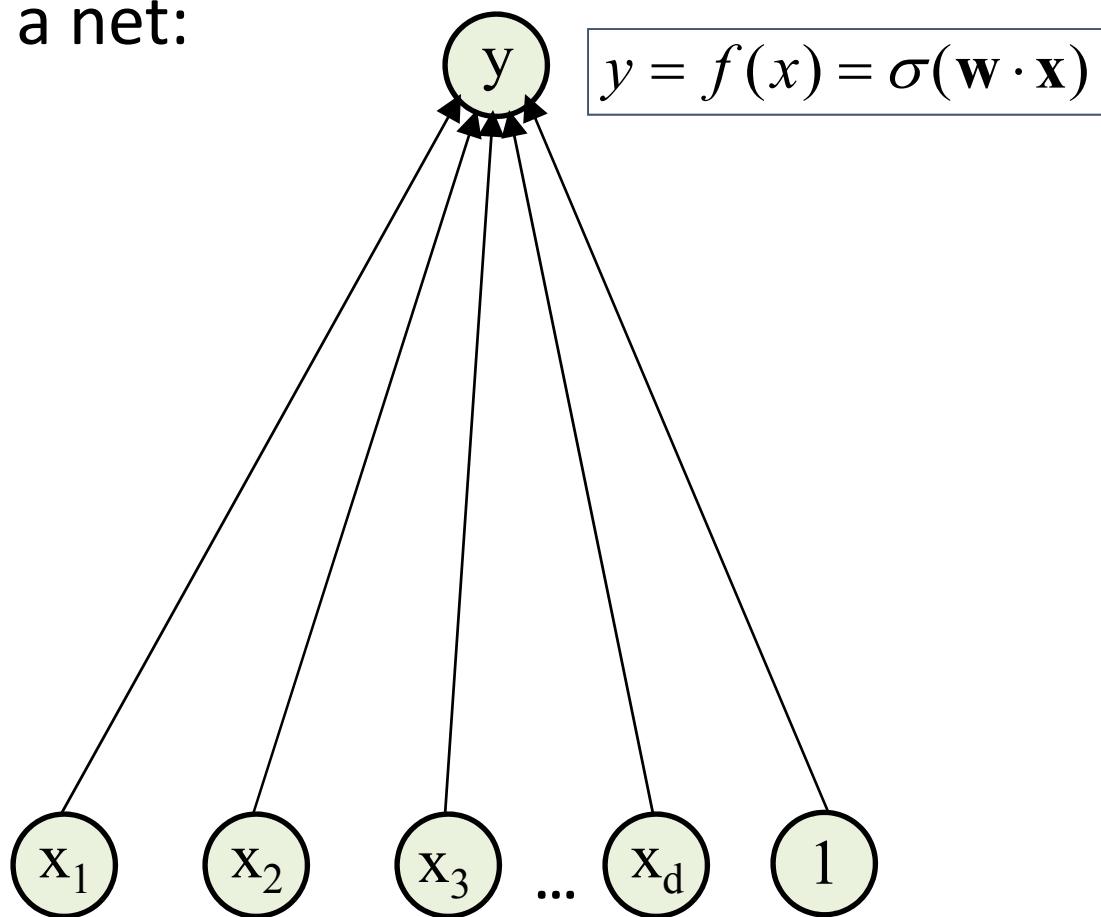
Learning Algorithm



*modified slides originally
by Adam Coates*

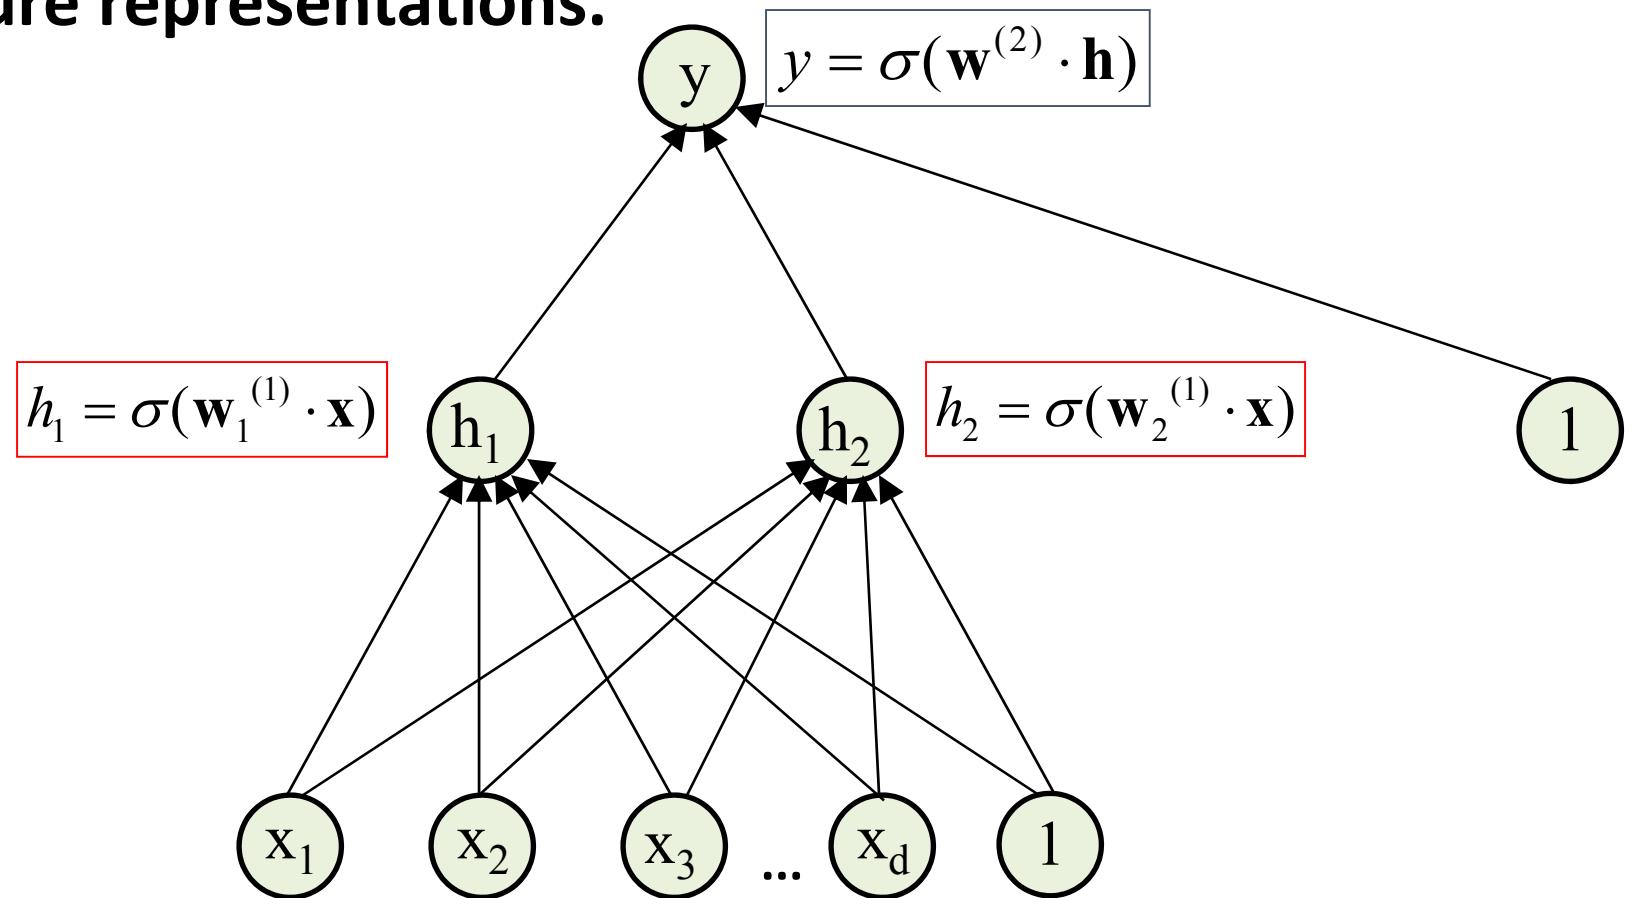
From “shallow” to “deep” mappings

Logistic regression: output is a **direct function of inputs**.
Think of it as a net:



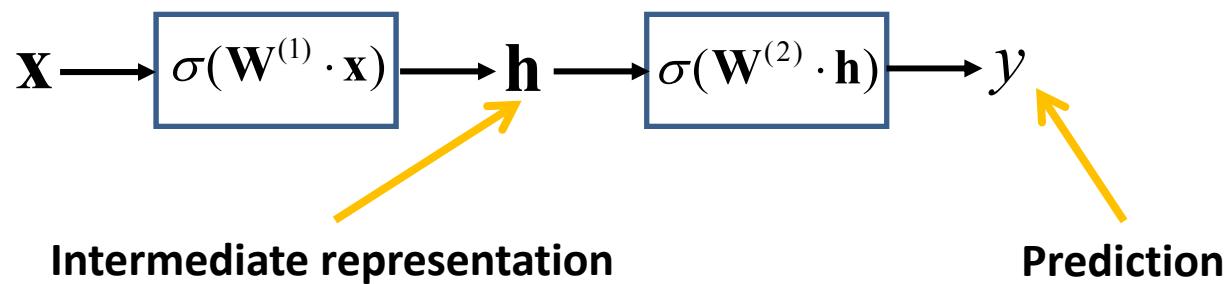
Neural network

Introduce latent nodes that play the role of **learned feature representations**.



Neural network

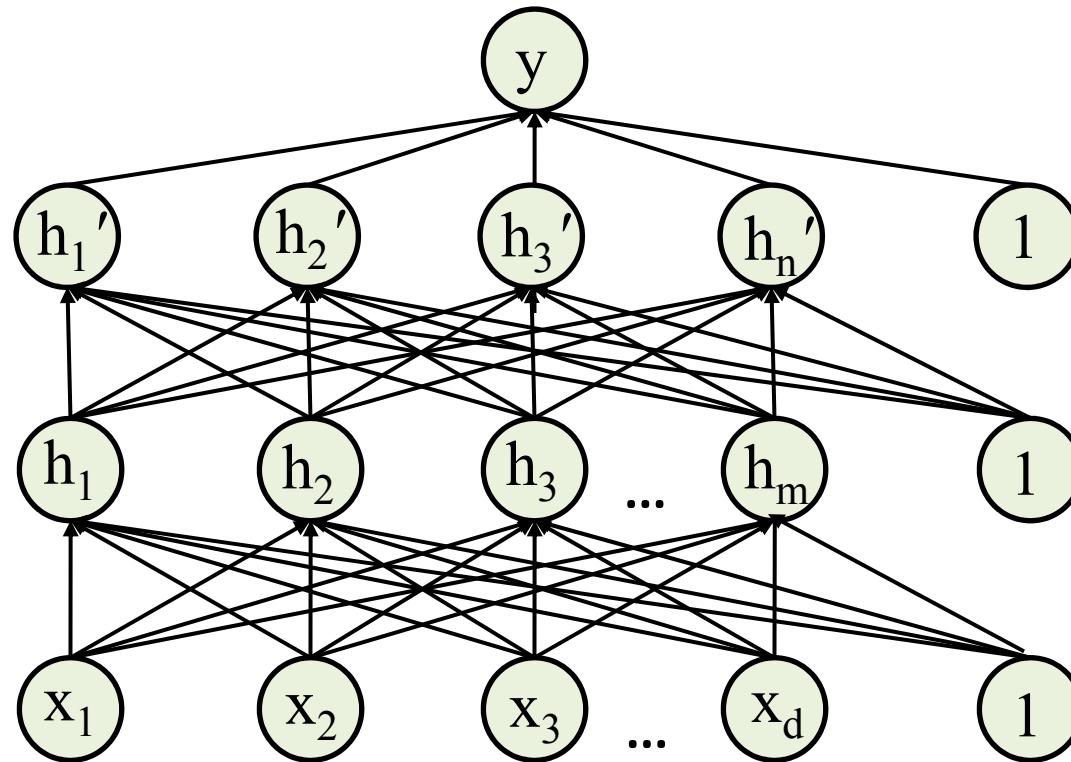
Same as logistic regression but now our output function has **multiple stages** ("layers", "modules").



where $\mathbf{W}^{(\cdot)} = \begin{bmatrix} \mathbf{w}_1^{(\cdot)} \\ \mathbf{w}_2^{(\cdot)} \\ \dots \\ \mathbf{w}_m^{(\cdot)} \end{bmatrix}$

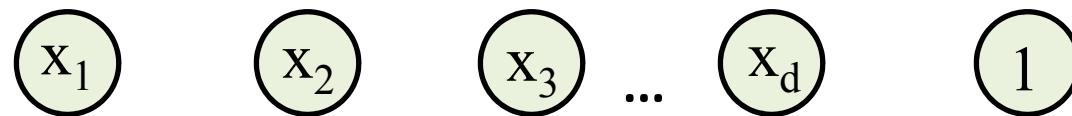
Neural network

Stack up several layers:



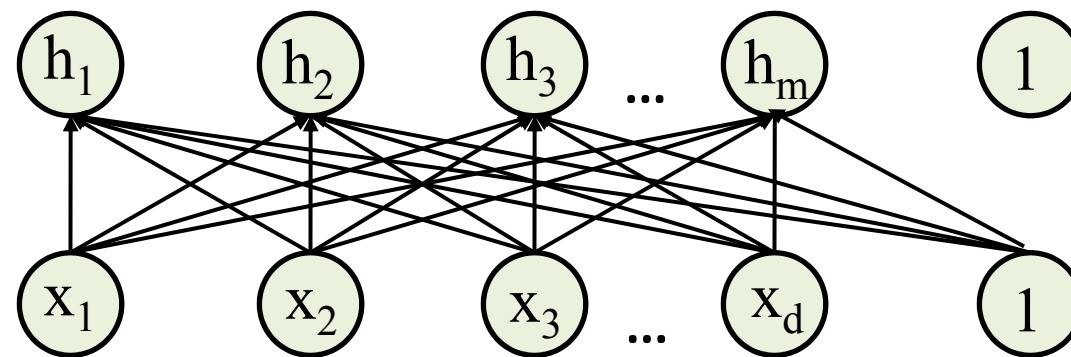
Forward propagation

Process to compute output:



Forward propagation

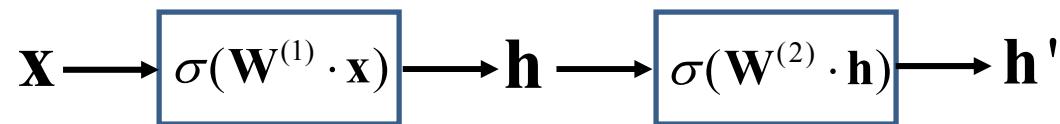
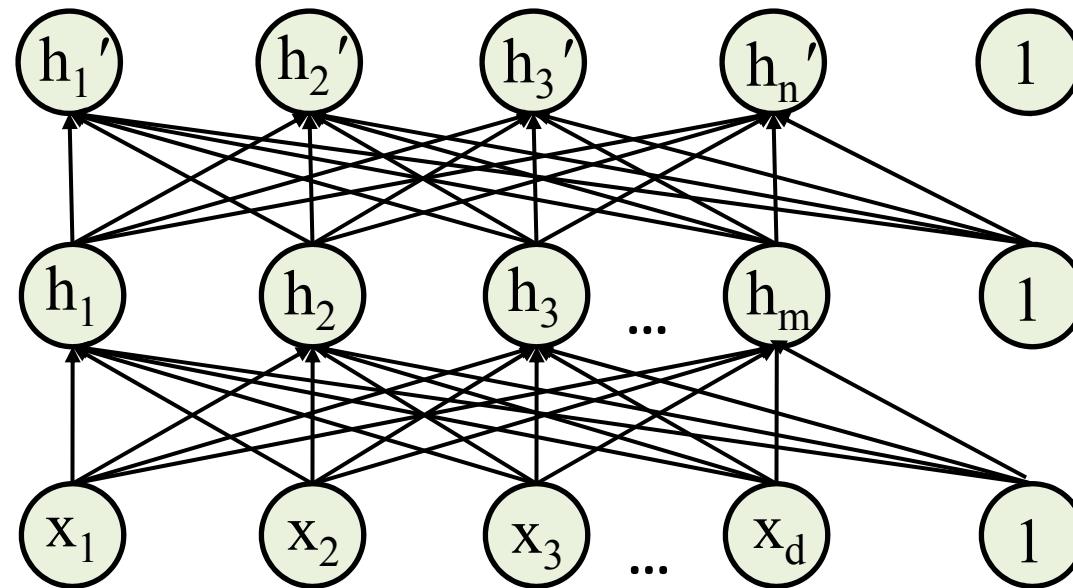
Process to compute output:



$$\mathbf{x} \rightarrow \sigma(\mathbf{W}^{(1)} \cdot \mathbf{x}) \rightarrow \mathbf{h}$$

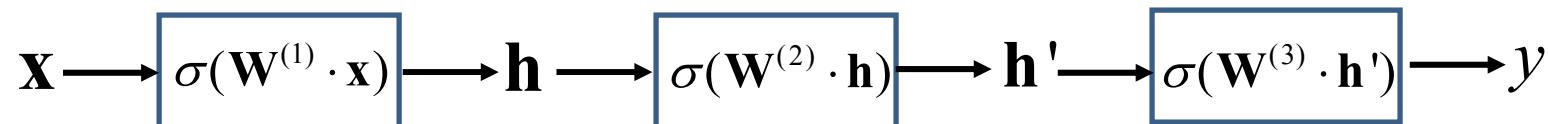
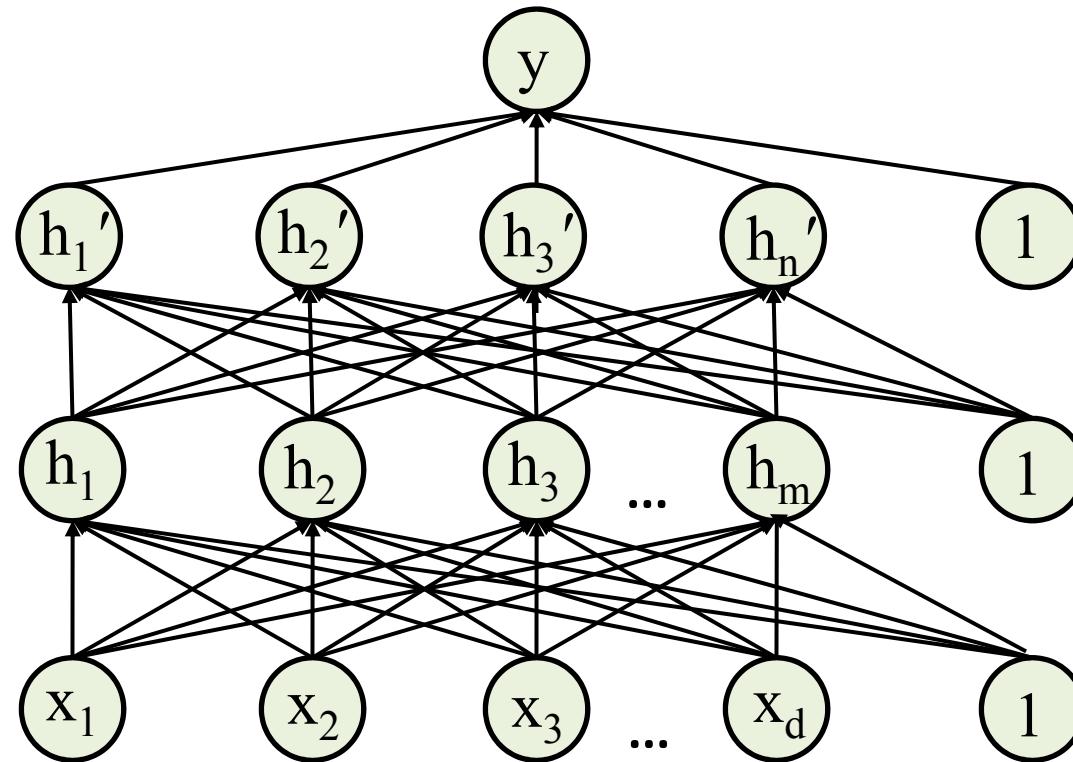
Forward propagation

Process to compute output:

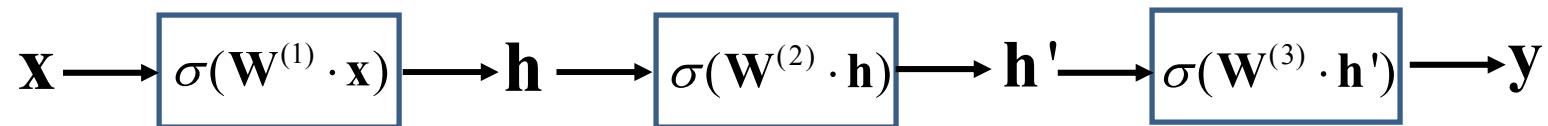
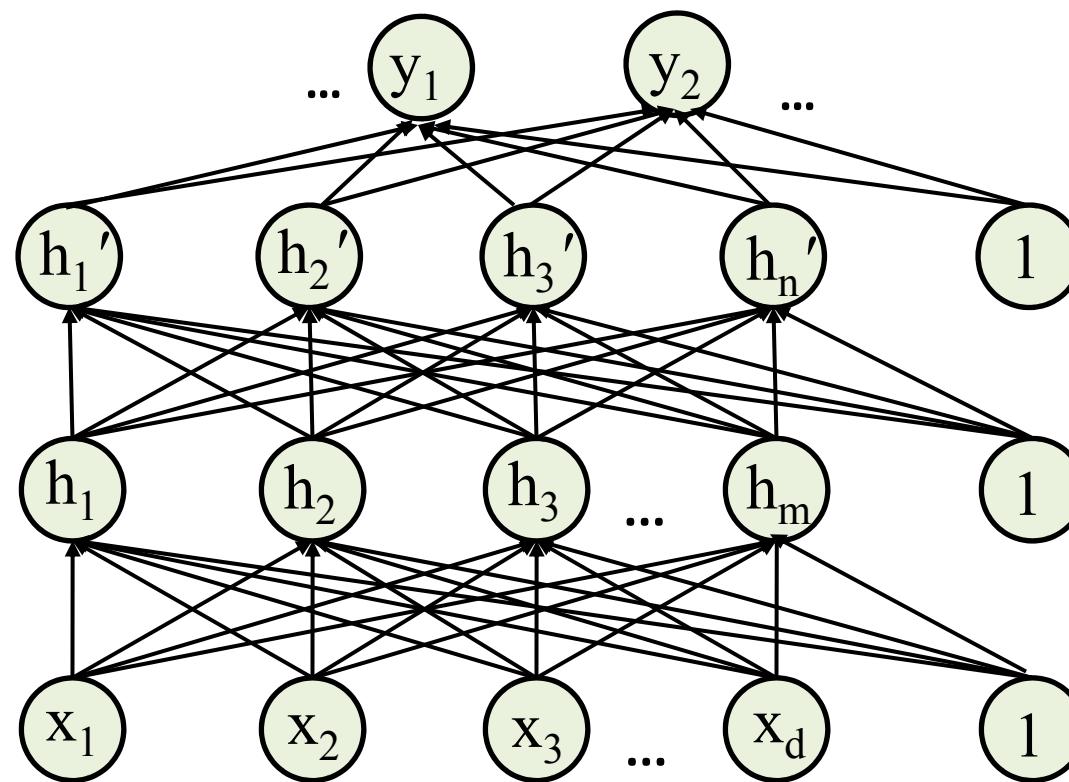


Forward propagation

Process to compute output:



Multiple outputs



How can we learn the parameters?

Use a loss function e.g., for classification:

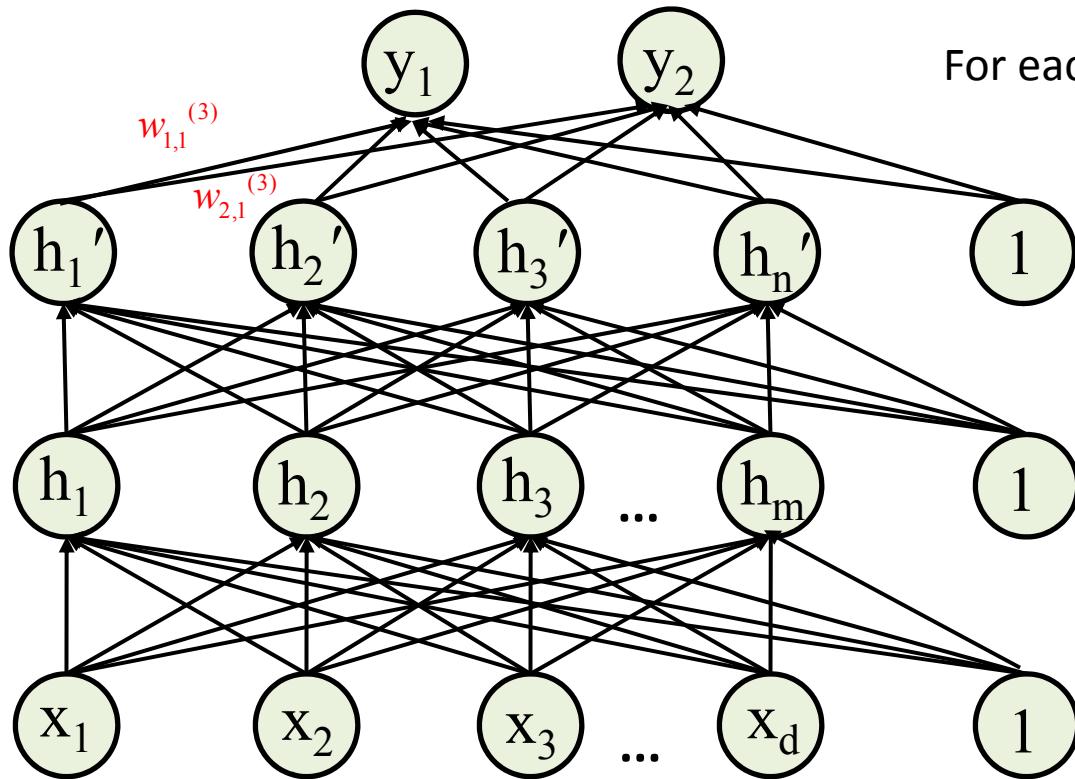
$$L(\mathbf{w}) = - \sum_i \sum_{output t} [\mathbf{y}_{i,t} == 1] \log f_t(\mathbf{x}_i) + [\mathbf{y}_{i,t} == 0] \log(1 - f_t(\mathbf{x}_i))$$

In case of regression i.e., for predicting continuous outputs:

$$L(\mathbf{w}) = \sum_i \sum_{output t} [\mathbf{y}_{i,t} - f_t(\mathbf{x}_i)]^2$$

Backpropagation

For each training example i :

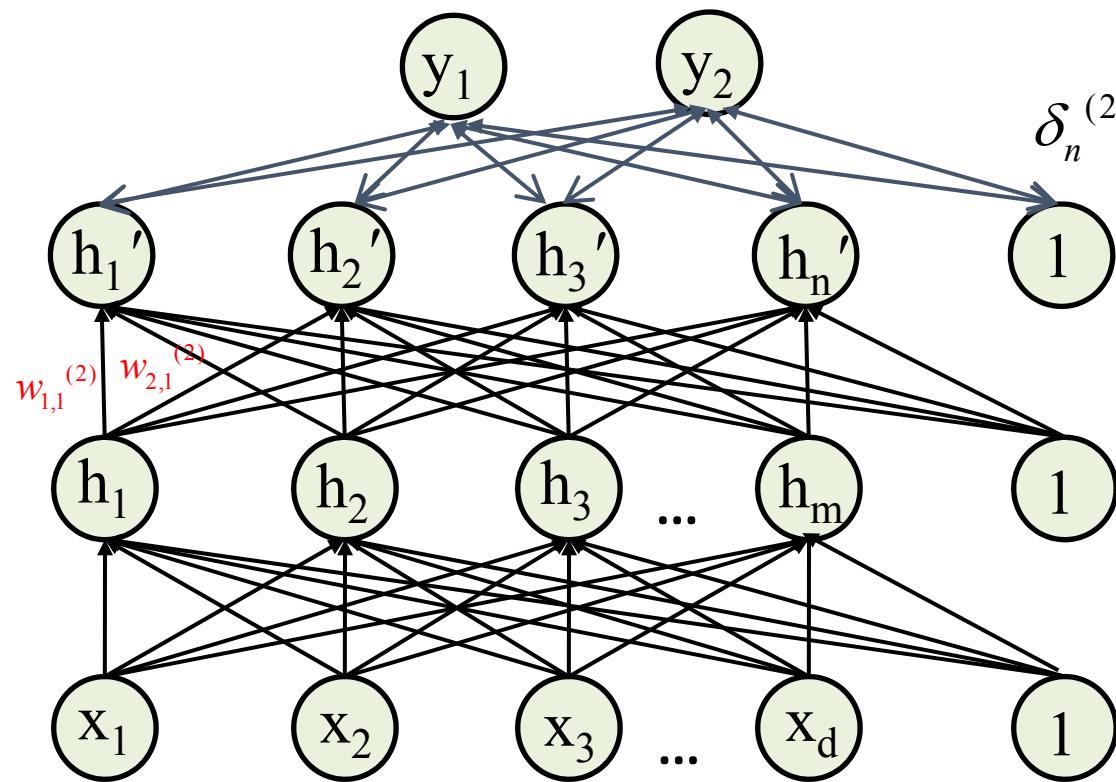


For each output:

$$\delta_t^{(3)} = y_t - f(\mathbf{x})$$

$$\frac{\partial L(\mathbf{w})}{\partial w_{t,n}^{(3)}} = \delta_t^{(3)} h_n$$

Backpropagation

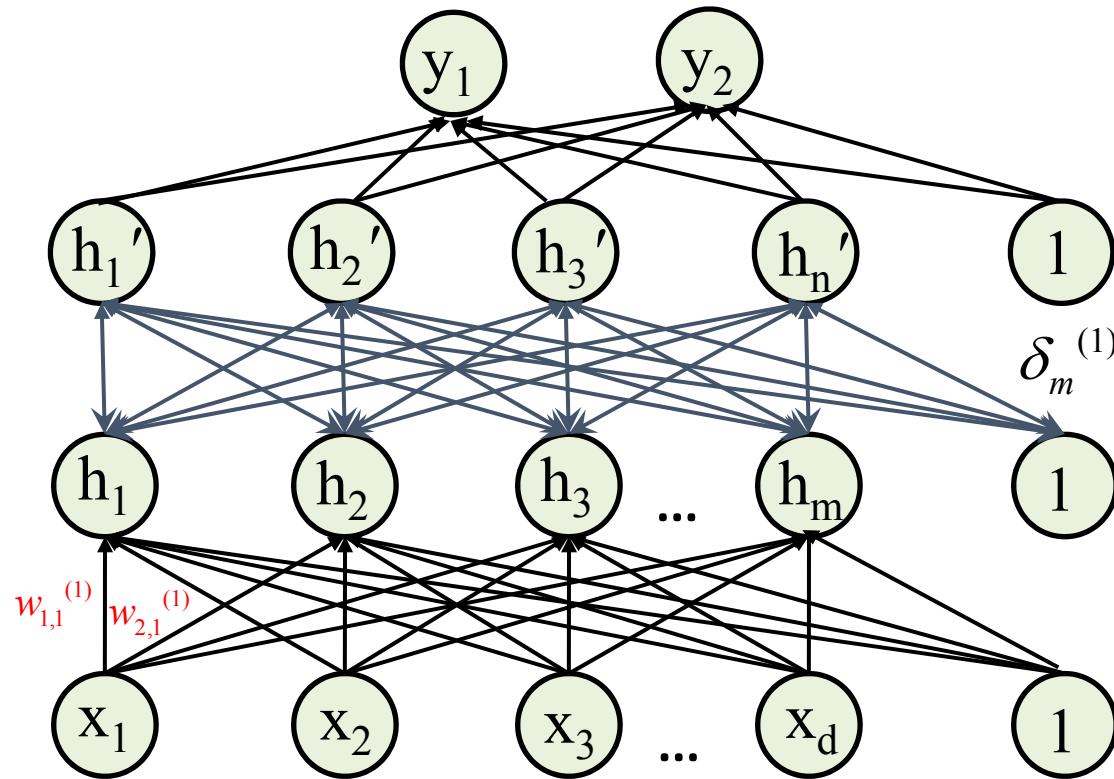


$$\delta_n^{(2)} = \sigma'(\mathbf{w}_n^{(2)} \cdot \mathbf{h}) \sum_t w_{t,n}^{(3)} \delta_t^{(3)}$$

Note: $\sigma'(\cdot) = \sigma(\cdot)[1 - \sigma(\cdot)]$

$$\frac{\partial L(\mathbf{w})}{\partial w_{n,m}^{(2)}} = \delta_n^{(2)} h_m$$

Backpropagation



$$\delta_m^{(1)} = \sigma'(\mathbf{w}_m^{(1)} \cdot \mathbf{x}) \sum_n w_{n,m}^{(2)} \delta_n^{(2)}$$

$$\frac{\partial L(\mathbf{w})}{\partial w_{m,d}^{(1)}} = \delta_m^{(1)} x_d$$

Is this magic?

All these are derivatives derived analytically using the
chain rule!

Gradient descent is expressed through **backpropagation**
of messages δ following the structure of the model

Training algorithm

For each training example [in a batch]

1. **Forward propagation** to compute outputs per layer
2. **Back propagate** messages δ from top to bottom layer
3. Multiply messages δ with inputs to compute **derivatives per layer**
4. **Accumulate the derivatives** from that training example

Apply the gradient descent rule

Yet, this does not work so easily...

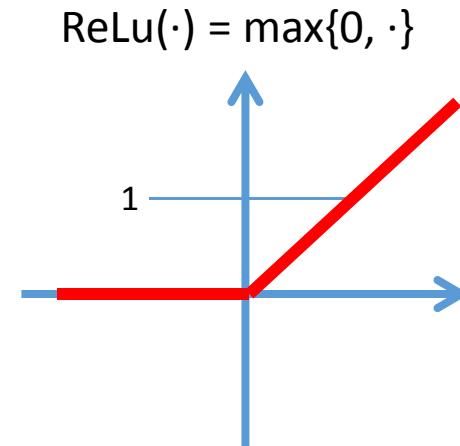
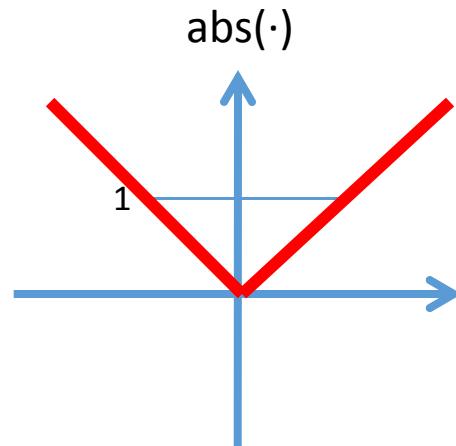
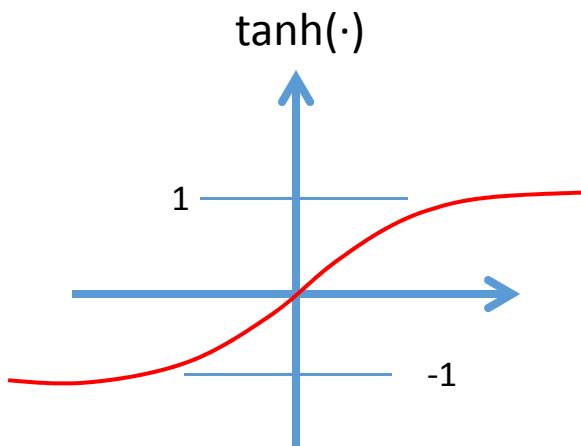


Yet, this does not work so easily...

- **Non-convex:** Local minima; convergence criteria.
- Optimization becomes difficult with **many layers**.
- Hard to diagnose and **debug malfunctions**.
- **Many things turn out to matter:**
 - Choice of nonlinearities.
 - Initialization of parameters.
 - Optimizer parameters: step size, schedule.

Non-linearities

- **Choice of functions inside network matters.**
 - Sigmoid function yields highly non-convex loss functions
 - Some other choices often used:

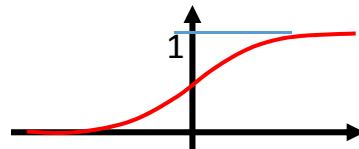


“Rectified Linear Unit”
→ Most popular.

[Nair & Hinton, 2010]

Initialization

- **Usually small random values.**
 - Try to choose so that typical input to a neuron avoids saturating



- **Initialization schemes for weights used as input to a node:**
 - tanh units: Uniform[-r, r]; sigmoid: Uniform[-4r, 4r].
 - See [[Glorot et al., AISTATS 2010](#)]

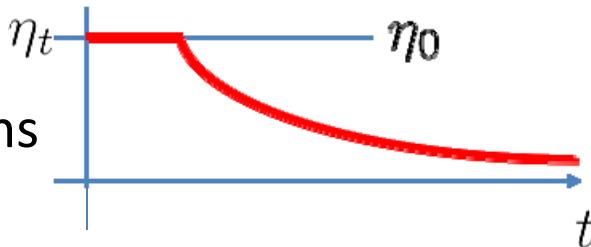
$$r = \sqrt{6 / (\text{fan-in} + \text{fan-out})}$$

Step size

- **Fixed step-size**
 - try many, choose the best...
 - pick size with least test error on a validation set after T iterations

- **Dynamic step size**

- decrease after T iterations



- if simply the objective is not decreasing much, cut step by half

Momentum

Modify stochastic/batch gradient descent:

Before : $\Delta\mathbf{w} = \eta \nabla_{\mathbf{w}} L(\mathbf{w})$, $w = w - \Delta\mathbf{w}$

With momentum : $\Delta\mathbf{w} = \mu \Delta\mathbf{w}_{previous} + \eta \nabla_{\mathbf{w}} L(\mathbf{w})$, $w = w - \Delta\mathbf{w}$

“Smooth” estimate of gradient from several steps of gradient descent:

- High-curvature directions cancel out, low-curvature directions “add up” and accelerate.
- **Other techniques: Adagrad, Adadelta, batch normalization...**

Momentum+L2 regularization

Modify stochastic/batch gradient descent:

Before : $\Delta\mathbf{w} = \eta \nabla_{\mathbf{w}} L(\mathbf{w})$, $w = w - \Delta\mathbf{w}$

With momentum : $\Delta\mathbf{w} = \mu \Delta\mathbf{w}_{previous} + \eta \nabla_{\mathbf{w}} L(\mathbf{w})$, $w = w - \Delta\mathbf{w}$

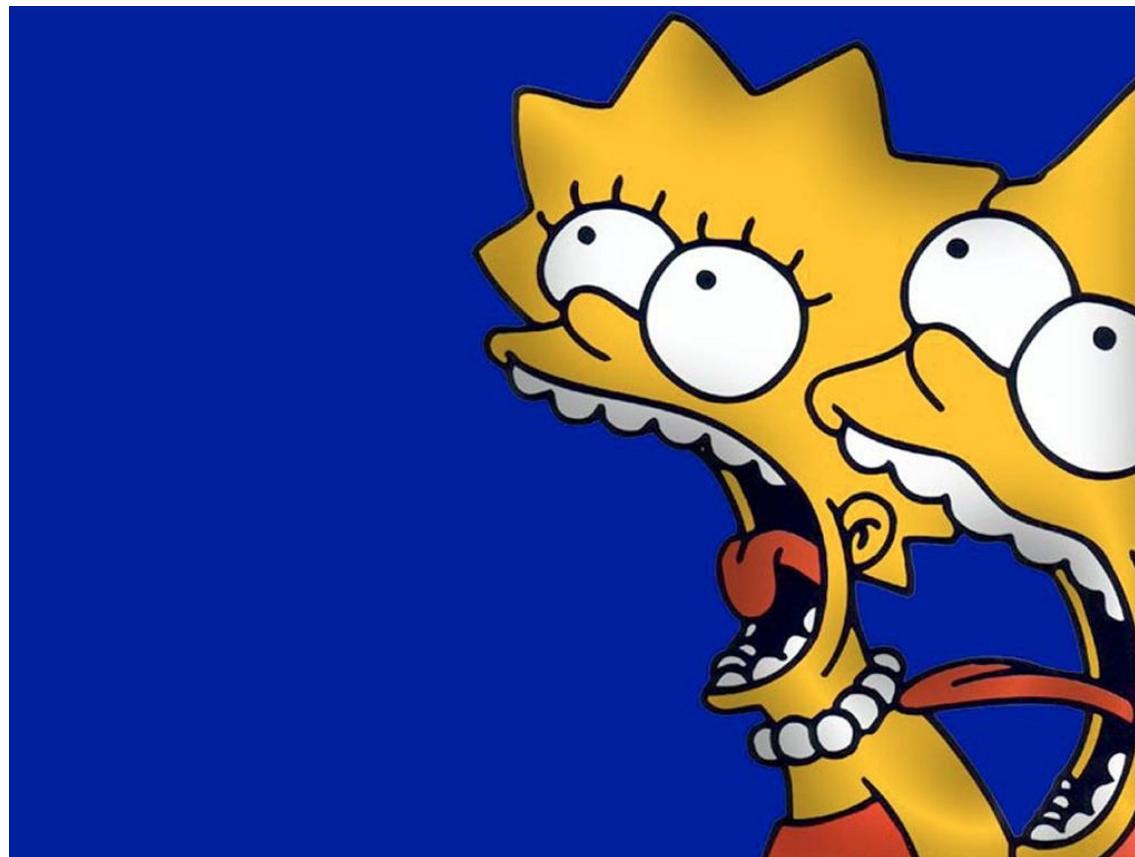
“Smooth” estimate of gradient from several steps of gradient descent:

- High-curvature directions cancel out, low-curvature directions “add up” and accelerate.
- **Other techniques: Adagrad, Adadelta, batch normalization...**

Add **L2 regularization** to the loss function:

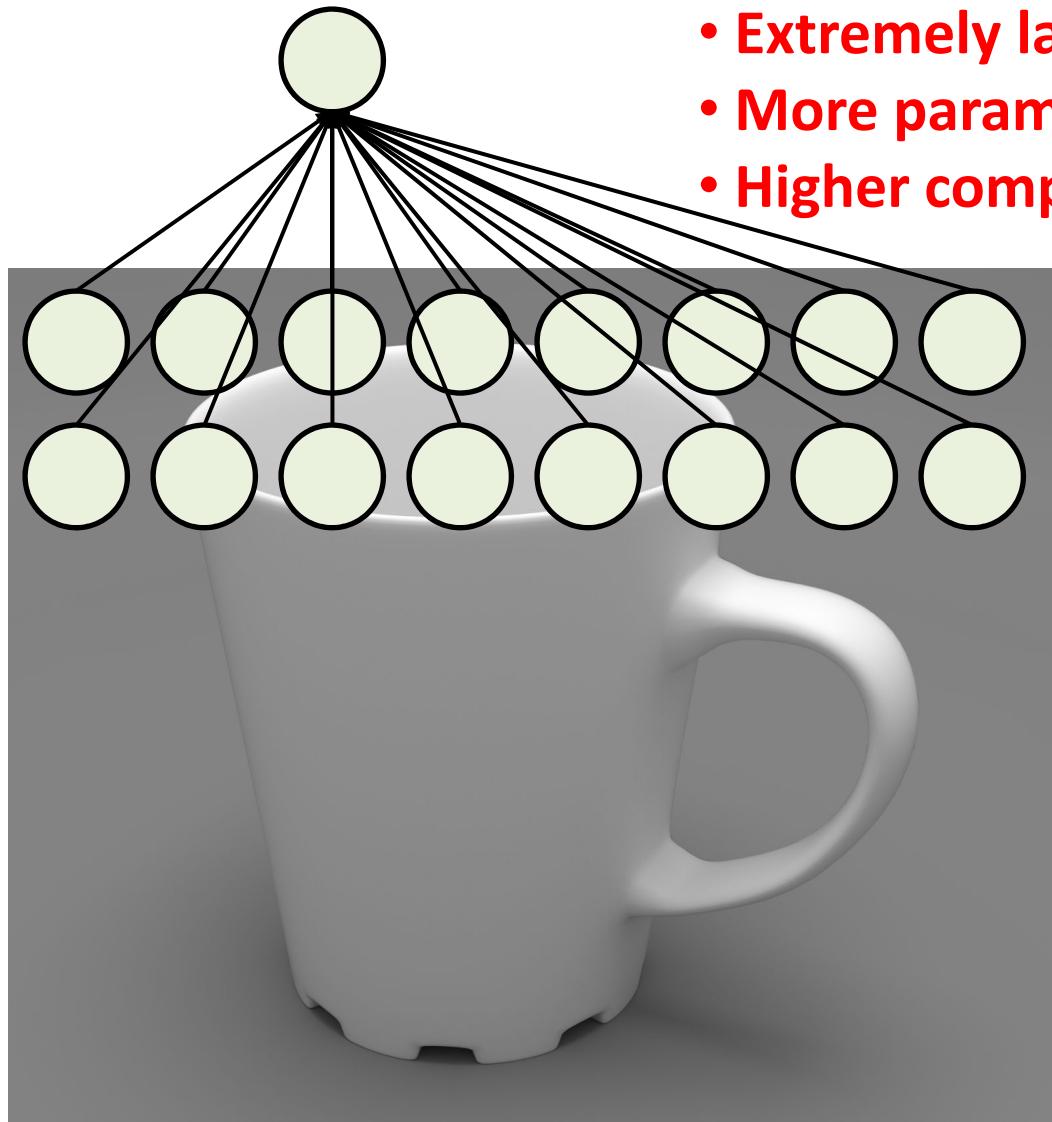
$$\Delta\mathbf{w} = \eta \nabla_{\mathbf{w}} (L(\mathbf{w}) + \lambda \|\mathbf{w}\|_2)$$

Yet, things will not still work well!



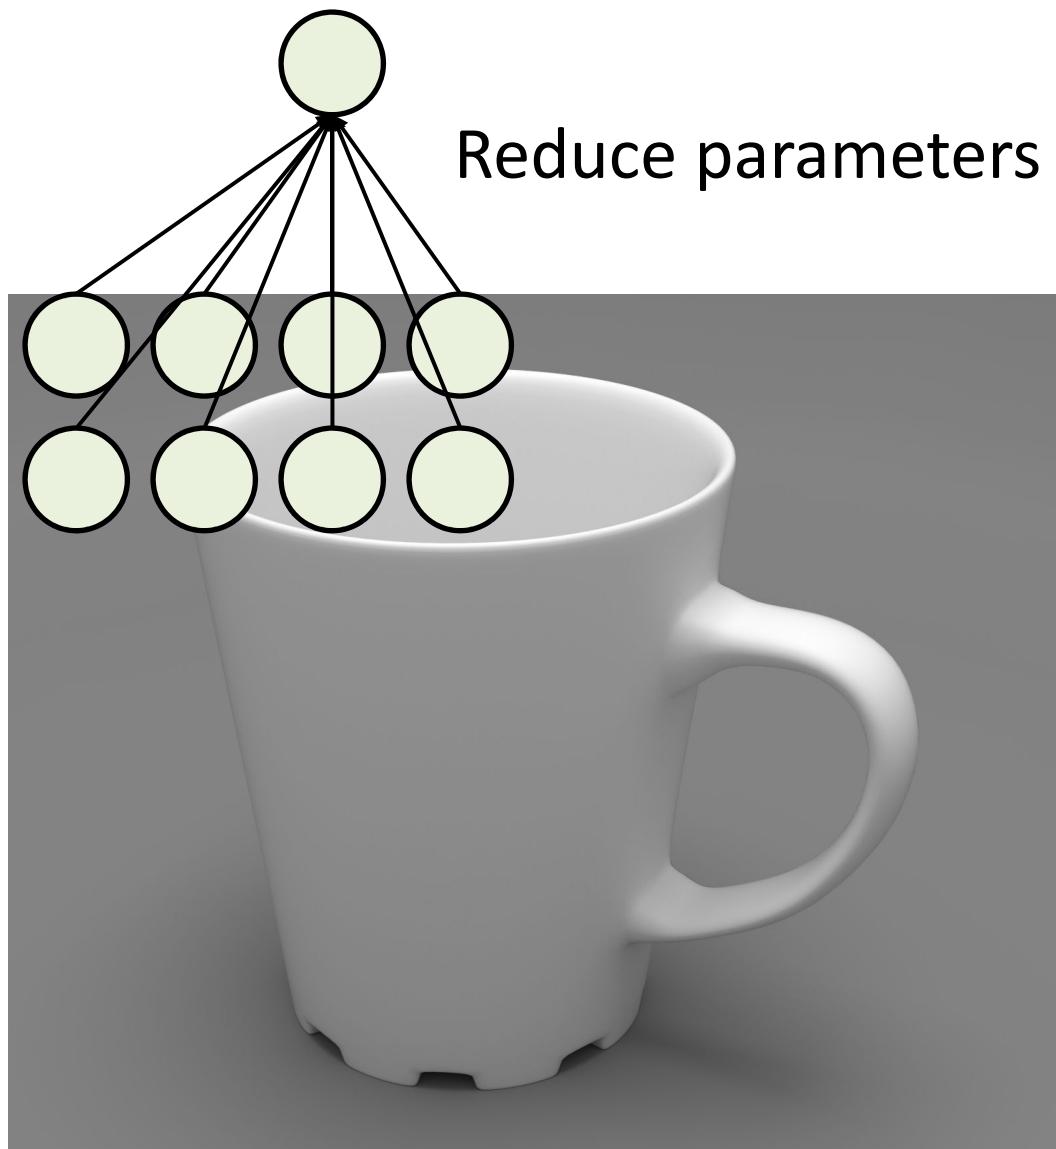
Main problem

- **Extremely large number of connections.**
- **More parameters to train.**
- **Higher computational expense.**

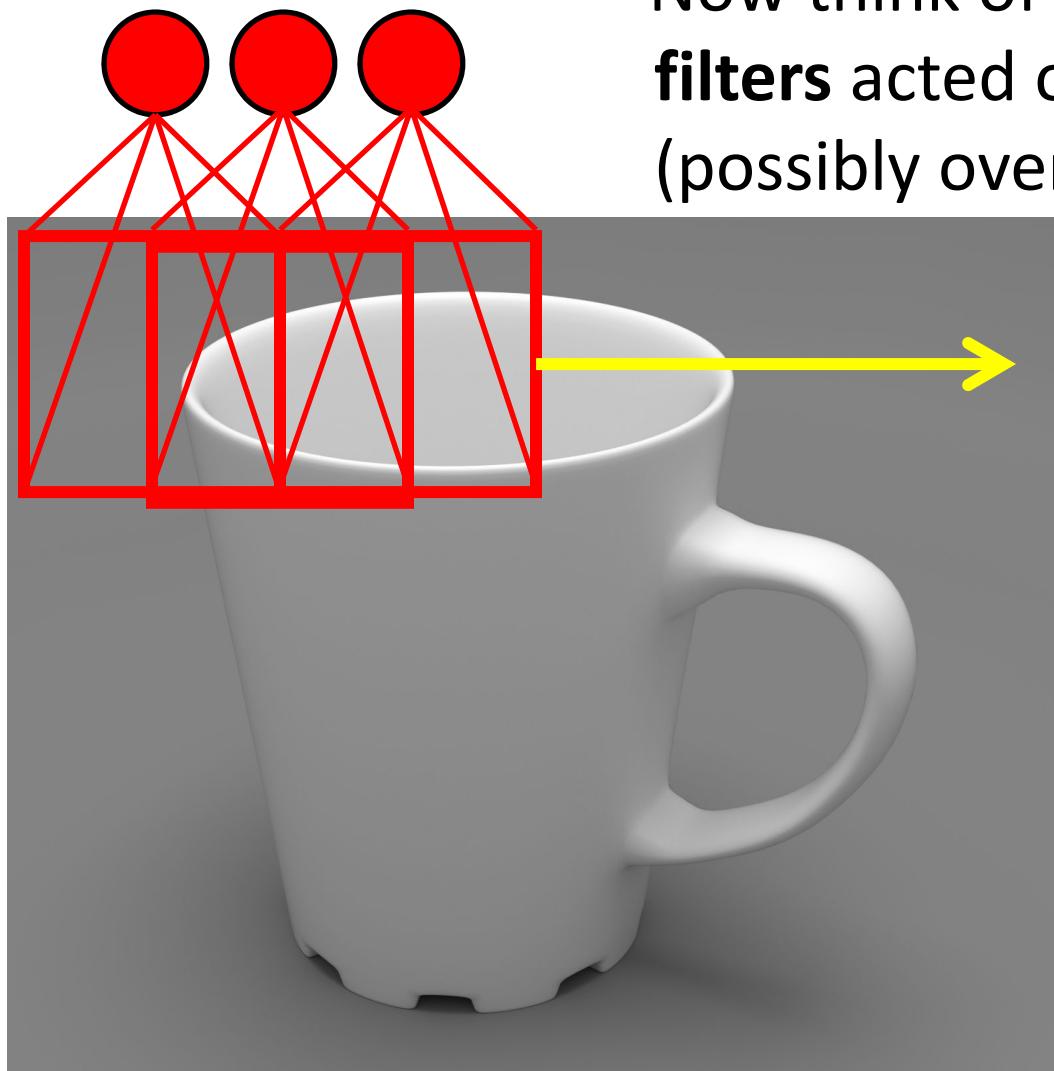


Local connectivity

Reduce parameters with **local connections!**



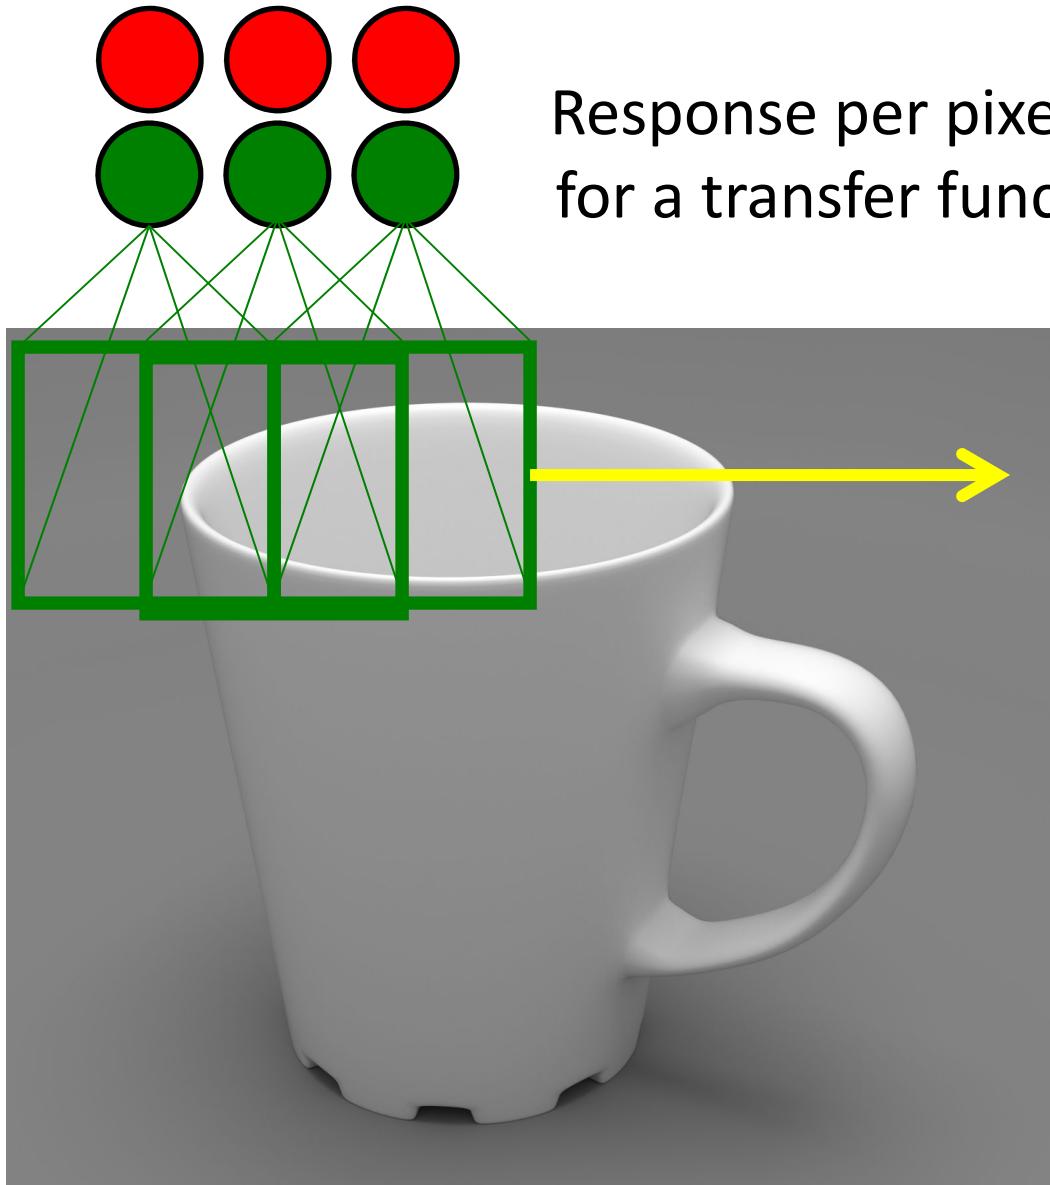
Neurons as convolution filters



Now think of neurons as convolutional **filters** acted on small adjacent (possibly overlapping) windows

Window size is called “**receptive field**” size and spacing is called “**step**” or “**stride**”

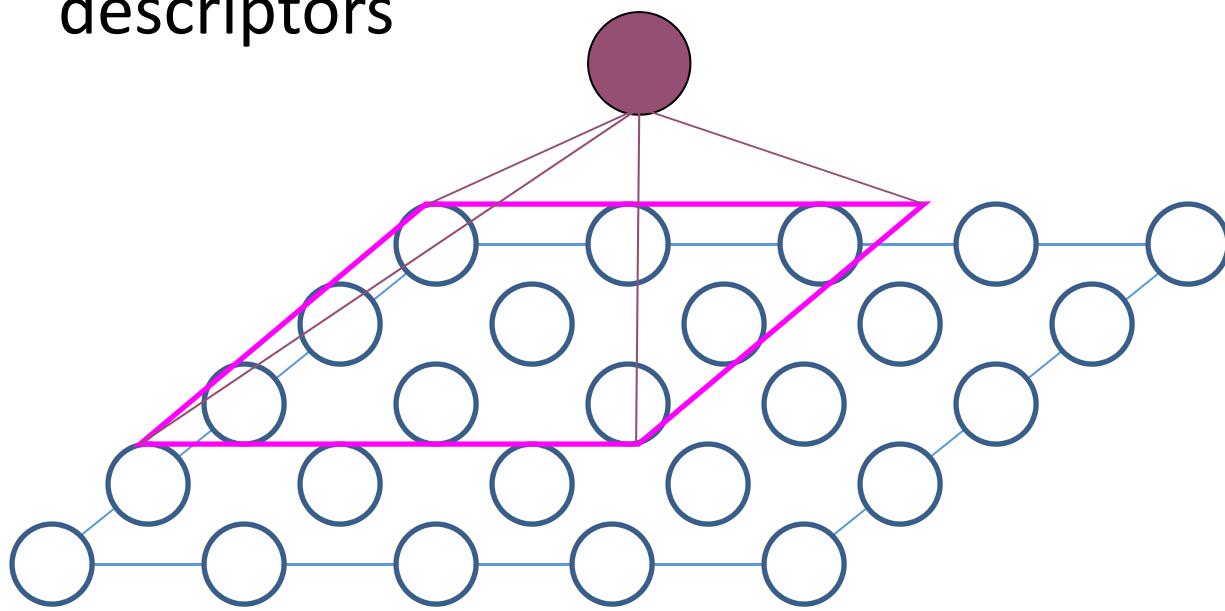
Can have many filters!



Response per pixel p , per filter f
for a transfer function g : $\mathbf{h}_{p,f} = g(\mathbf{w}_f \cdot \mathbf{x}_p)$

Pooling

Apart from hidden layers dedicated to convolution, we can have layers dedicated to extract **locally invariant** descriptors



Max pooling:

$$h_{p',f} = \max_p(\mathbf{x}_p)$$

Mean pooling:

$$h_{p',f} = \text{avg}_p(\mathbf{x}_p)$$

Fixed filter (e.g., Gaussian):

$$h_{p',f} = w_{gaussian} \cdot \mathbf{x}_p$$

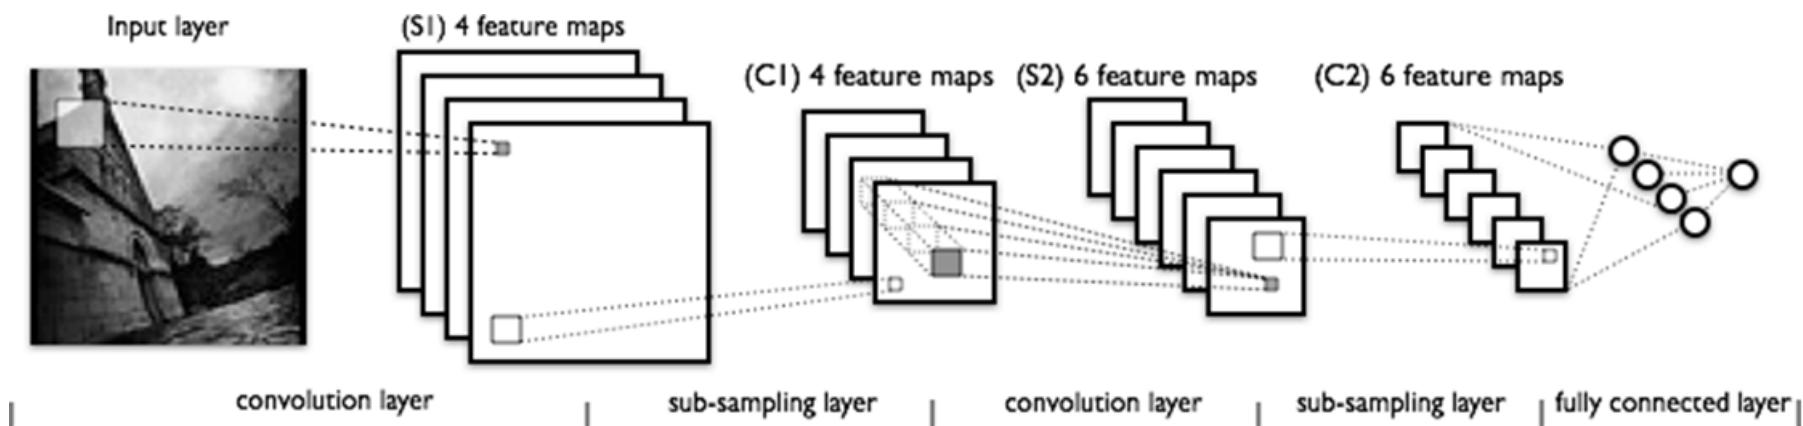
Progressively reduce the resolution of the image, so that the next convolutional filters are applied on larger scales

[Scherer et al., ICANN 2010]
[Boureau et al., ICML 2010]

A mini convolutional neural network

Interchange convolutional and pooling (subsampling) layers.

In the end, **unwrap all feature maps into a single feature vector** and pass it through the classical (**fully connected**) neural network.



Source: <http://deeplearning.net/tutorial/lenet.html>

AlexNet

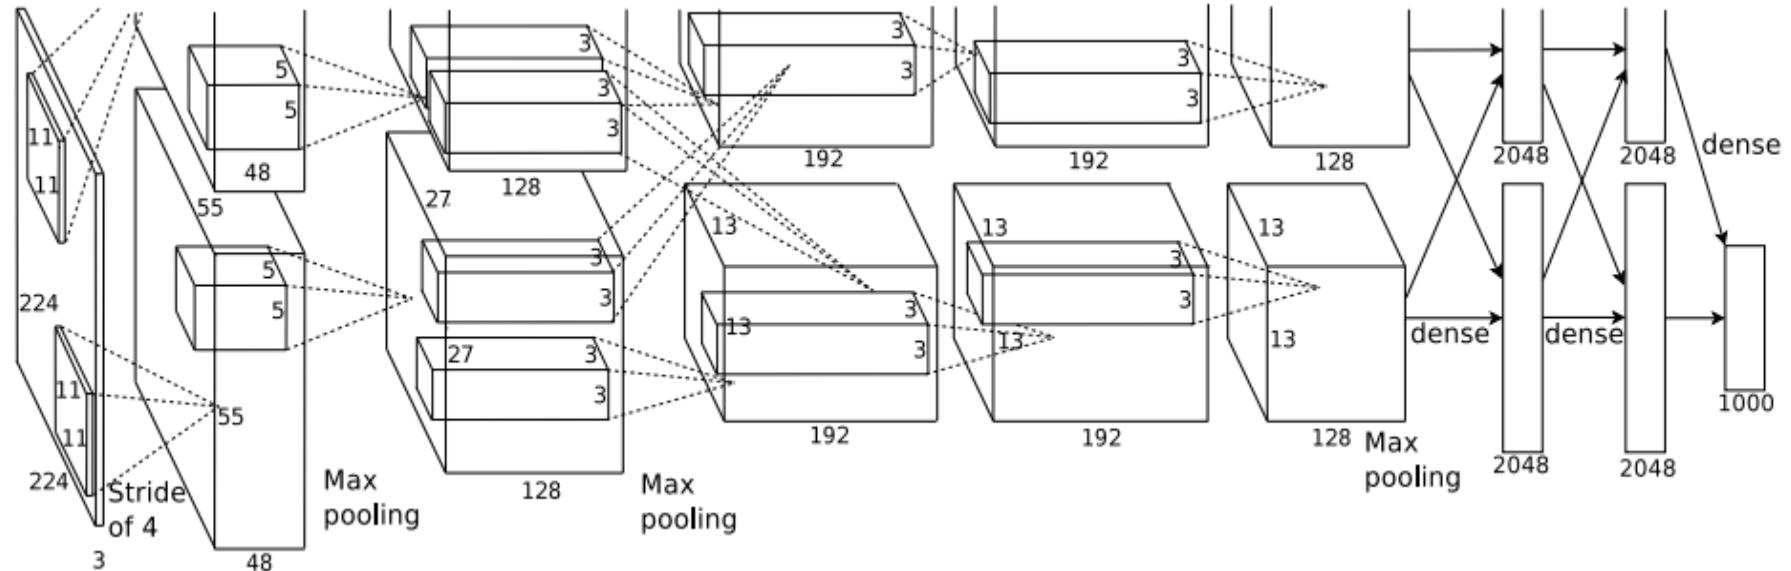
Proposed architecture from [Krizhevsky et al., NIPS 2012](#):

Convolutional layers with Rectified linear units

Max-pooling layers

Stochastic gradient descent on GPU with momentum, L2 regularization, dropout

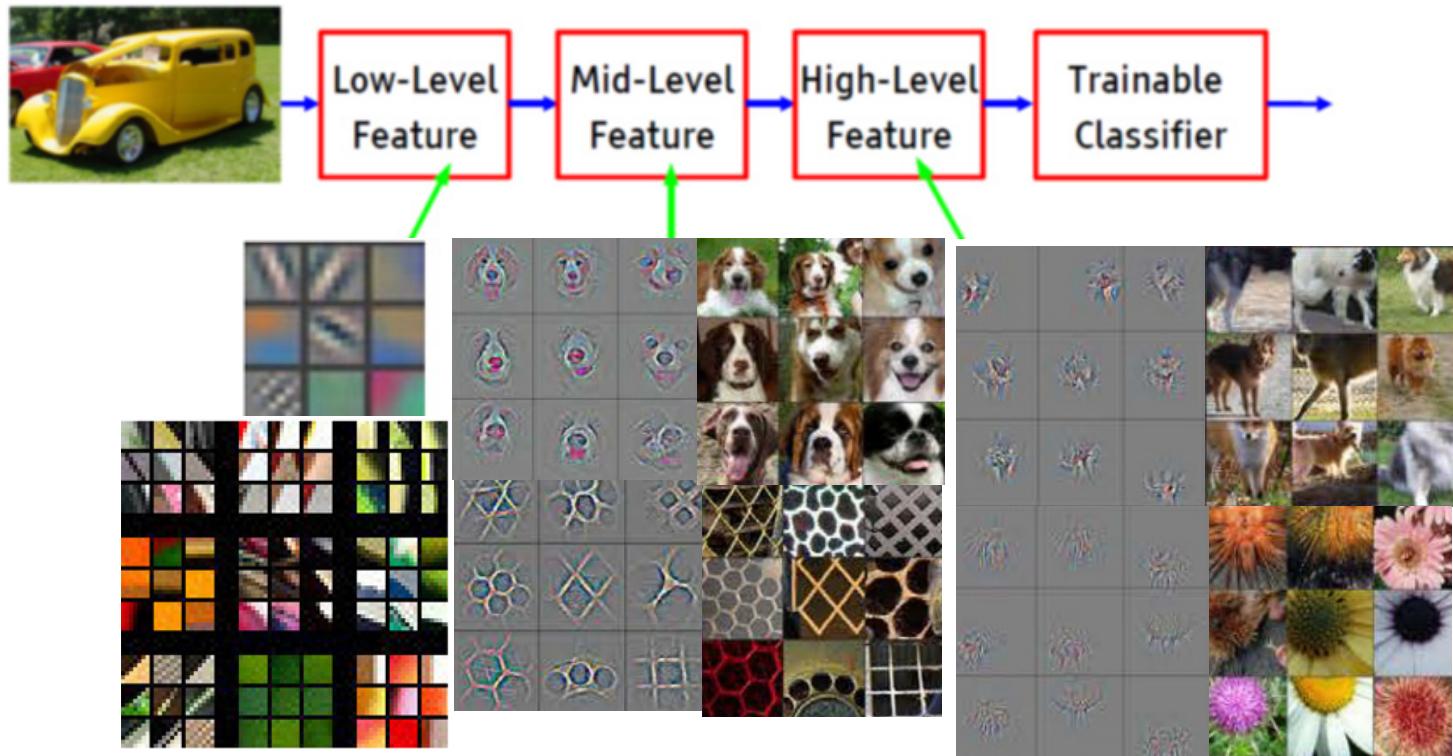
Applied to image classification (ImageNet competition – top runner & game changer)



Krizhevsky et al., NIPS 2012

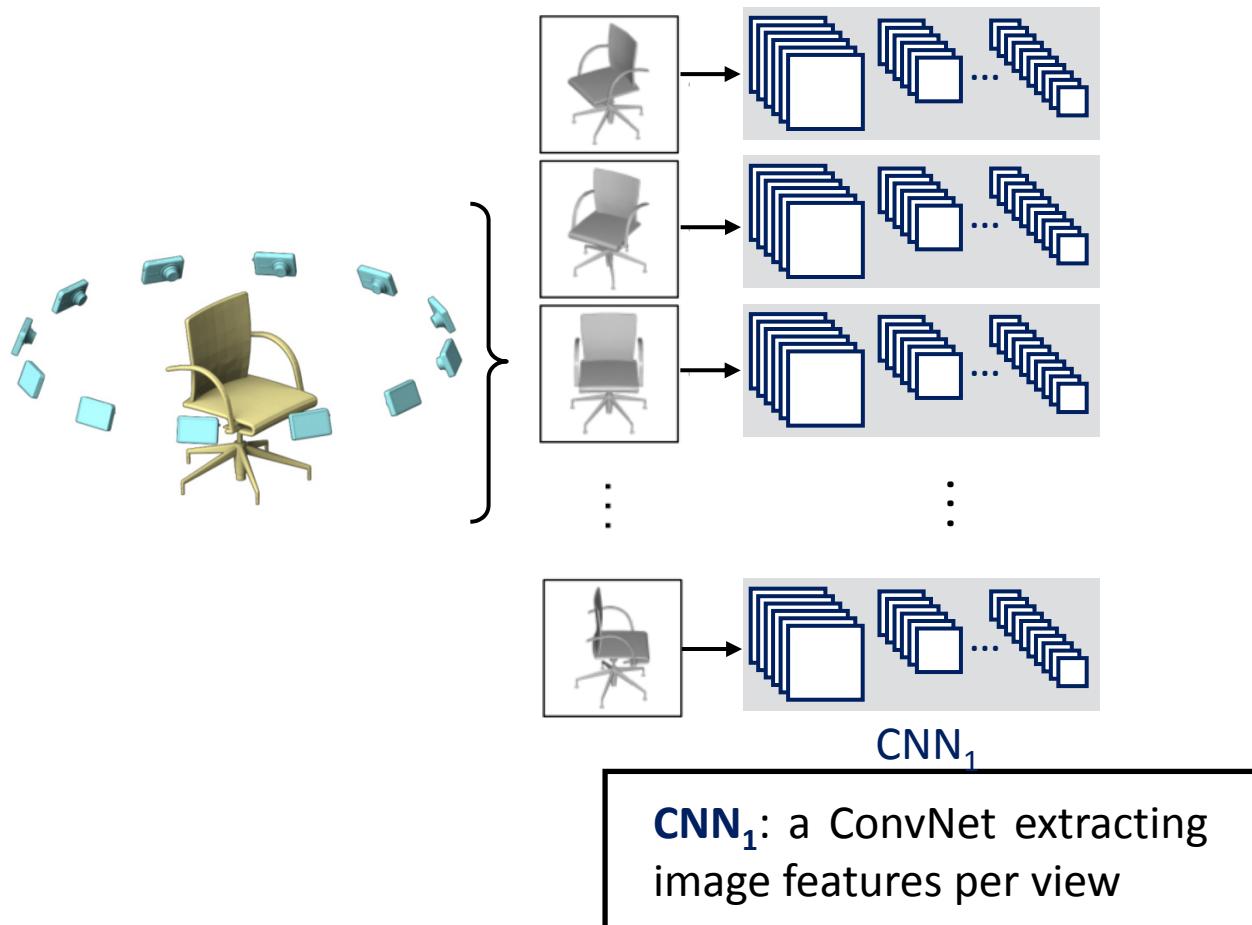
Learned representations

Think of convolution filters as optimized **feature templates** capturing various hierarchical patterns (edges, local structures, sub-parts, parts...)



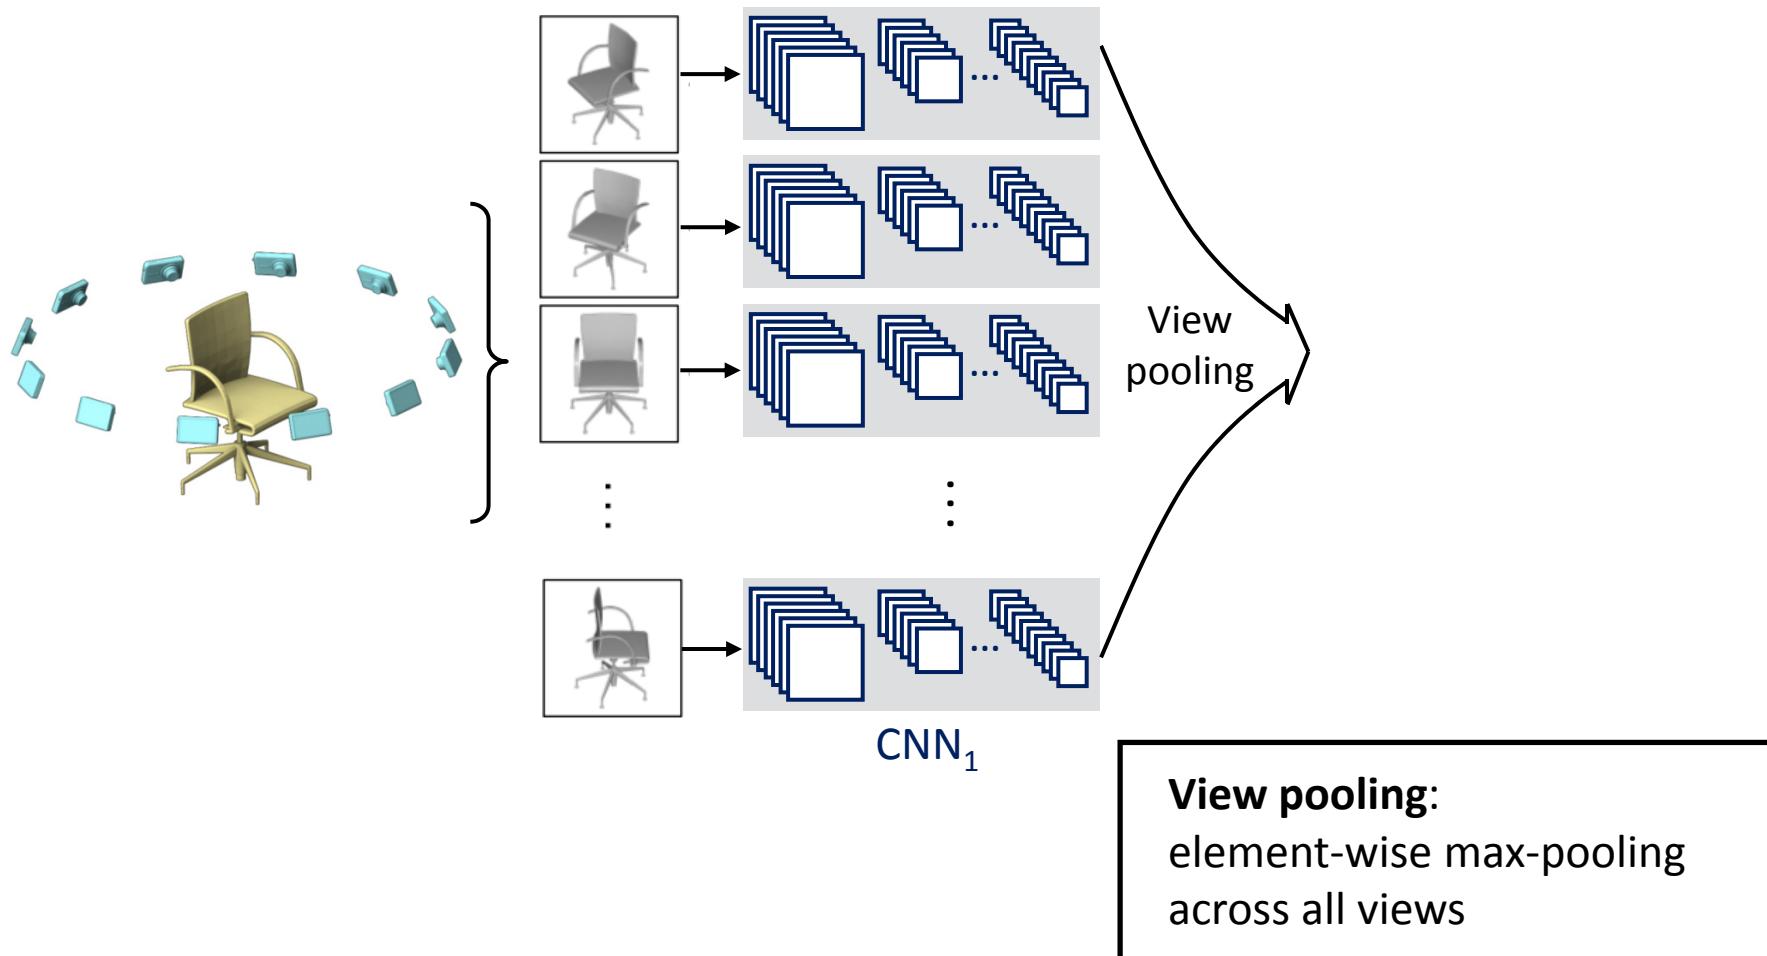
see Matthew D. Zeiler and Rob Fergus, Visualizing and Understanding Convolutional Networks, 2014

Multi-view CNNs for shape analysis



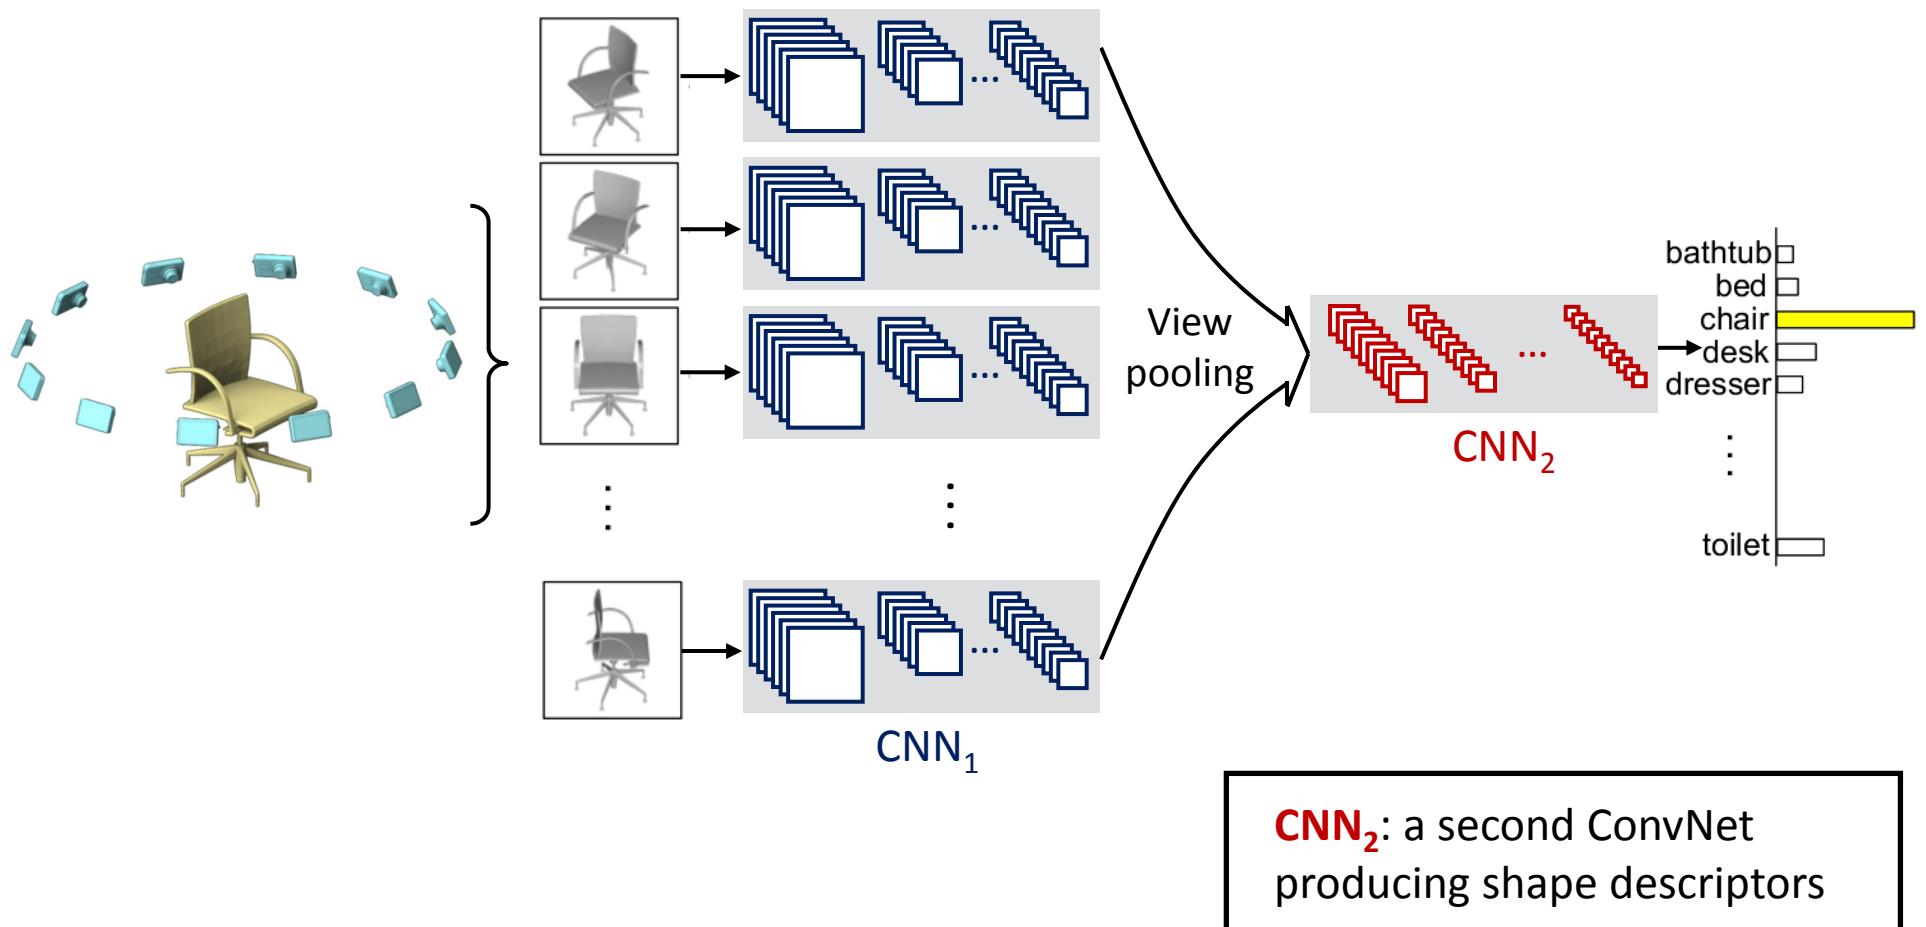
*Image from Hang Su, Subhransu Maji, Evangelos Kalogerakis, Erik Learned-Miller
Multi-view Convolutional Neural Networks for 3D Shape Recognition, ICCV 2015*

Multi-view CNNs for shape analysis



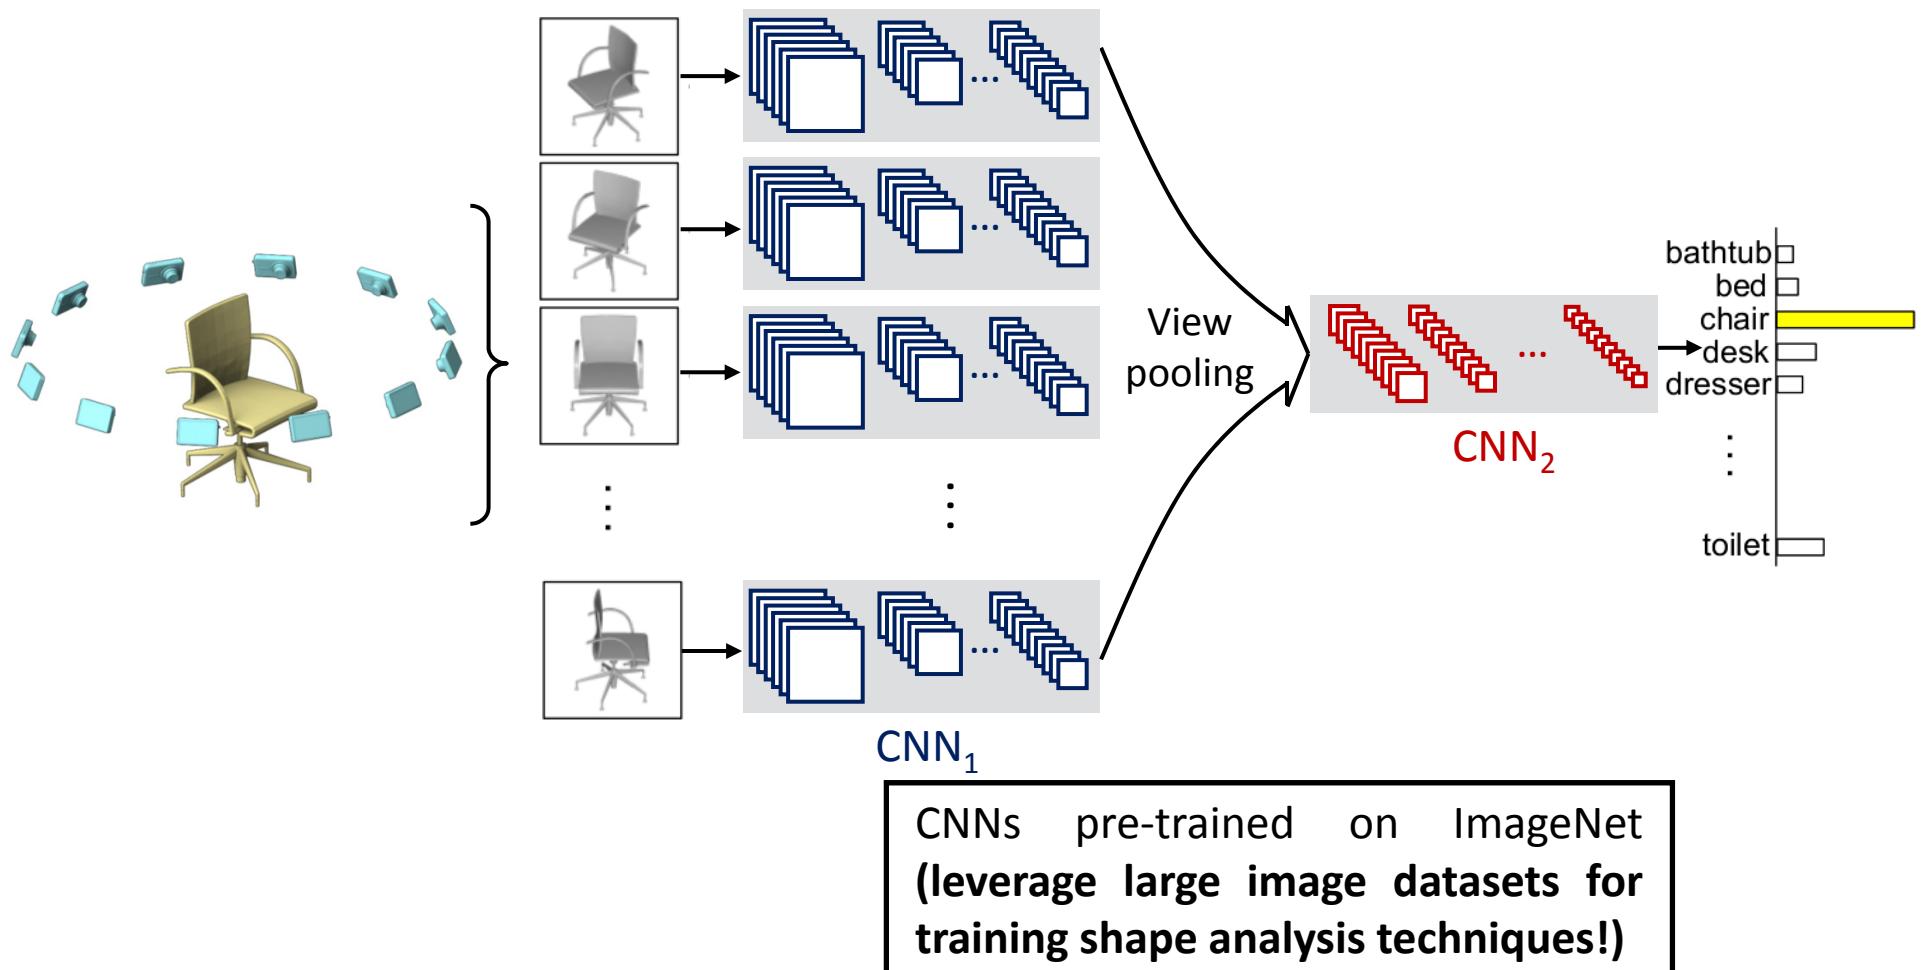
*Image from Hang Su, Subhransu Maji, Evangelos Kalogerakis, Erik Learned-Miller
Multi-view Convolutional Neural Networks for 3D Shape Recognition, ICCV 2015*

Multi-view CNNs for shape analysis



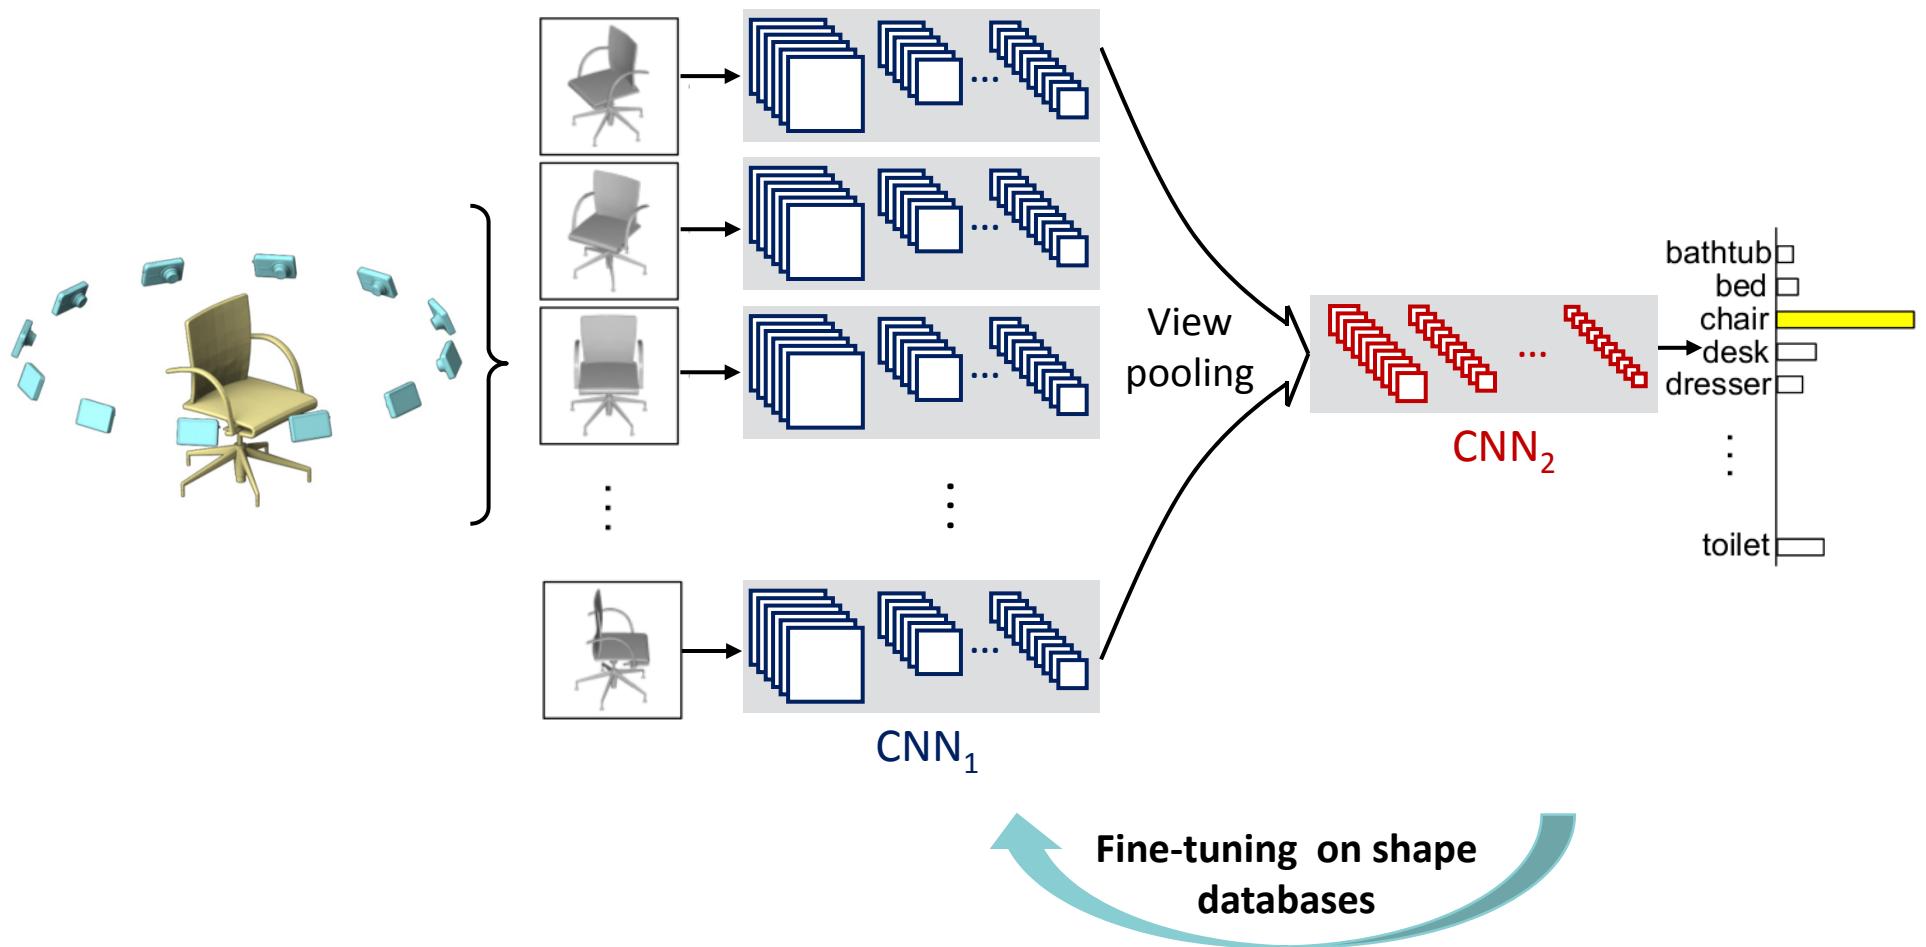
*Image from Hang Su, Subhransu Maji, Evangelos Kalogerakis, Erik Learned-Miller
Multi-view Convolutional Neural Networks for 3D Shape Recognition, ICCV 2015*

Multi-view CNNs for shape analysis



*Image from Hang Su, Subhransu Maji, Evangelos Kalogerakis, Erik Learned-Miller
Multi-view Convolutional Neural Networks for 3D Shape Recognition, ICCV 2015*

Multi-view CNNs for shape analysis

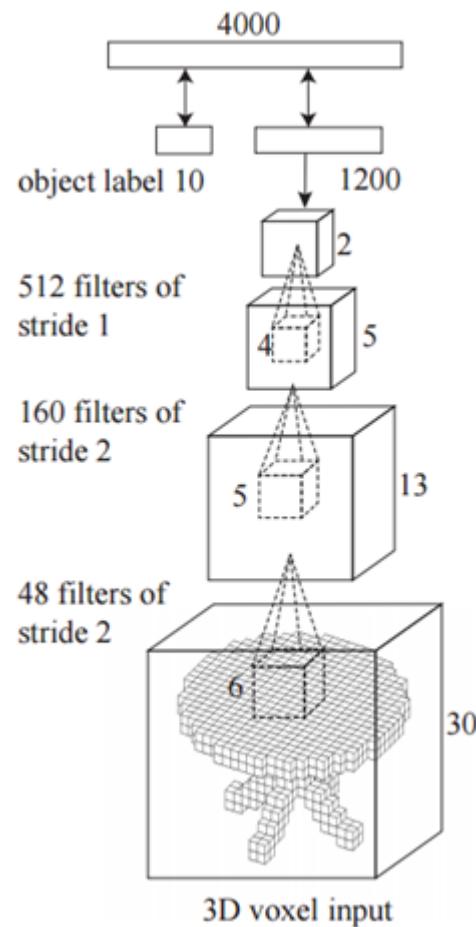


*Image from Hang Su, Subhransu Maji, Evangelos Kalogerakis, Erik Learned-Miller
Multi-view Convolutional Neural Networks for 3D Shape Recognition, ICCV 2015*

Volumetric CNNs

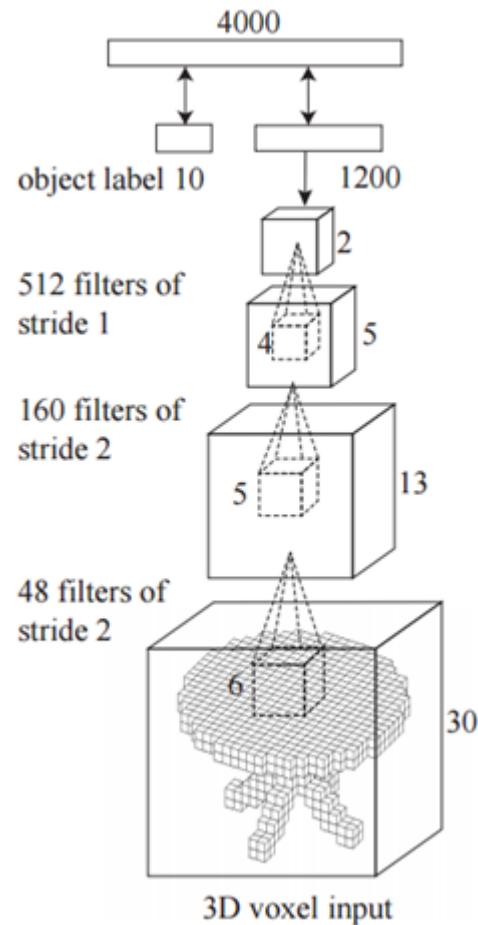
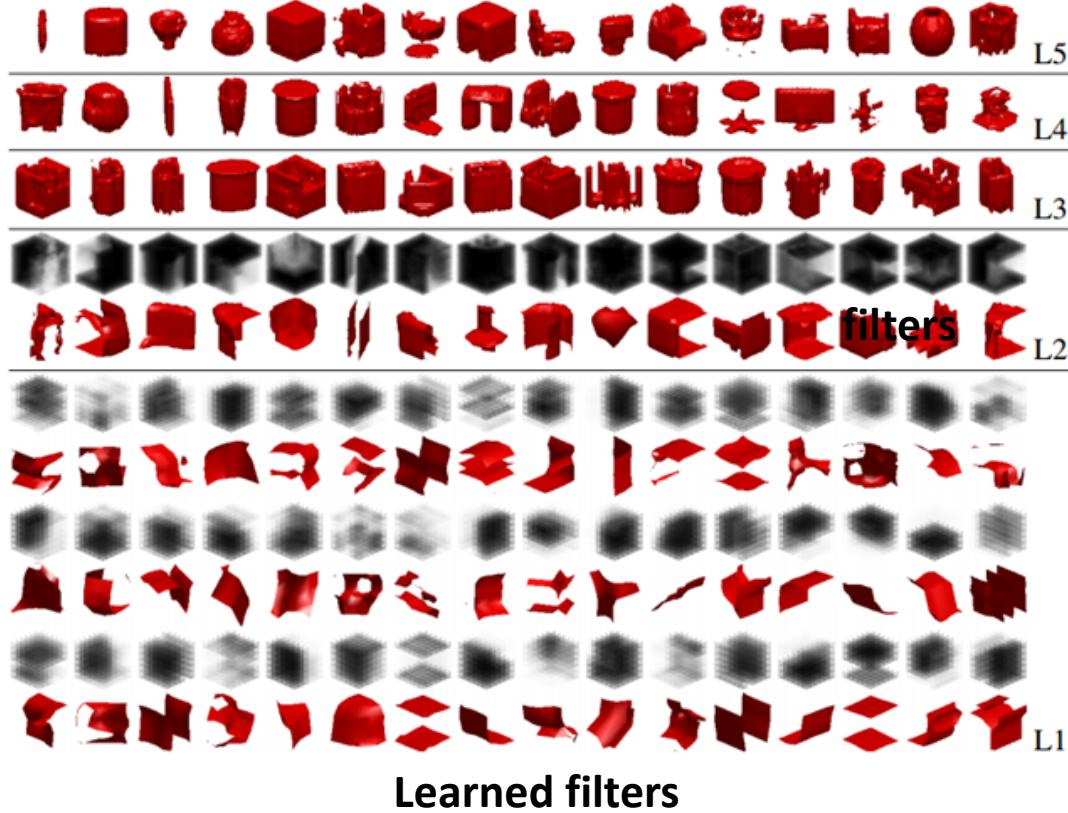
Key idea: represent a shape as a volumetric image with binary voxels.

Learn filters operating on these volumetric data.



*Image from Z. Wu, S. Song, A. Khosla, F. Yu, L. Zhang, X. Tang and J. Xiao
3D ShapeNets: A Deep Representation for Volumetric Shapes, 2015*

Volumetric CNNs



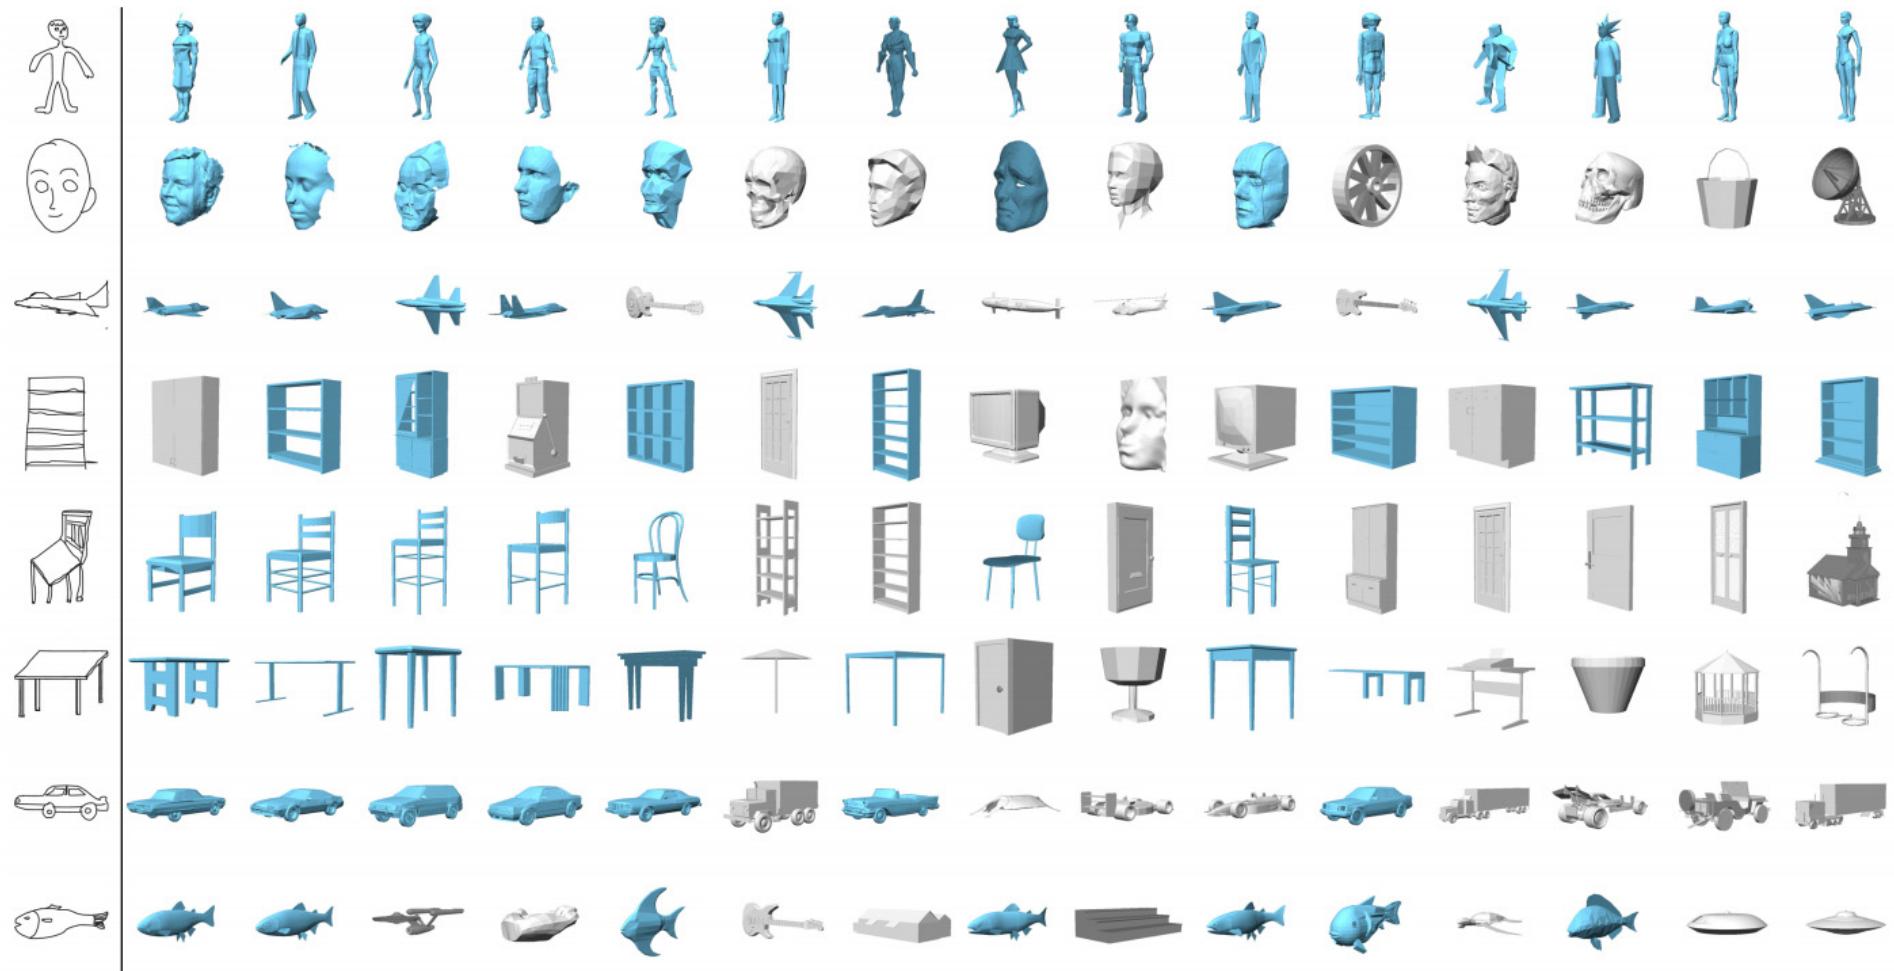
*Image from Z. Wu, S. Song, A. Khosla, F. Yu, L. Zhang, X. Tang and J. Xiao
3D ShapeNets: A Deep Representation for Volumetric Shapes, 2015*

Sketch-based 3D Shape Retrieval using CNNs



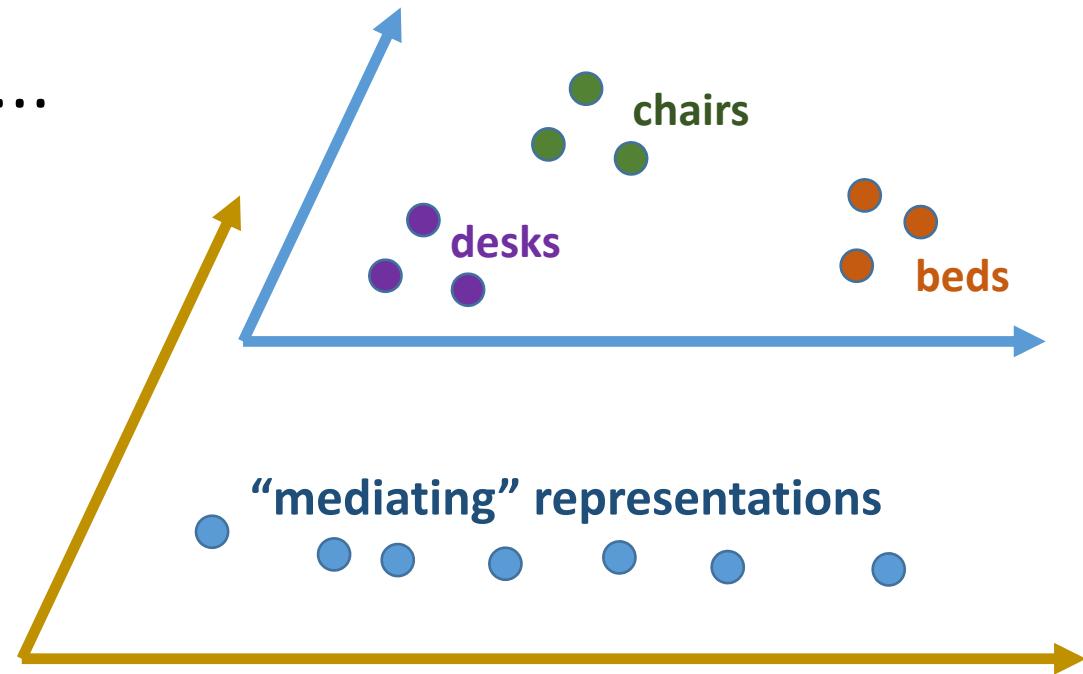
*Image from Fang Wang, Le Kang, Yi Li,
Sketch-based 3D Shape Retrieval using Convolutional Neural Networks, 2015*

Sketch-based 3D Shape Retrieval using CNNs



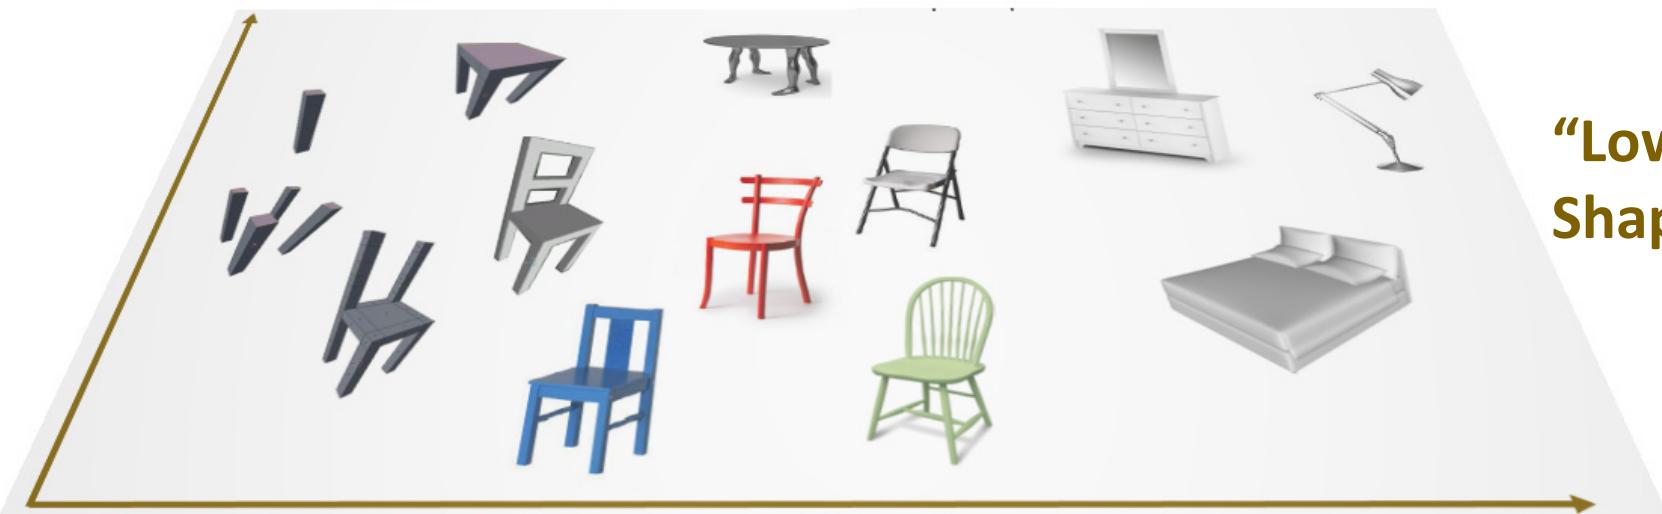
*Image from Fang Wang, Le Kang, Yi Li,
Sketch-based 3D Shape Retrieval using Convolutional Neural Networks, 2015*

So far...

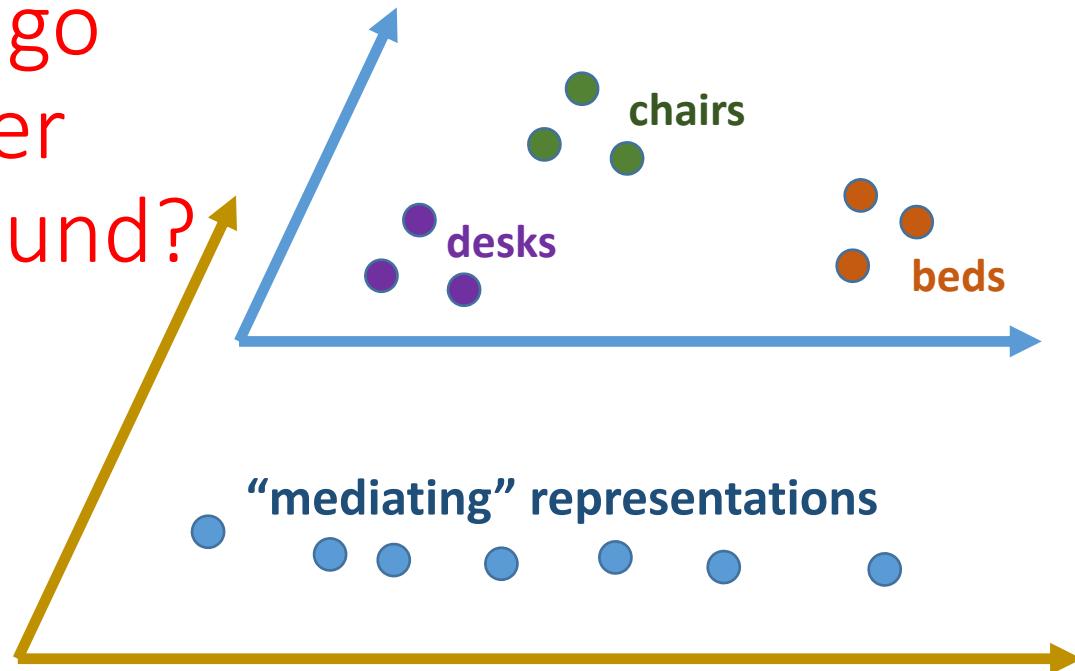


"High-level" Space

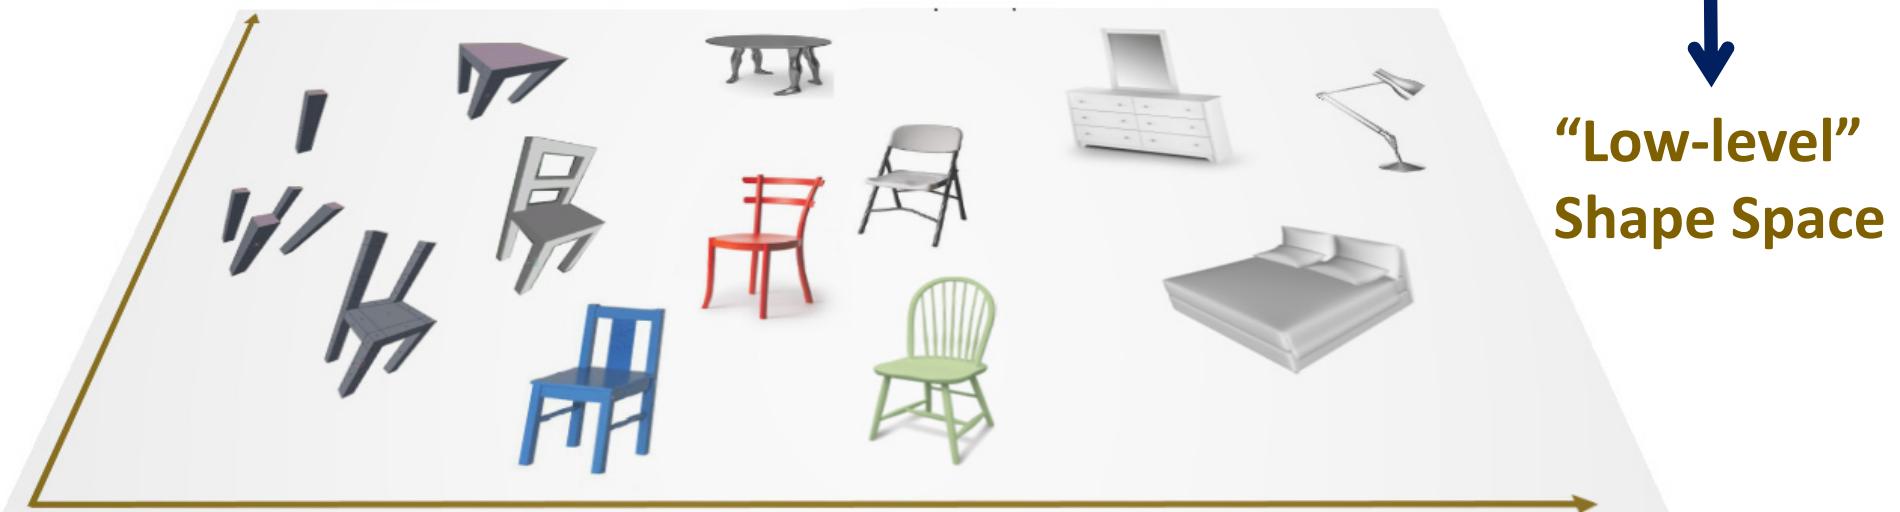
"Low-level" Shape Space



Can we go
the other
way around?

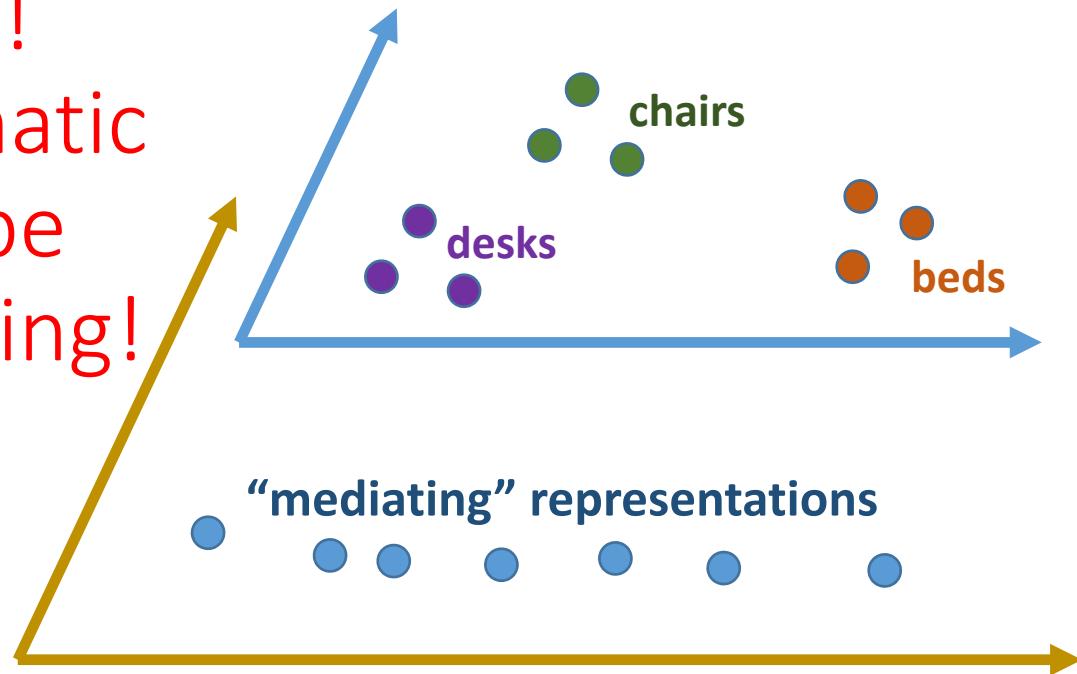


"High-level"
Space

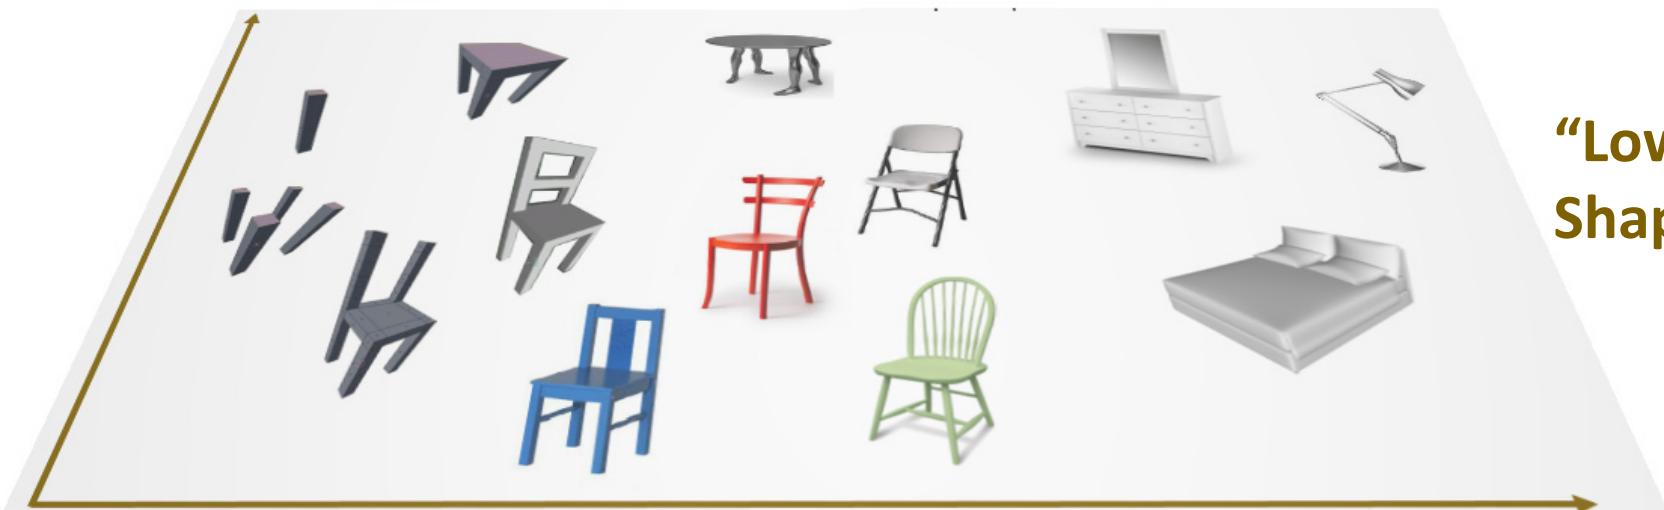


"Low-level"
Shape Space

YES!
Automatic
Shape
modeling!

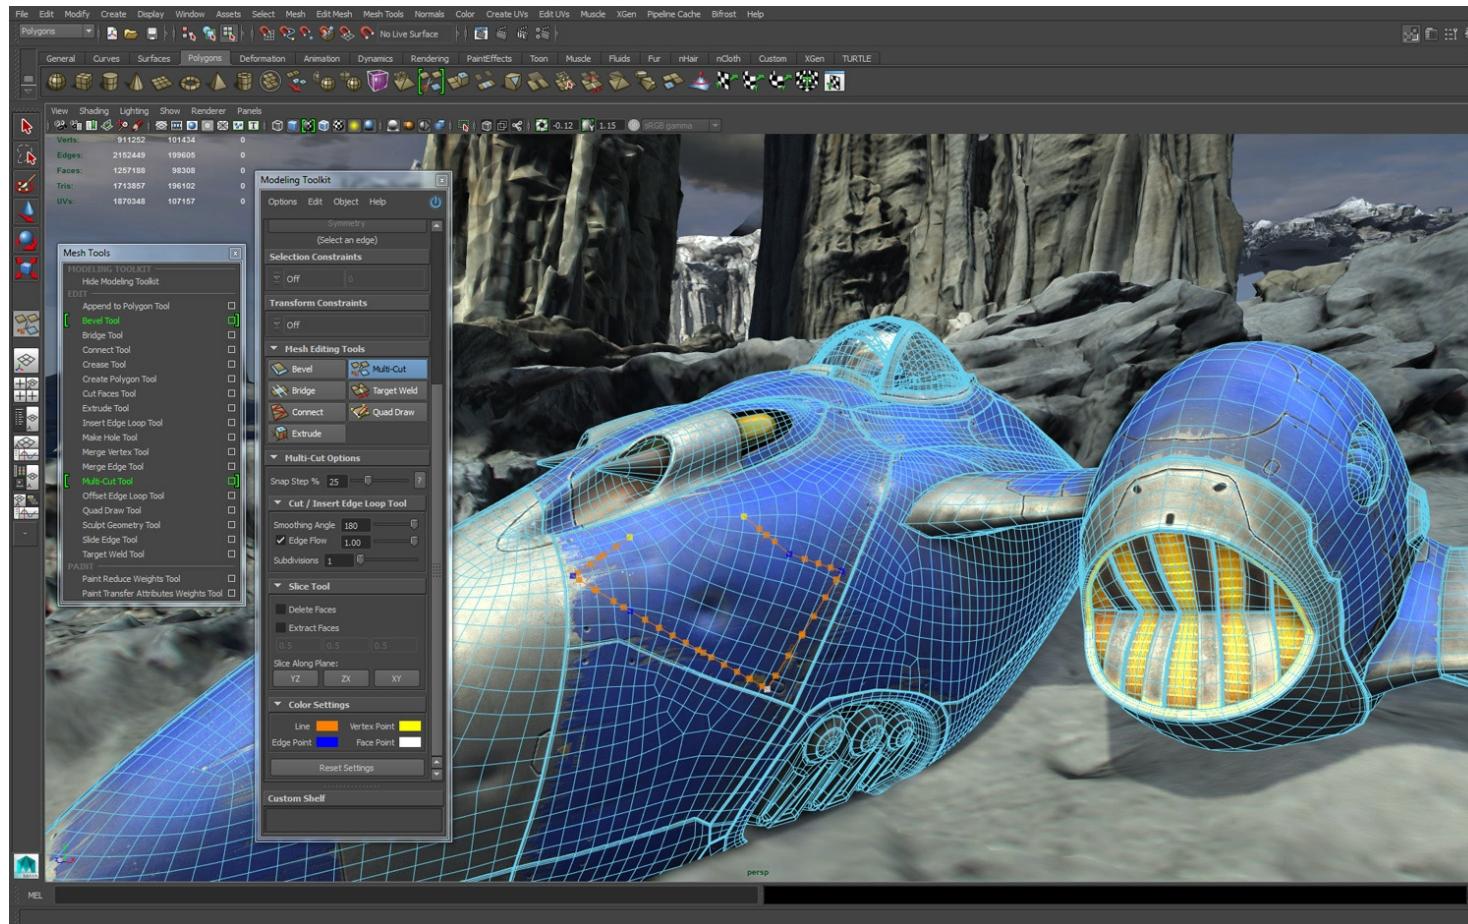


"High-level"
Space



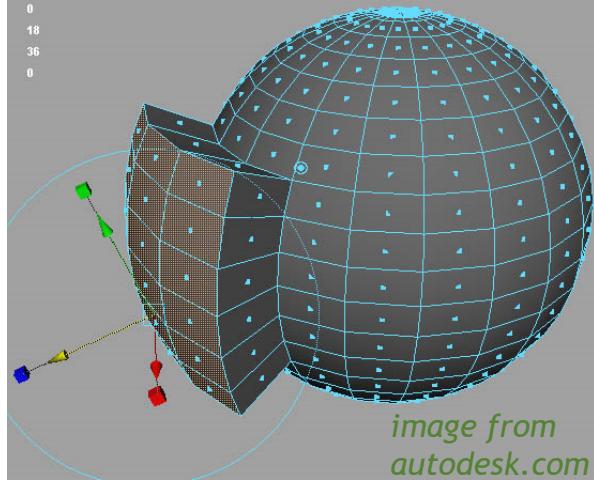
"Low-level"
Shape Space

Why automatic geometric modeling? Because it is not easy!

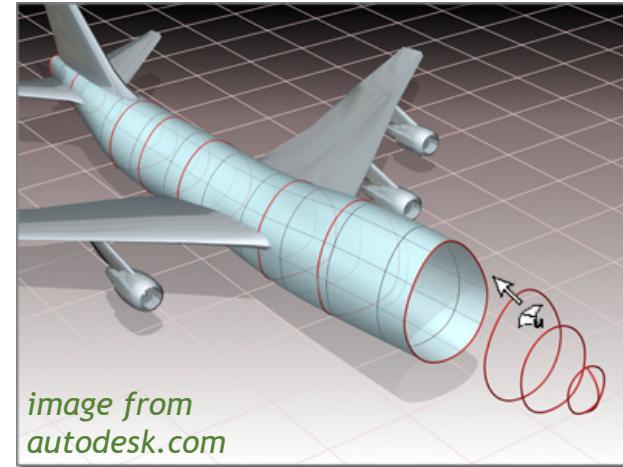


Autodesk Maya 2015

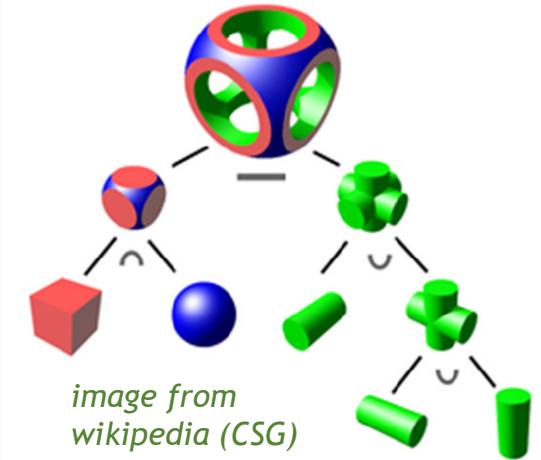
“Traditional” Geometric Modeling



Manipulating polygons



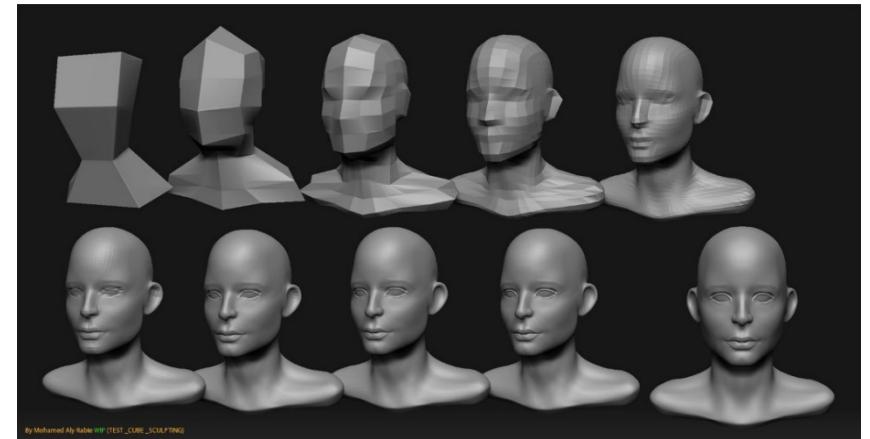
Manipulating curves



Manipulating 3D primitives



Manipulating control points, cages image from Blender



Digital Sculpting image from Mohamed Aly Rable

“Traditional” Geometric Modeling

Impressive results at the hands of **experienced users**

Operations requires **exact and accurate input**

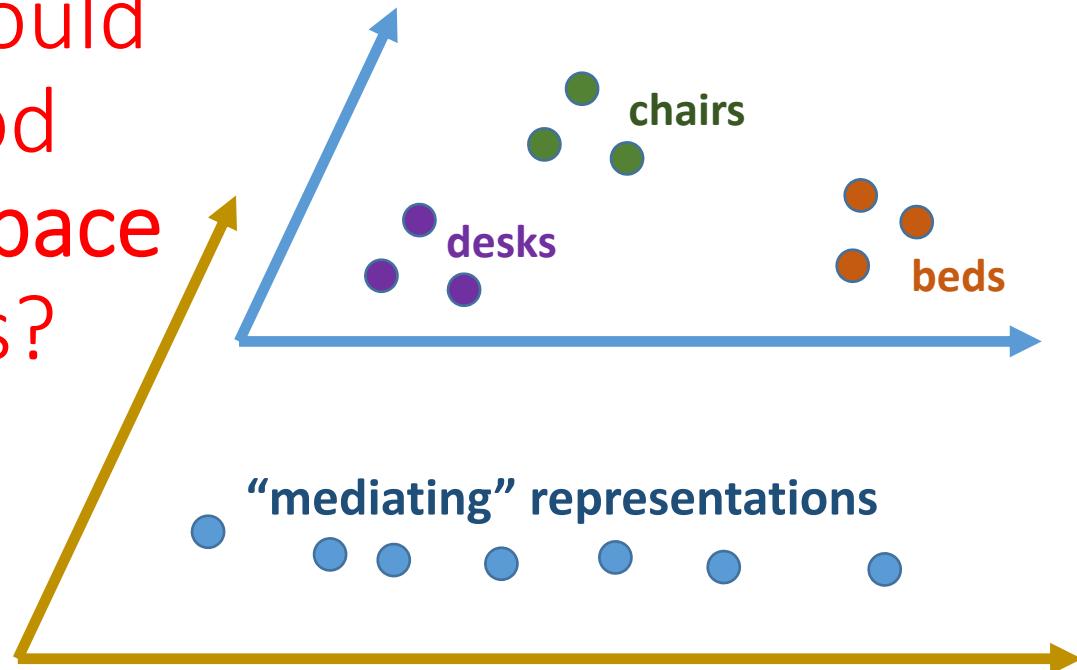
Creating compelling 3D models **takes lots of time**

Tools usually have **steep learning curves**

An alternative approach...

- Users provide **high-level, possibly approximate** input
 - Computers learn to generate **low-level, accurate** geometry
- **Machine learning!**

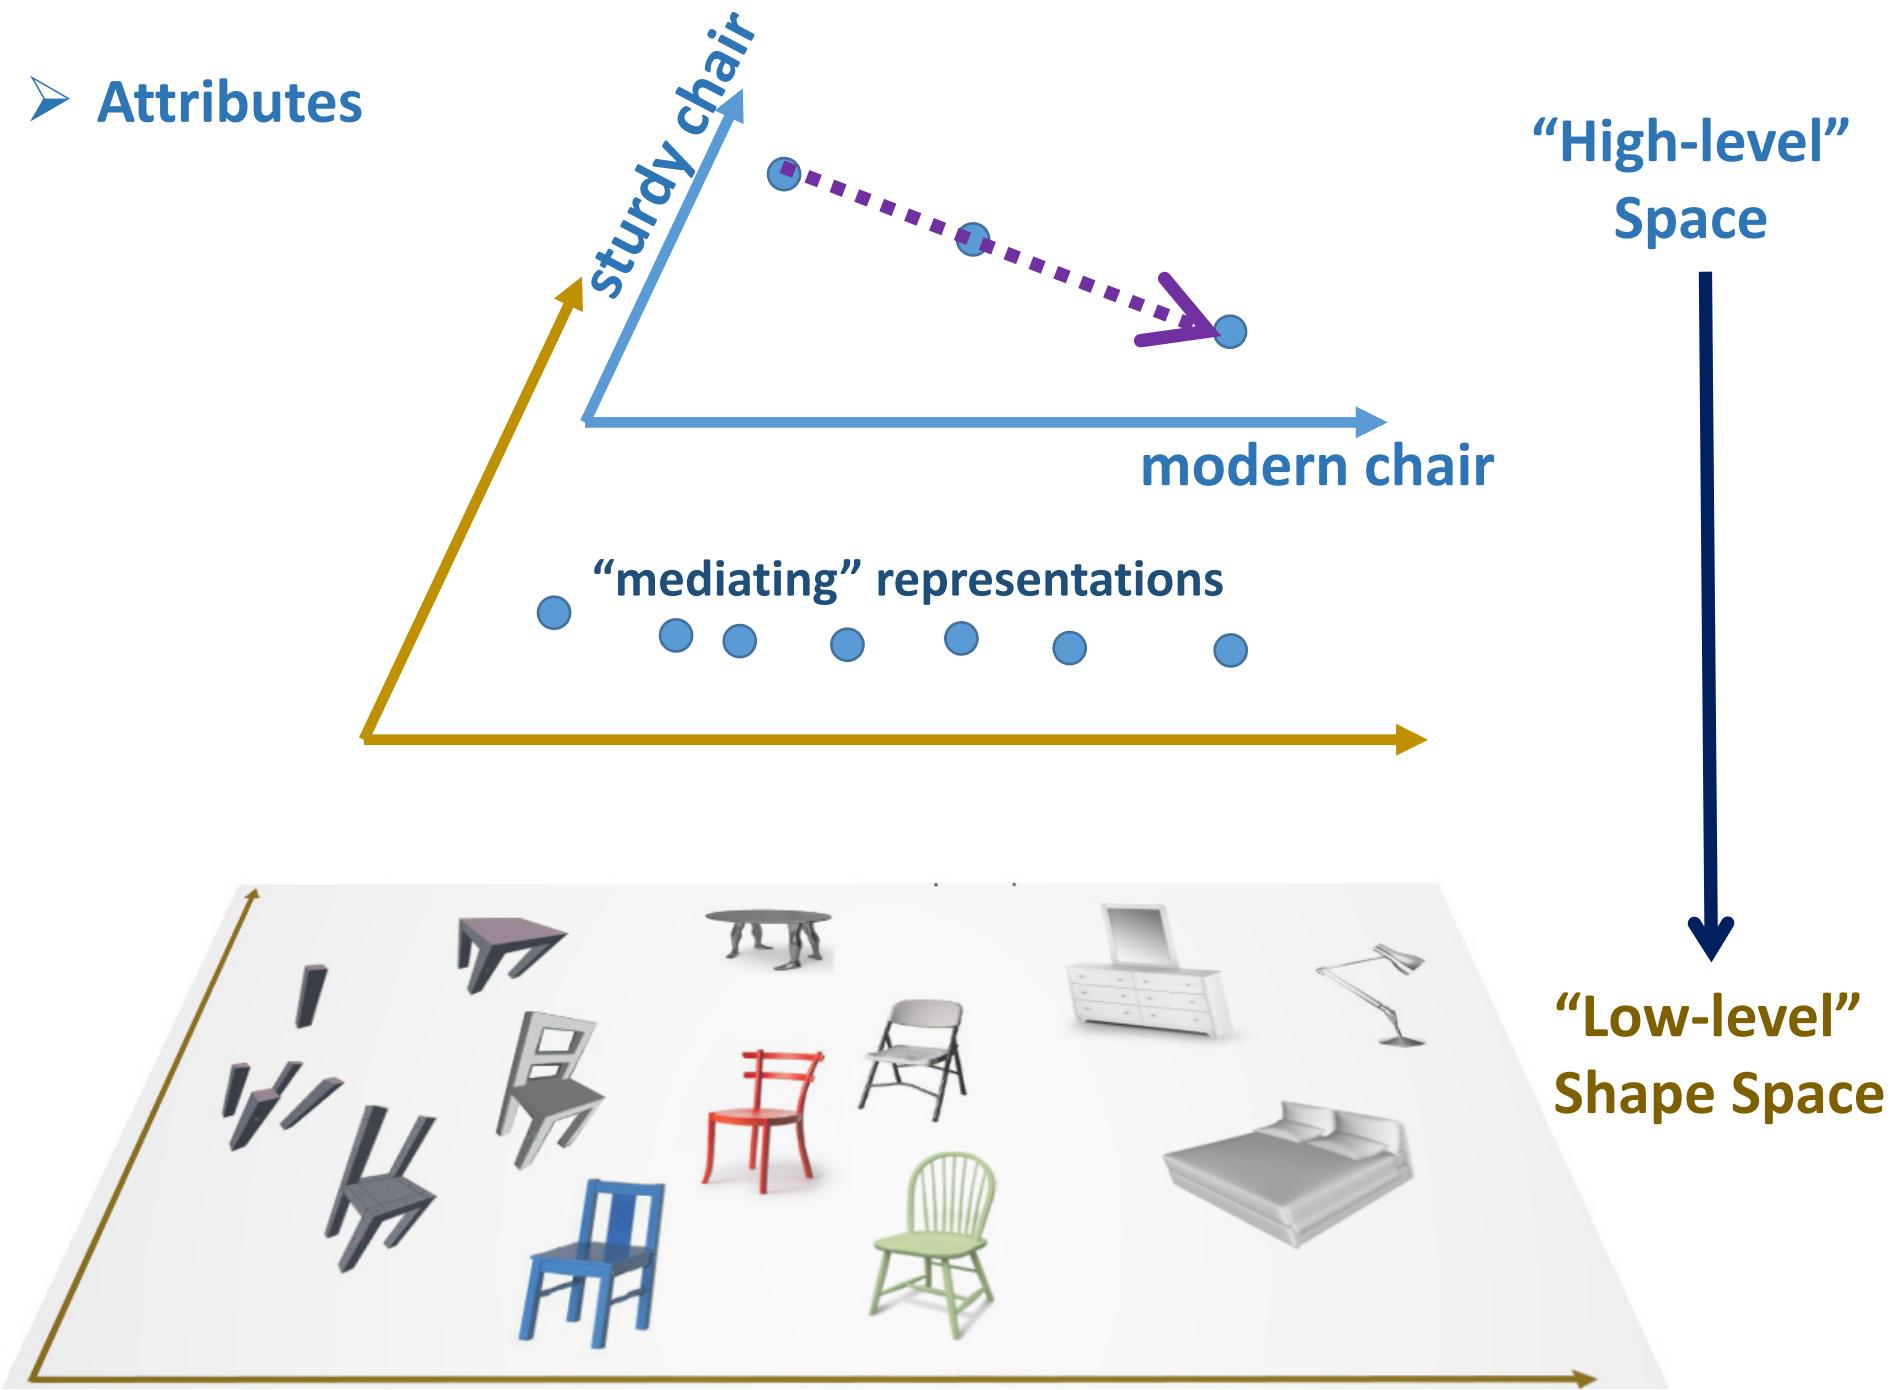
What would
be a good
design space
for users?



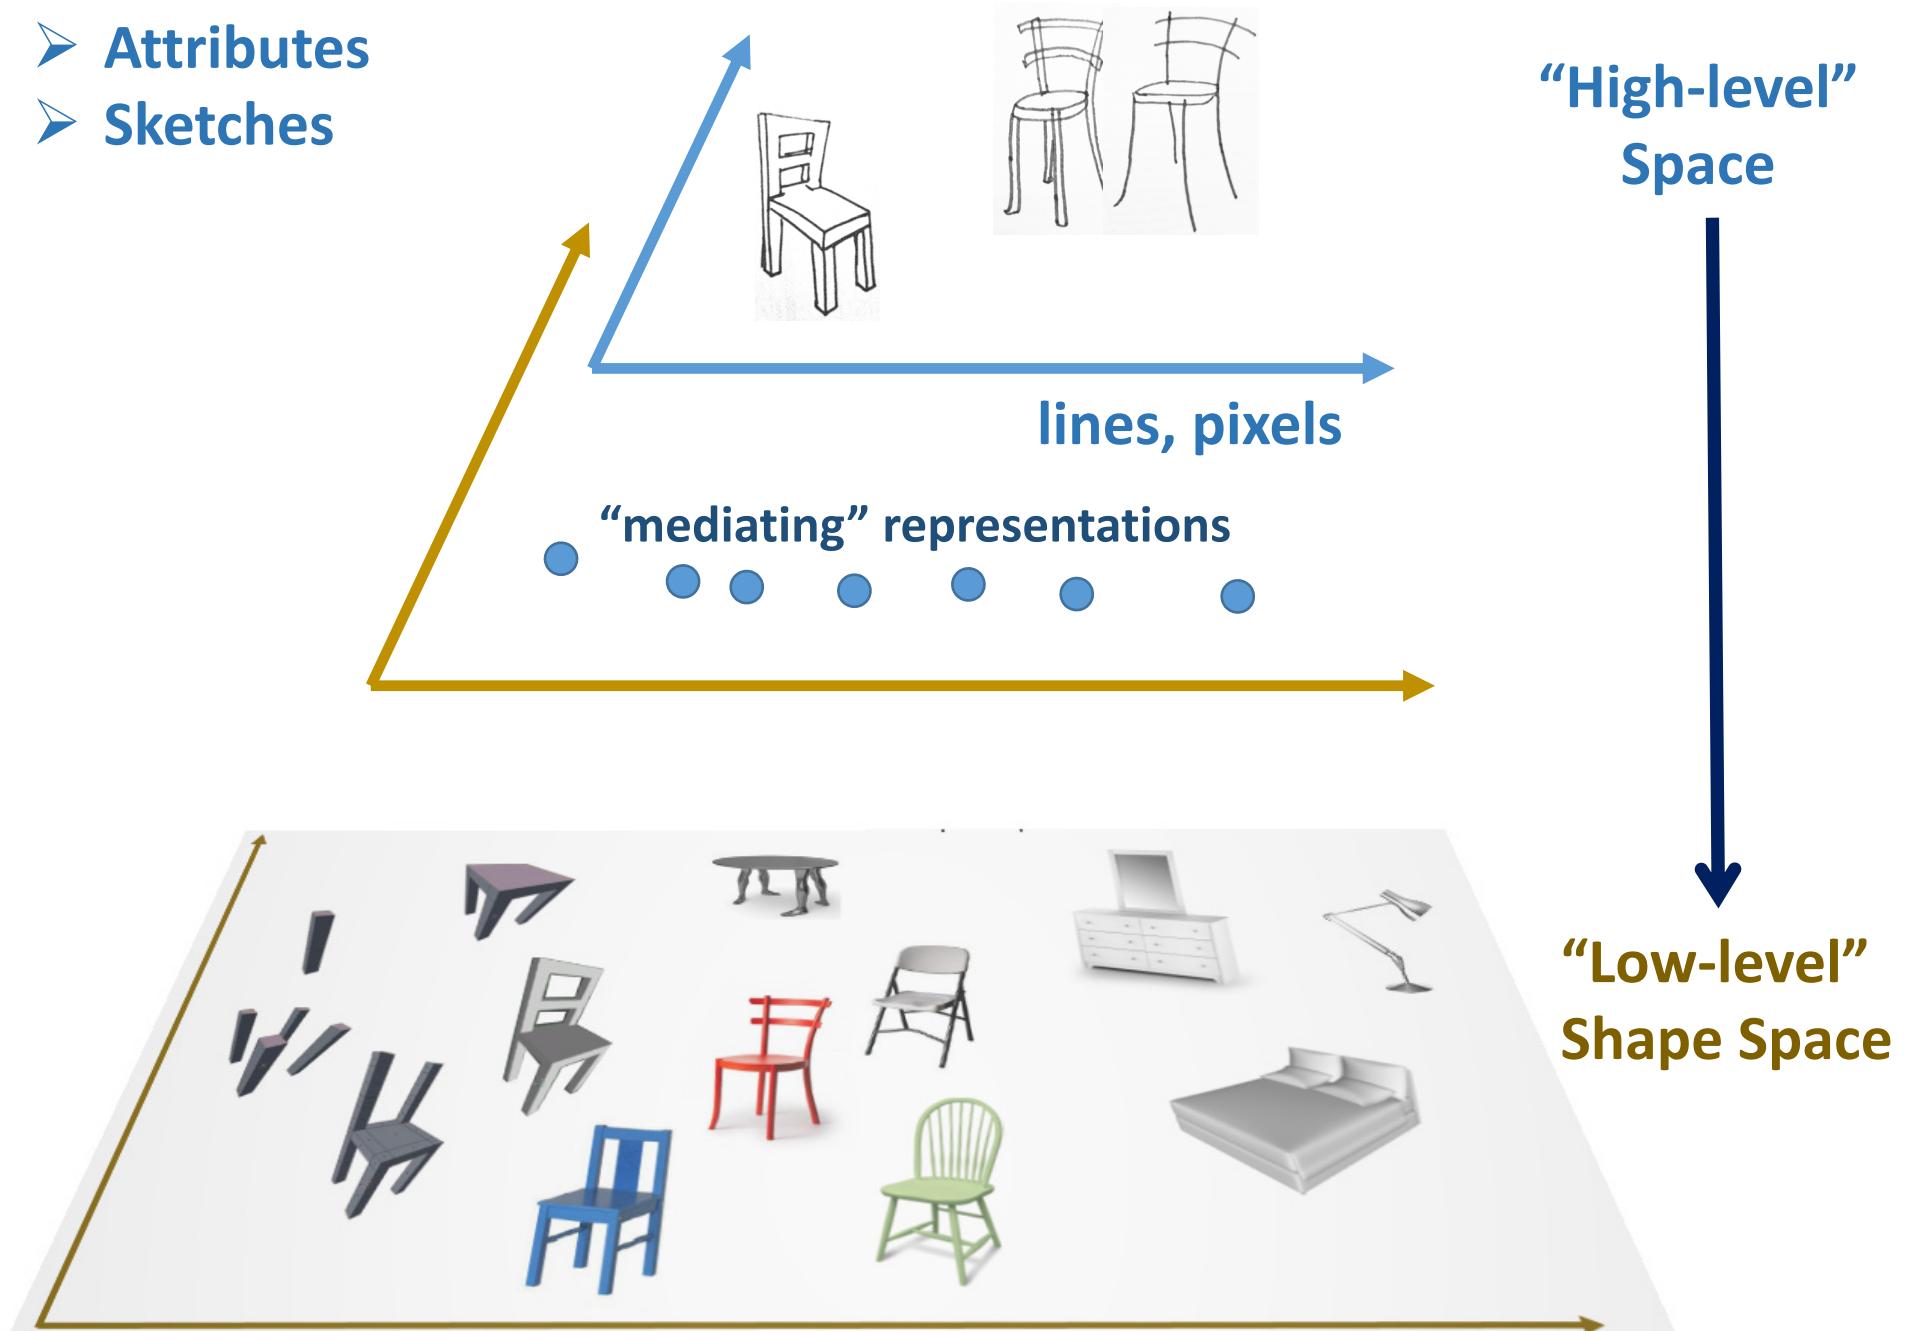
"High-level"
Space

"Low-level"
Shape Space

➤ Attributes



- Attributes
- Sketches

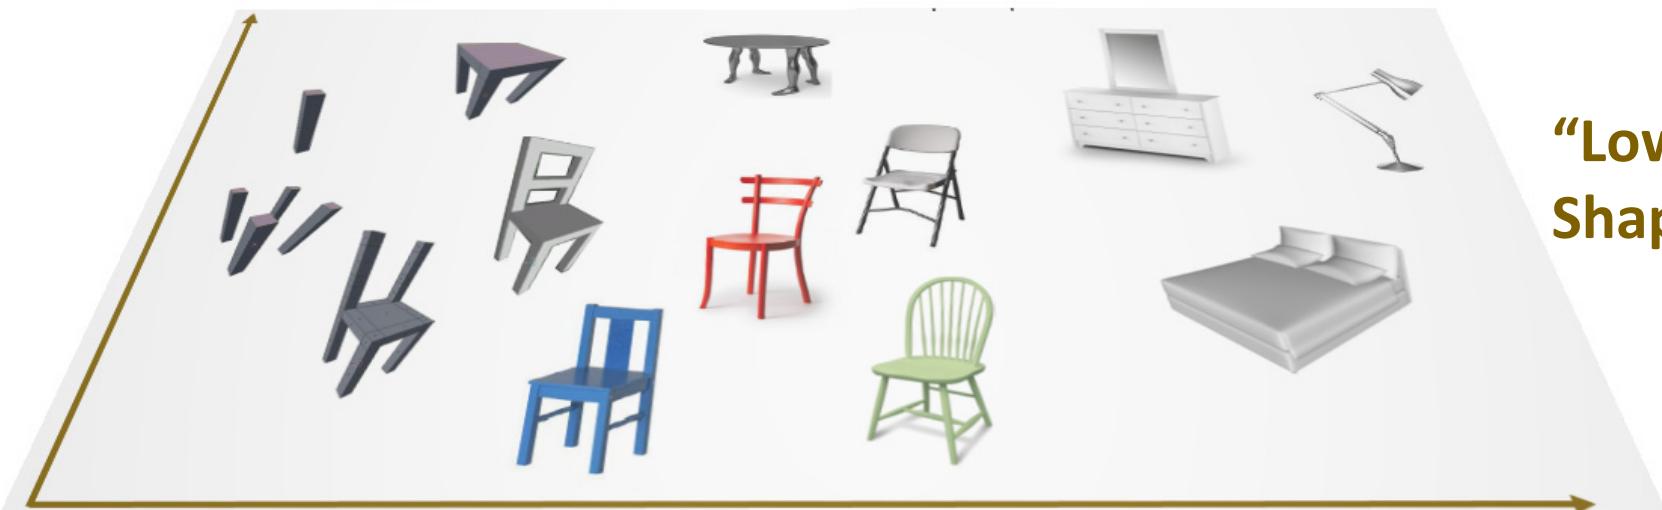


- Attributes
- Sketch
- Natural language
- Gestures
- Brain signals
- ...
- Learn it from data!

“mediating” representations

?

“High-level” Space



“Low-level” Shape Space

Machine learning for Geometric Modeling

- Learn mappings from **design** (“high-level”) to “low-level” space: $\mathbf{y} = f(\mathbf{x})$
- Learn which shapes are **probable** (“plausible”) given input: $P(y | f(\mathbf{x}))$

“Plausible” chairs



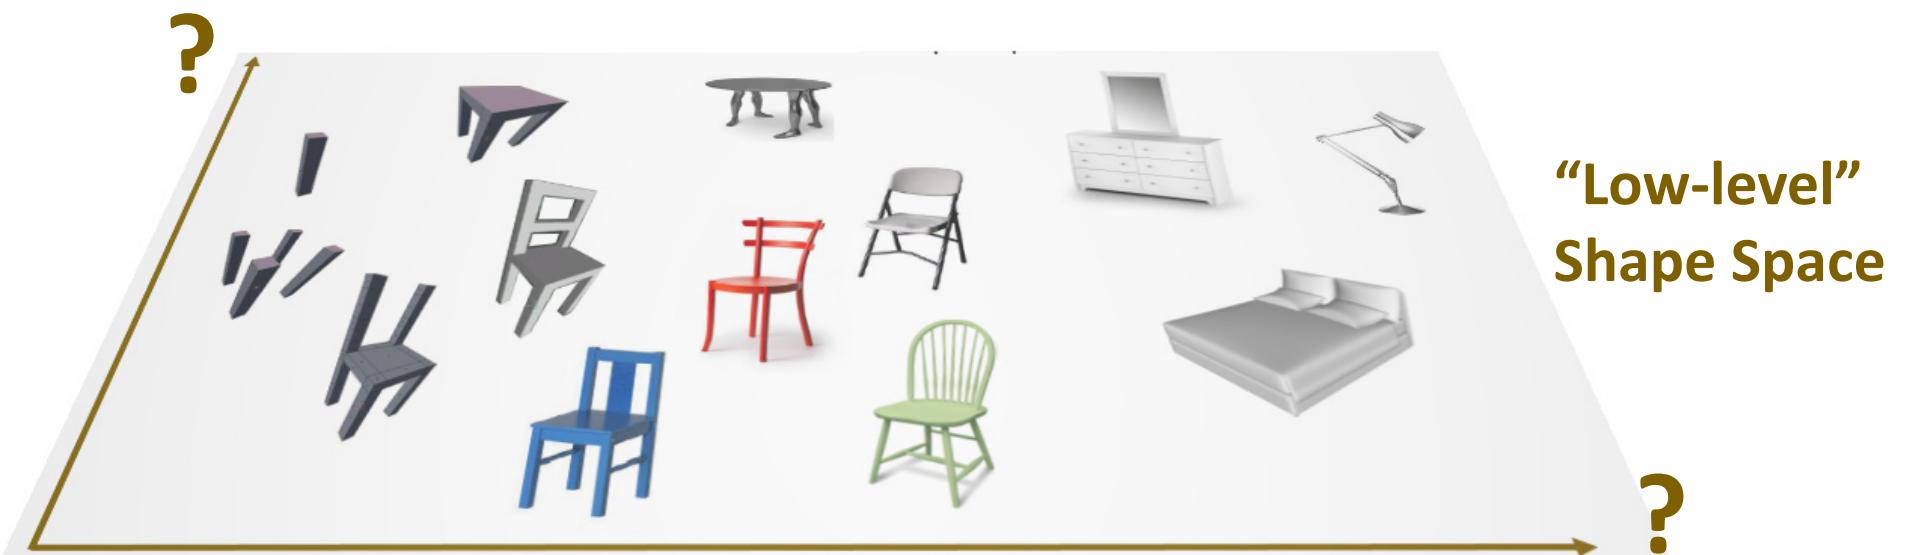
“Plausible” chairs



(not a binary or even an objective choice!)

The representation challenge

How do we represent the **shape space**?

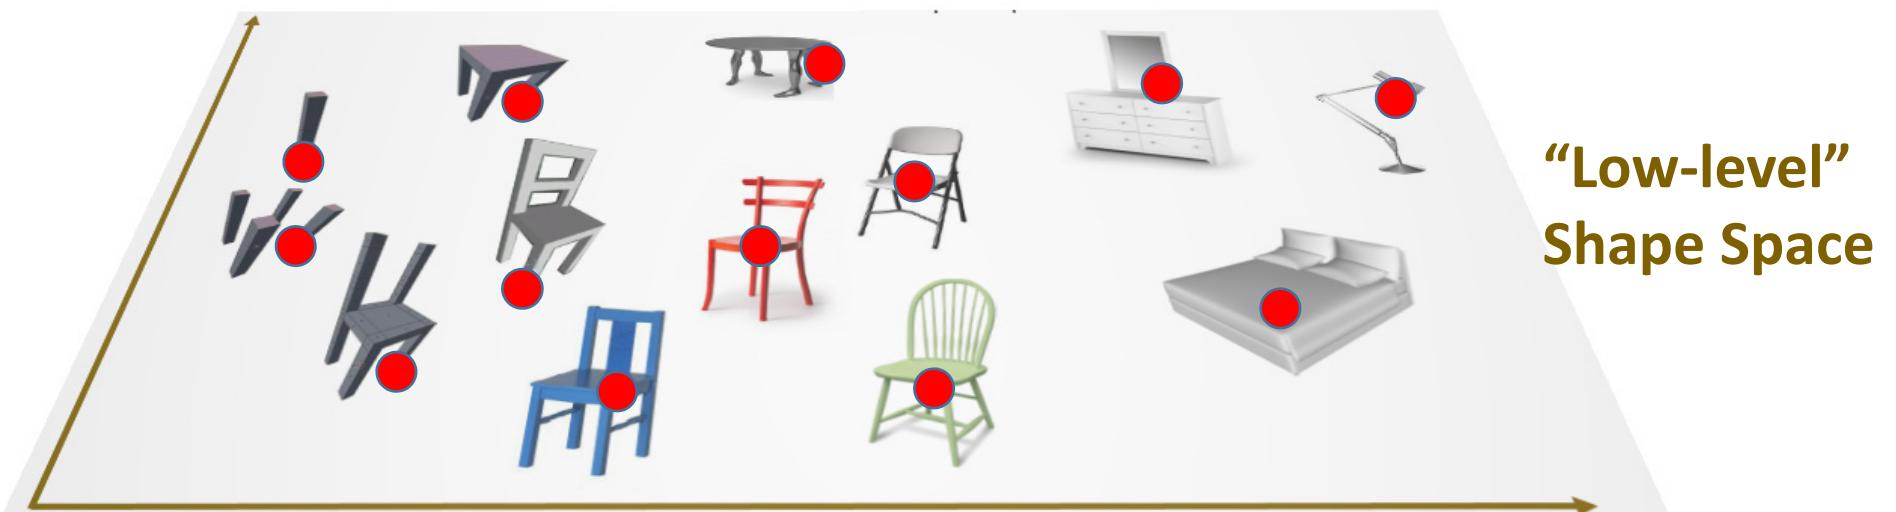


“Low-level” shape space representation

Can we use the polygon meshes as-is for our shape space?

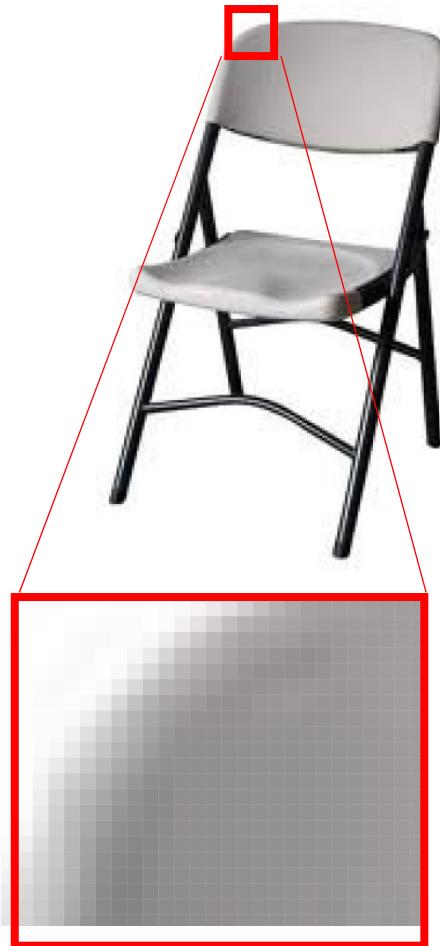
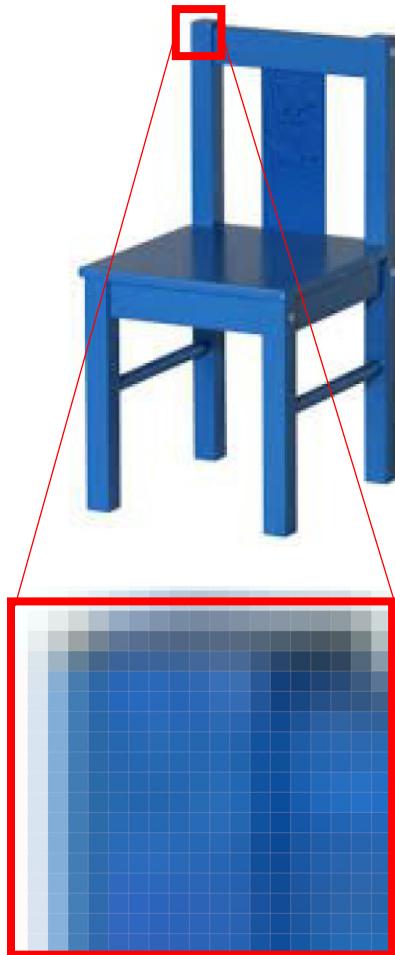
No. Take the first vertex on each mesh. Where is it?

Meshes have different number of vertices, faces etc



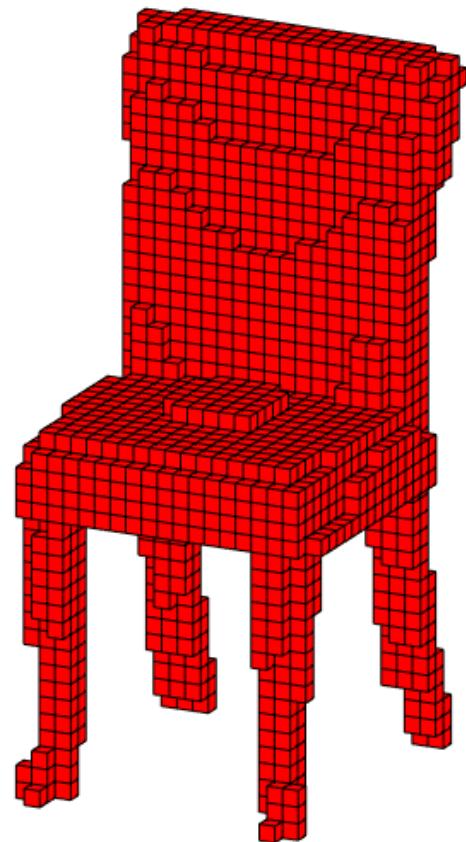
The “computer vision” approach

Learn mappings to pixels & multiple views!



The “volumetric” approach

Learn mappings to voxels!



The “correspondences” approach

Find point correspondences between 3D surface points. Can do alignment. Can we always have dense correspondences?

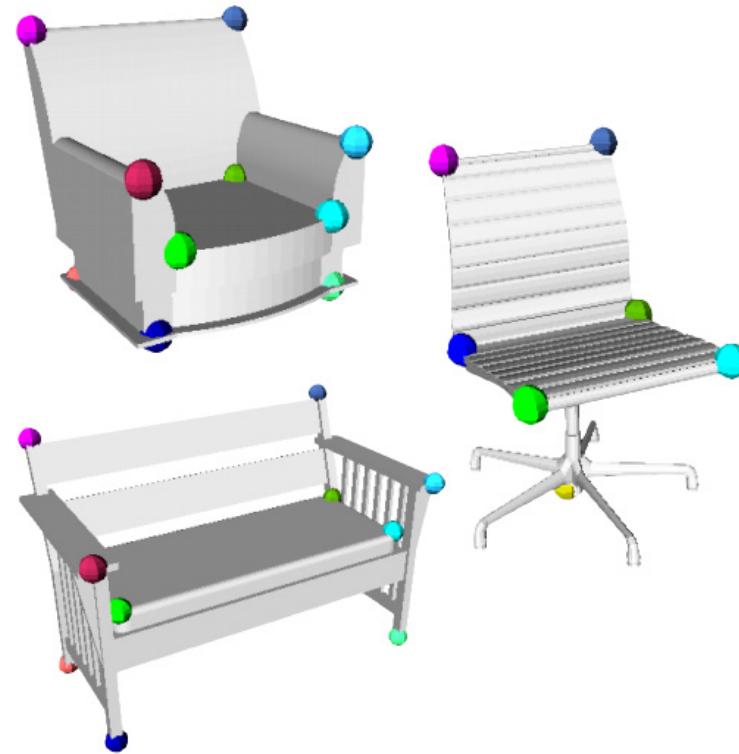


Image from Vladimir G. Kim, Wilmot Li, Niloy J. Mitra, Siddhartha Chaudhuri, Stephen DiVerdi, and Thomas Funkhouser, “Learning Part-based Templates from Large Collections of 3D Shapes”, 2013

The “abstractions” approach

Parameterize shapes with primitives (cuboids, cylinders etc)
How can we capture surface detail?

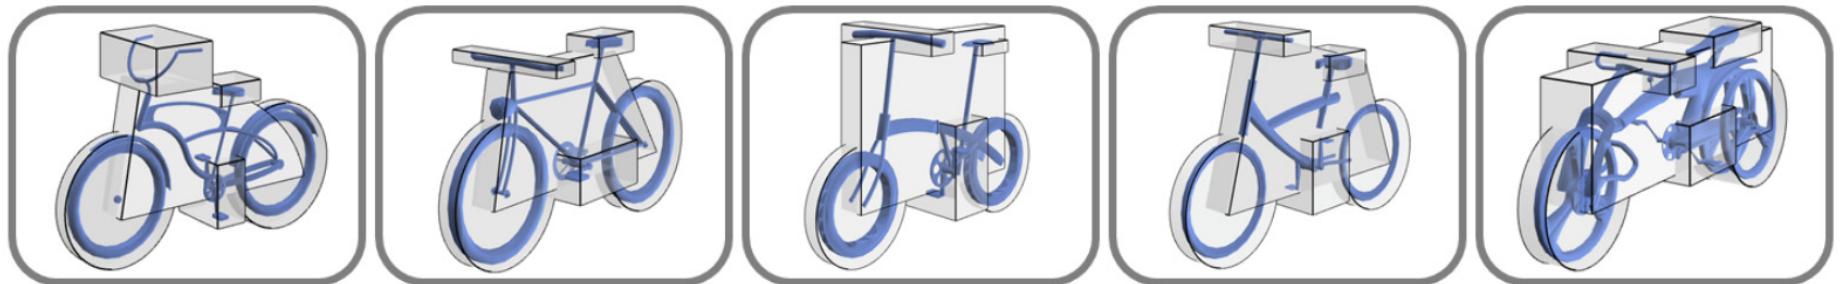


Image from E. Yumer., L. Kara, Co-Constrained Handles for Deformation in Shape Collections, 2014

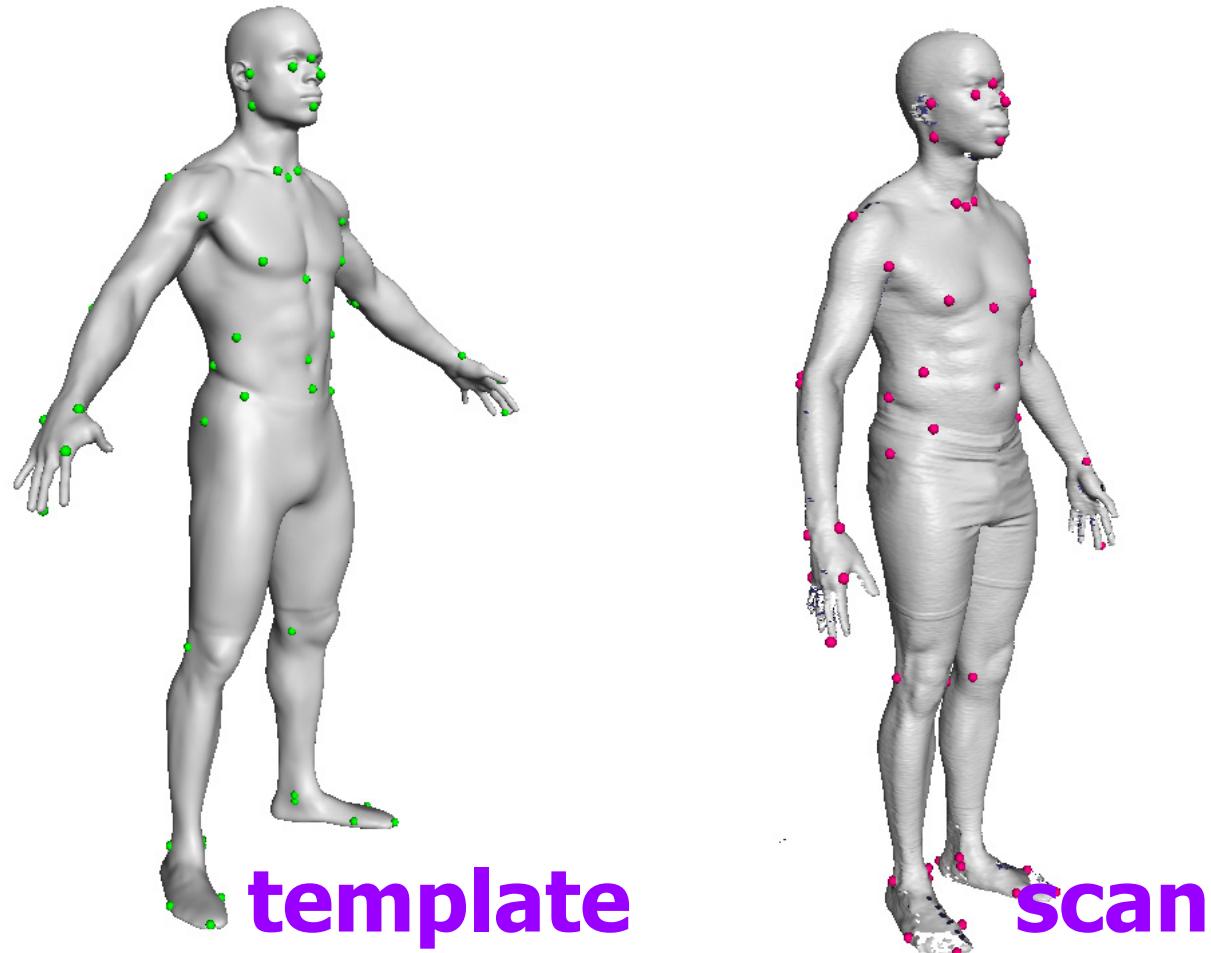
Case study: the space of human bodies

Training shapes: 125 male + 125 female scanned bodies



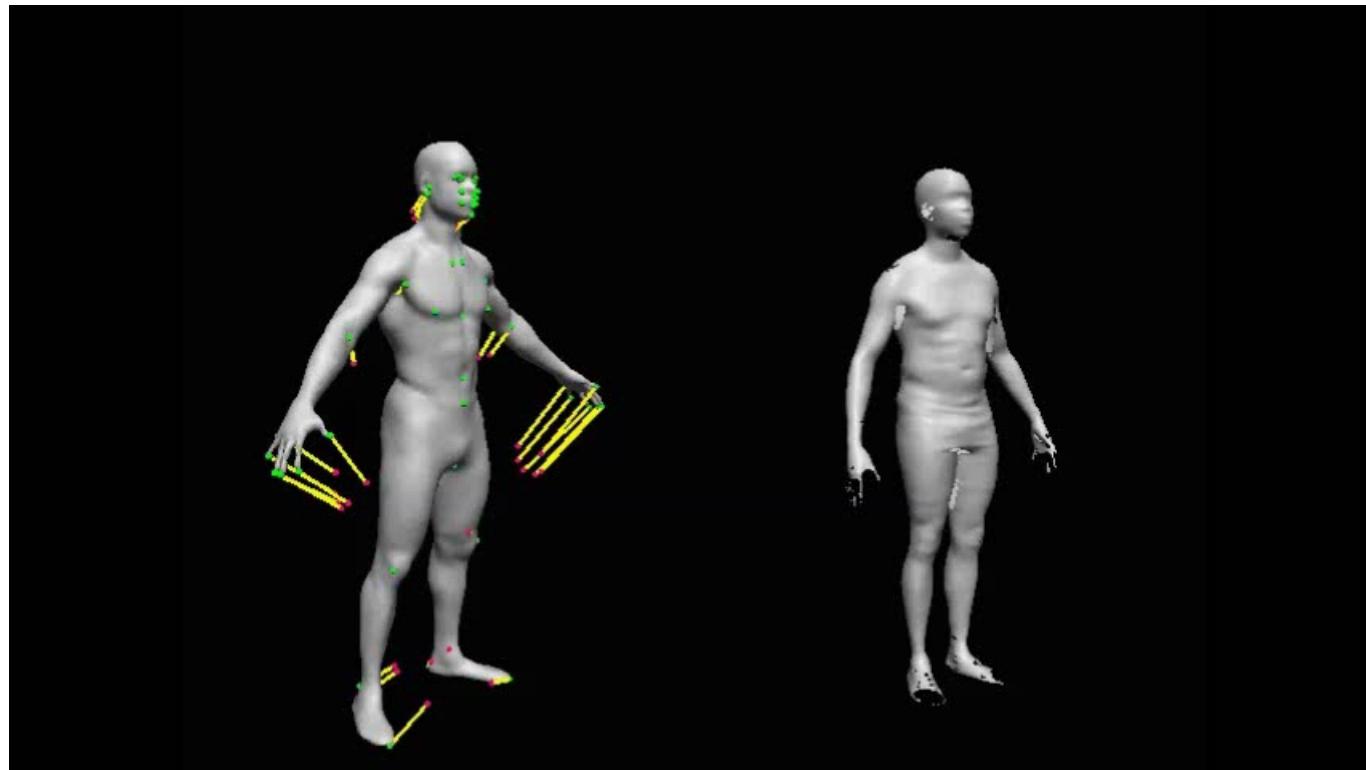
Slides from Brett Allen, Brian Curless, Zoran Popović, Exploring the space of human body shapes, 2003

Matching algorithm



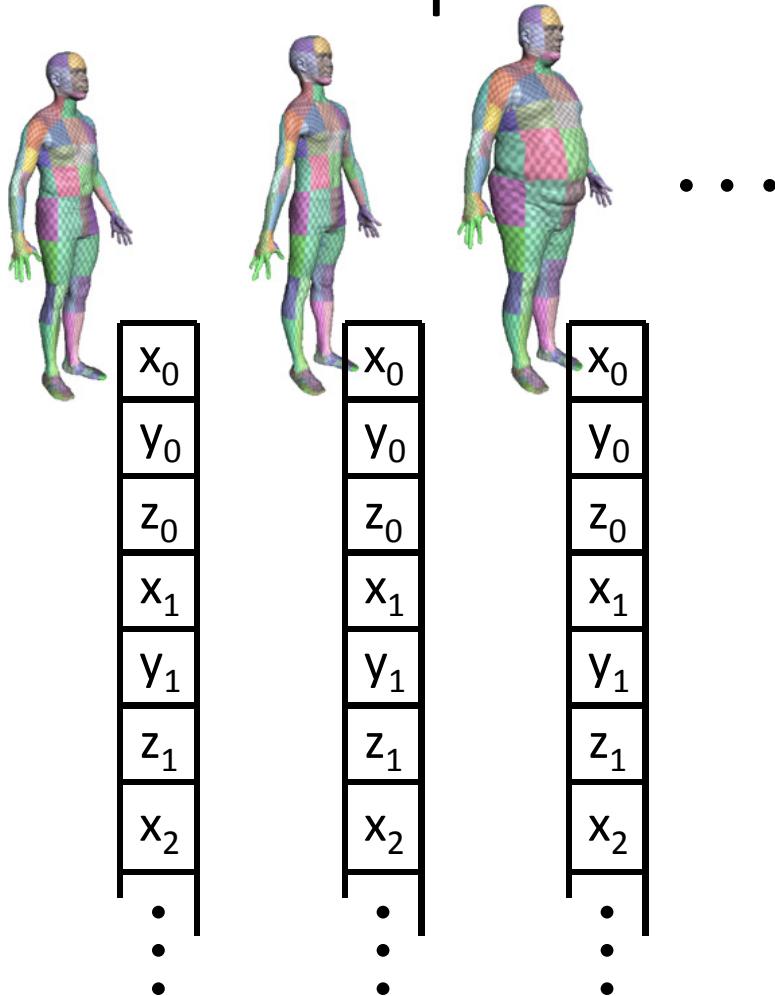
Slides from Brett Allen, Brian Curless, Zoran Popović, Exploring the space of human body shapes, 2003

Matching algorithm



*Slides from Brett Allen, Brian Curless, Zoran Popović, Exploring the space of human body shapes, 2003
to access the video: <http://grail.cs.washington.edu/projects/digital-human/pub/allen04exploring.html>*

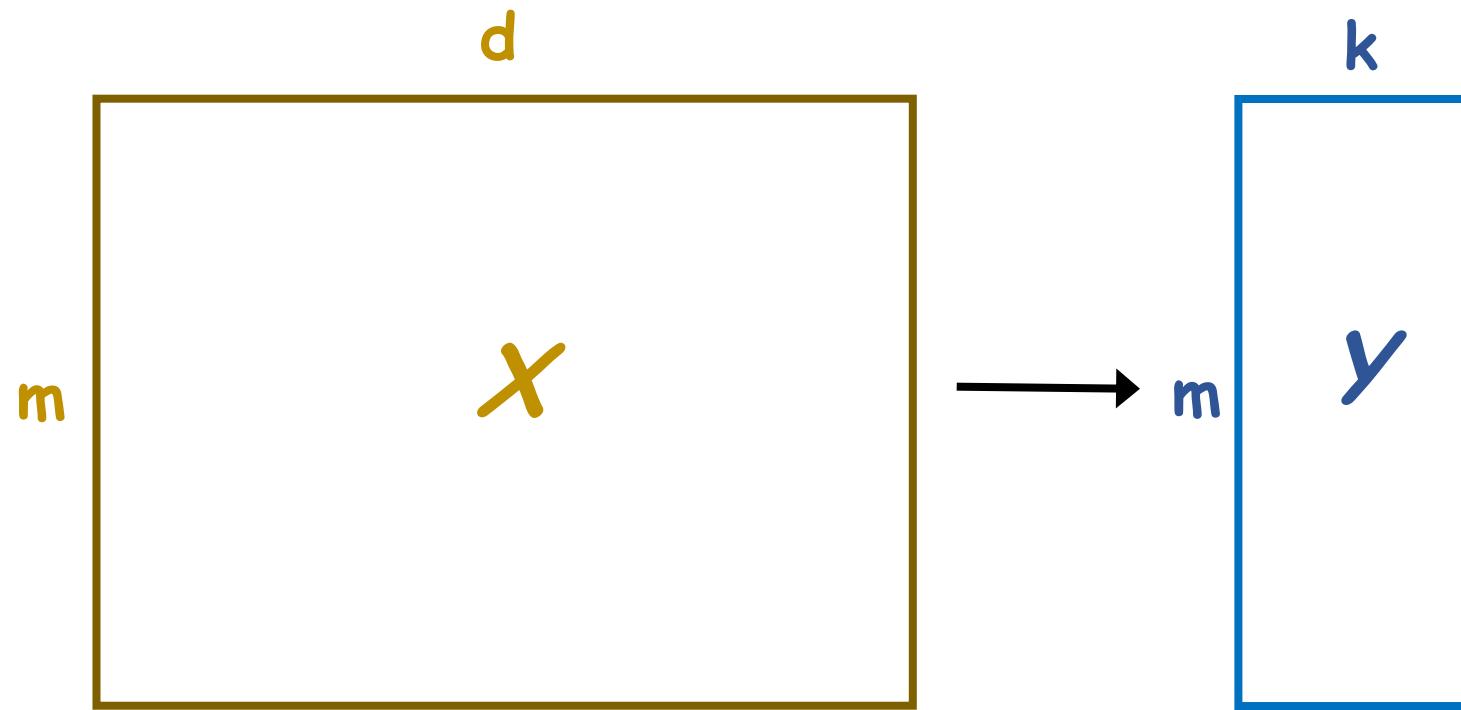
Principal Component Analysis



Slides from Brett Allen, Brian Curless, Zoran Popović, Exploring the space of human body shapes, 2003

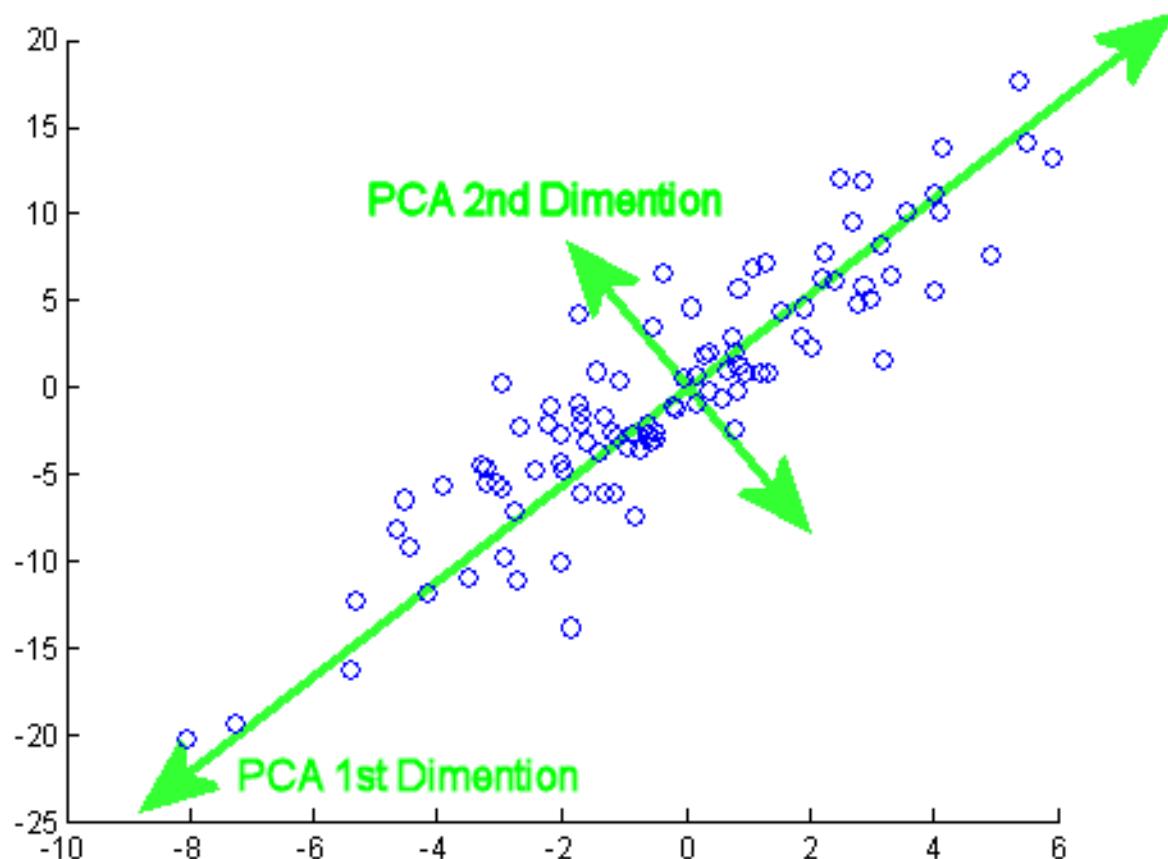
Dimensionality Reduction

Summarization of data with many (d) variables by a smaller set of (k) derived latent variables.

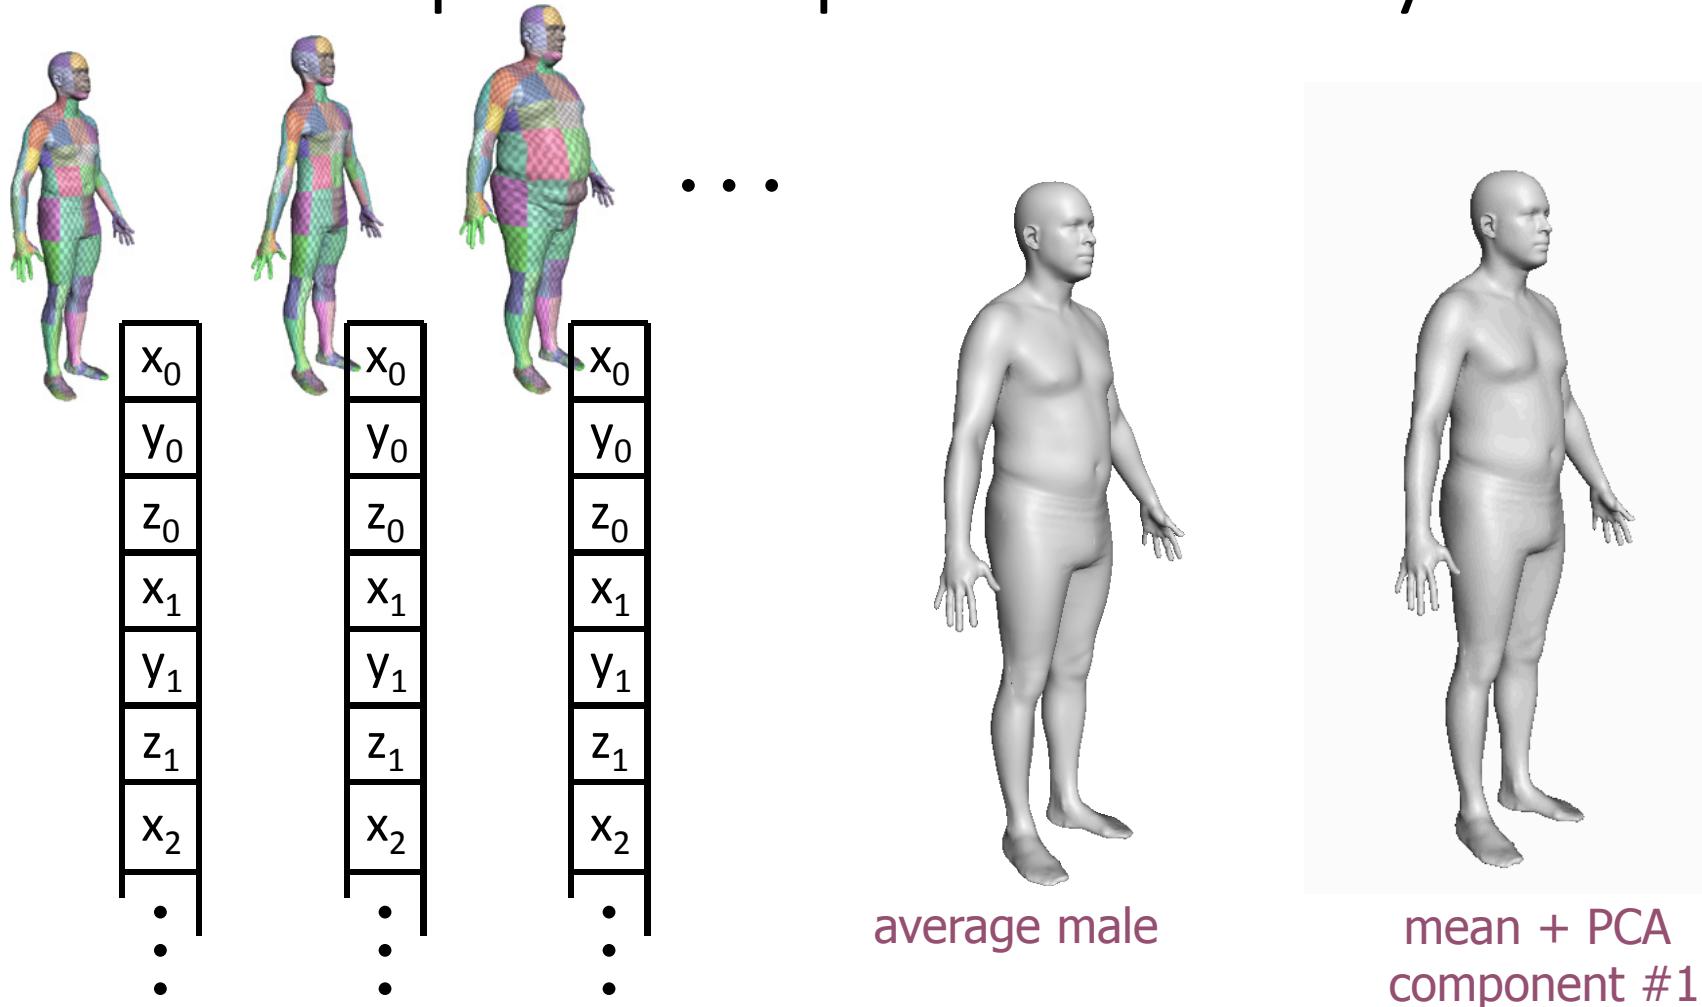


Principal Component Analysis

Each principal axis is a linear combination of the original variables

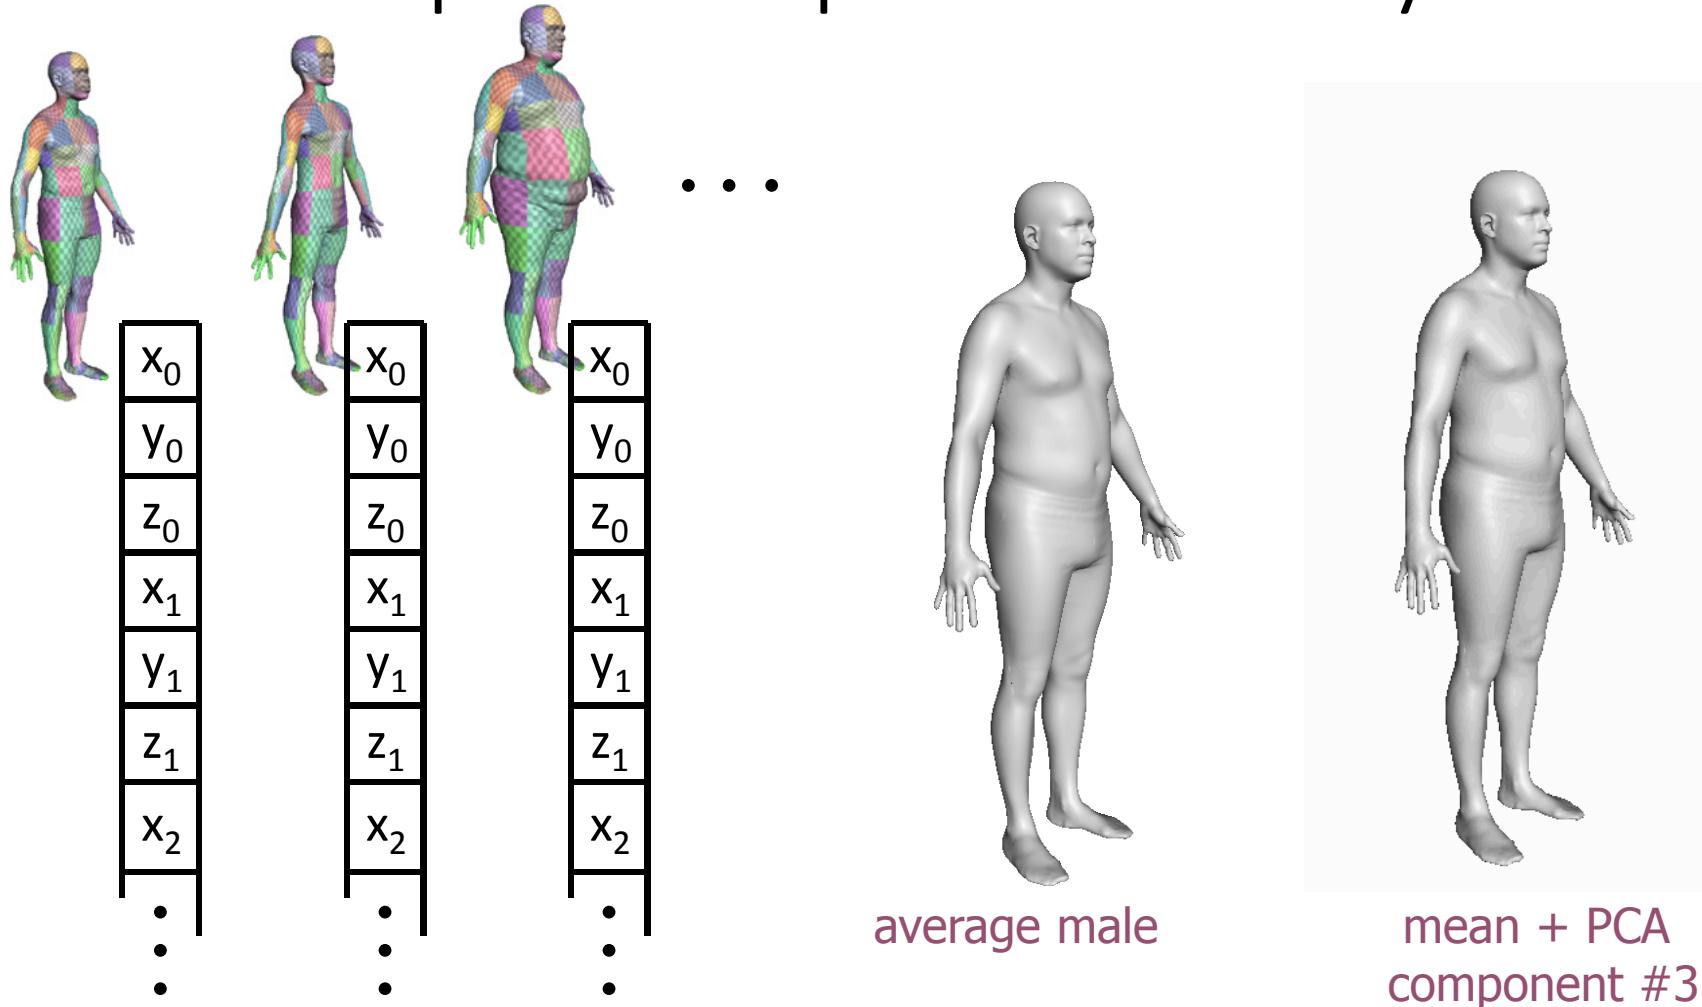


Principal Component Analysis



Slides from Brett Allen, Brian Curless, Zoran Popović, Exploring the space of human body shapes, 2003
to access the video: <http://grail.cs.washington.edu/projects/digital-human/pub/allen04exploring.html>

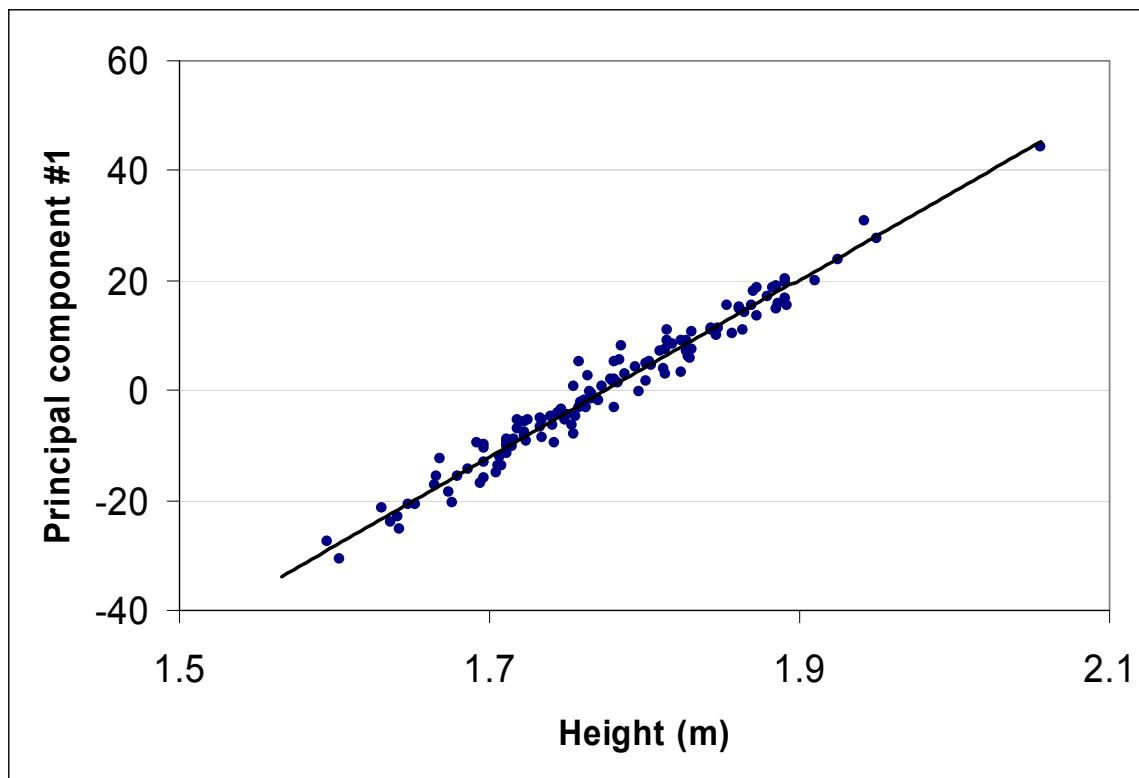
Principal Component Analysis



Slides from Brett Allen, Brian Curless, Zoran Popović, Exploring the space of human body shapes, 2003
to access the video: <http://grail.cs.washington.edu/projects/digital-human/pub/allen04exploring.html>

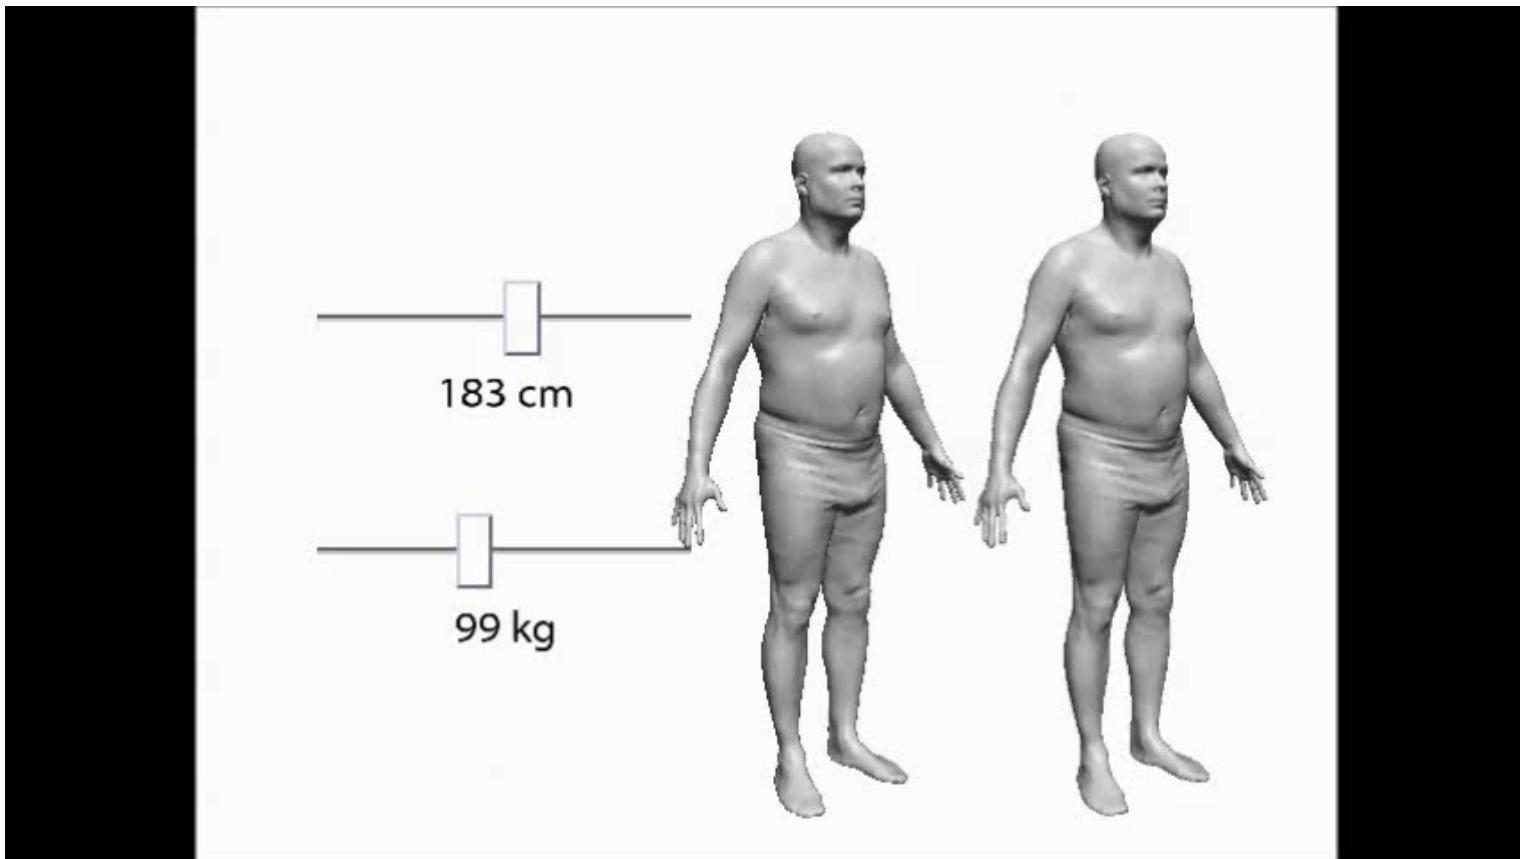
Fitting to attributes

Correlate PCA space with known attributes:



Slides from Brett Allen, Brian Curless, Zoran Popović, Exploring the space of human body shapes, 2003

Fitting to attributes



Slides from Brett Allen, Brian Curless, Zoran Popović, *Exploring the space of human body shapes*, 2003
to access the video: <http://grail.cs.washington.edu/projects/digital-human/pub/allen04exploring.html>

Case study: a probabilistic model for component-based synthesis

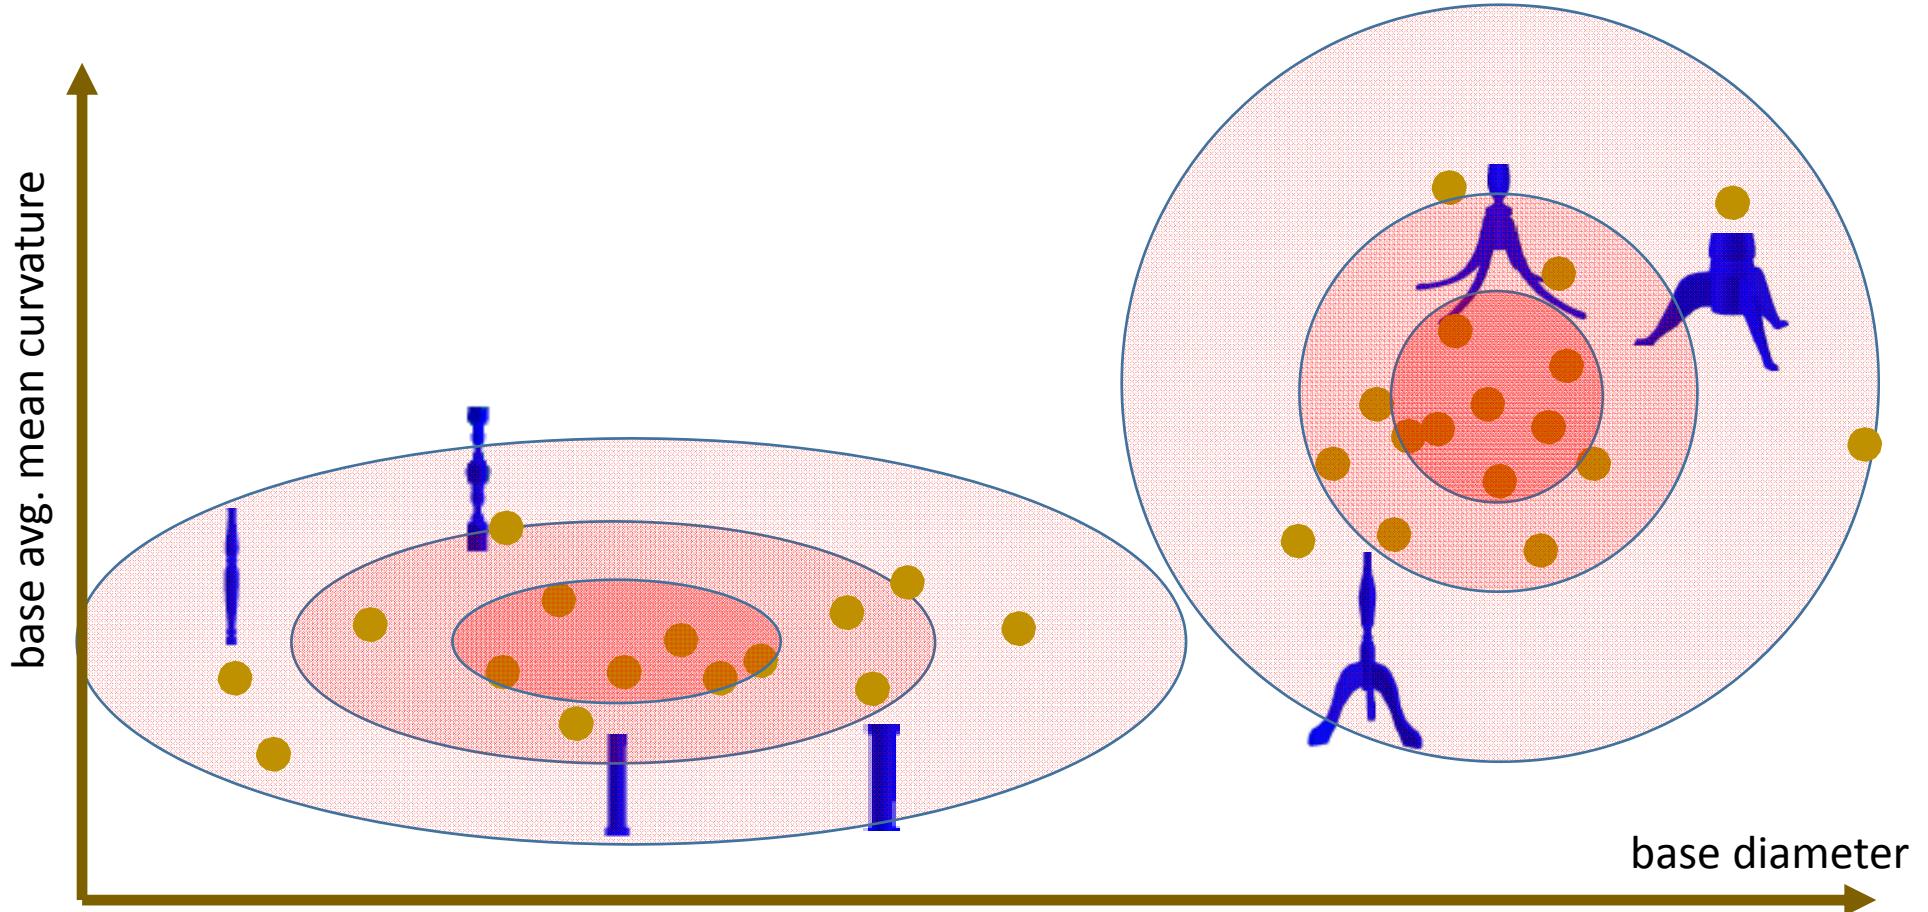
Given some training segmented shapes:



... and more

Case study: a probabilistic model for component-based synthesis

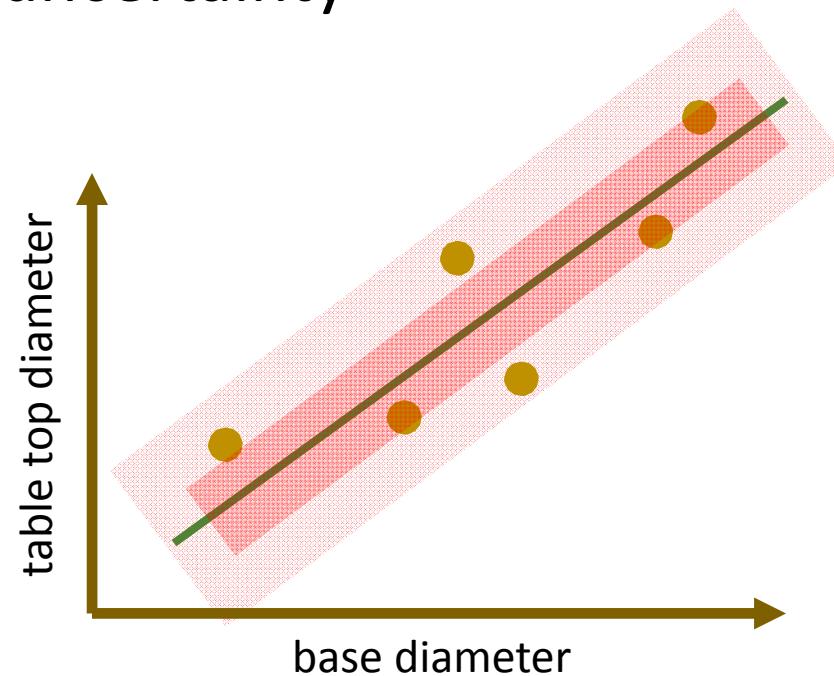
Describe shape space of parts with a probability distribution



*Slides from Evangelos Kalogerakis, Siddhartha Chaudhuri, Daphne Koller, Vladlen Koltun
A Probabilistic Model for Component-Based Synthesis, 2012*

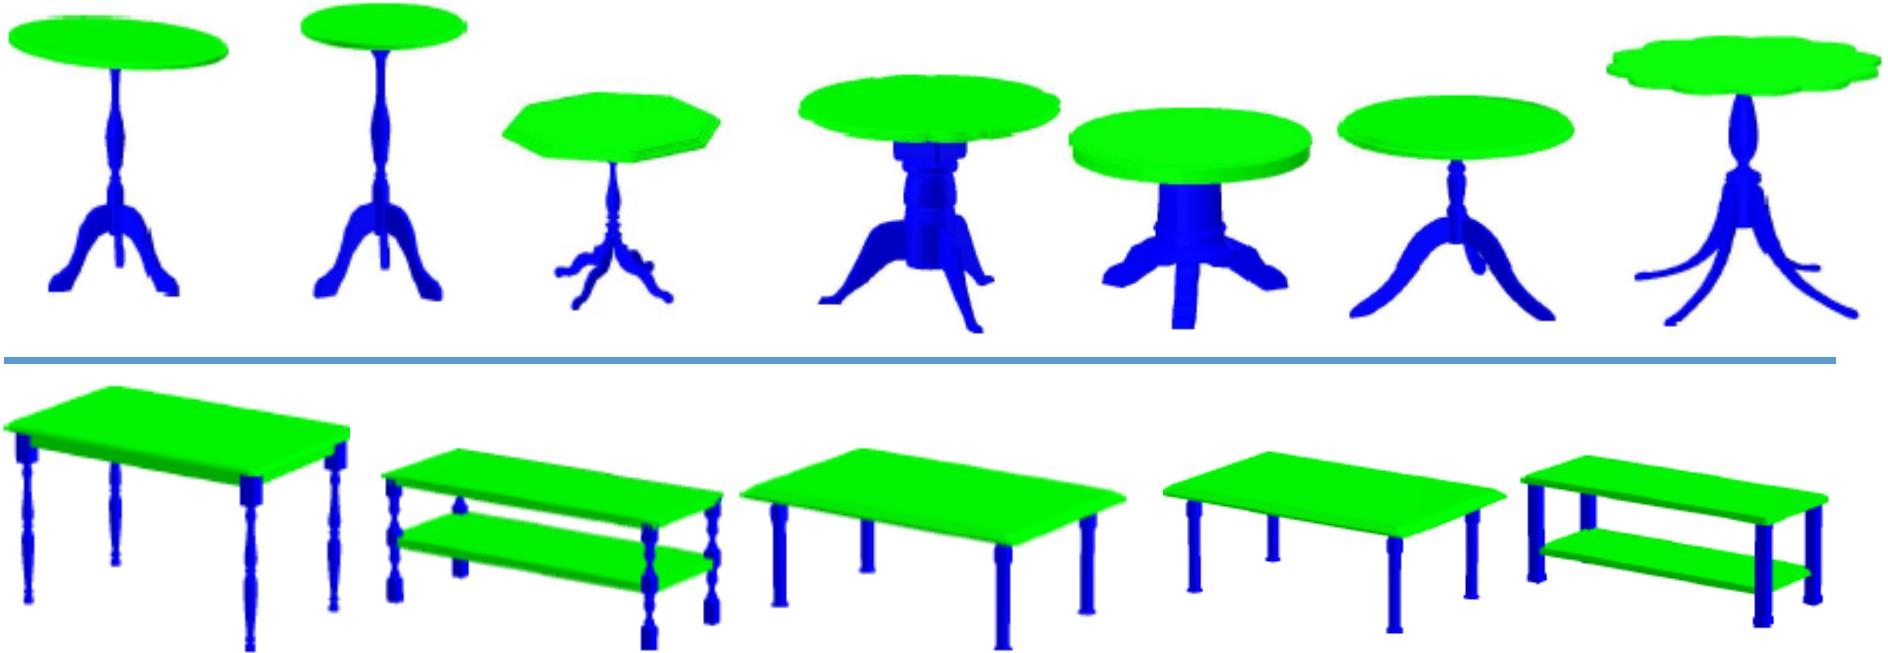
Case study: a probabilistic model for component-based synthesis

Learn relationships between different part parameters within each cluster e.g. diameter of table top is related to scale of base plus some uncertainty



Case study: a probabilistic model for component-based synthesis

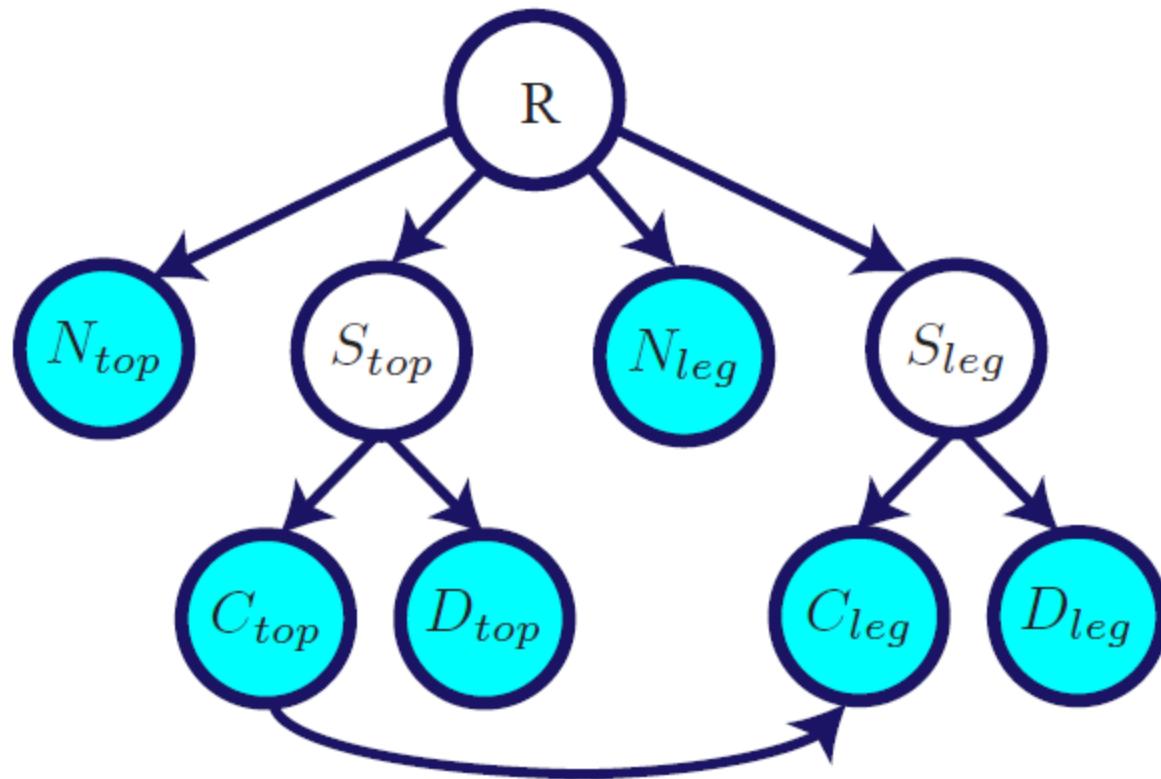
Learn relationships between part clusters e.g. circular table tops are associated with bases with split legs



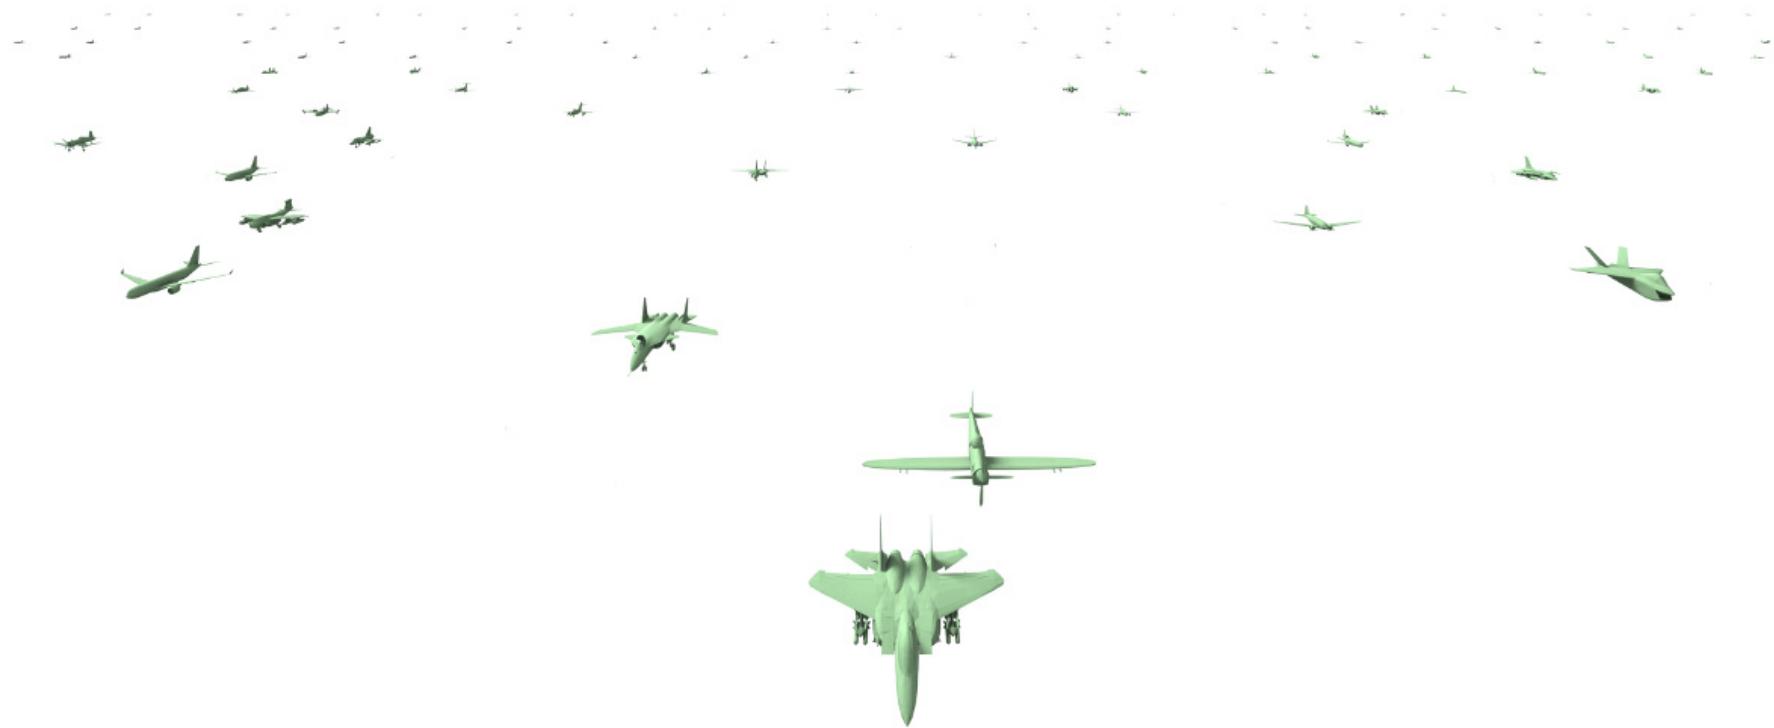
*Slides from Evangelos Kalogerakis, Siddhartha Chaudhuri, Daphne Koller, Vladlen Koltun
A Probabilistic Model for Component-Based Synthesis, 2012*

Case study: a probabilistic model for component-based synthesis

Represent all these relationships within a structured probability distribution (probabilistic graphical model)

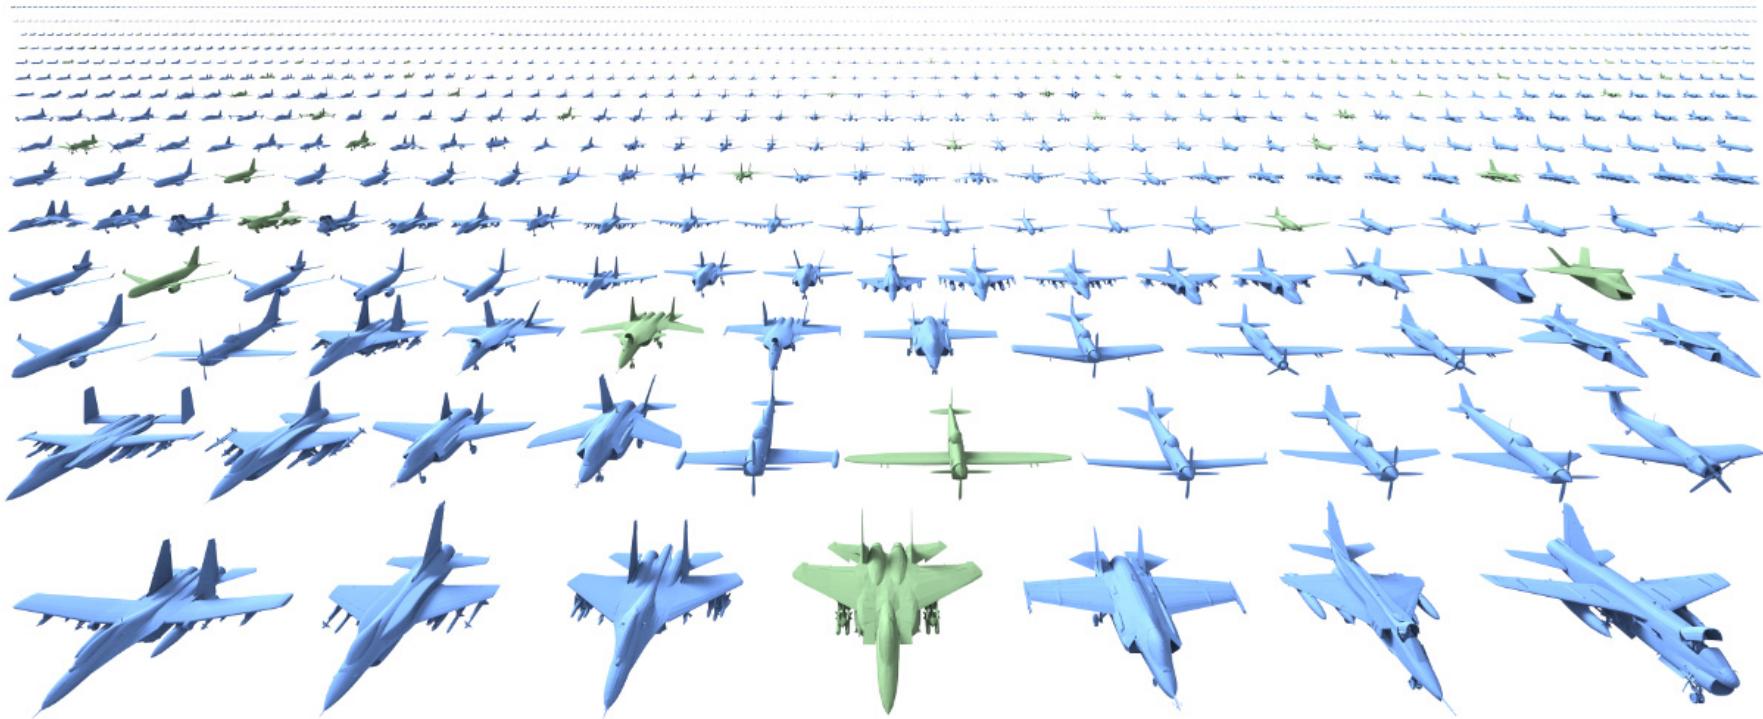


Shape Synthesis - Airplanes



*Slides from Evangelos Kalogerakis, Siddhartha Chaudhuri, Daphne Koller, Vladlen Koltun
A Probabilistic Model for Component-Based Synthesis, 2012*

Shape Synthesis - Airplanes



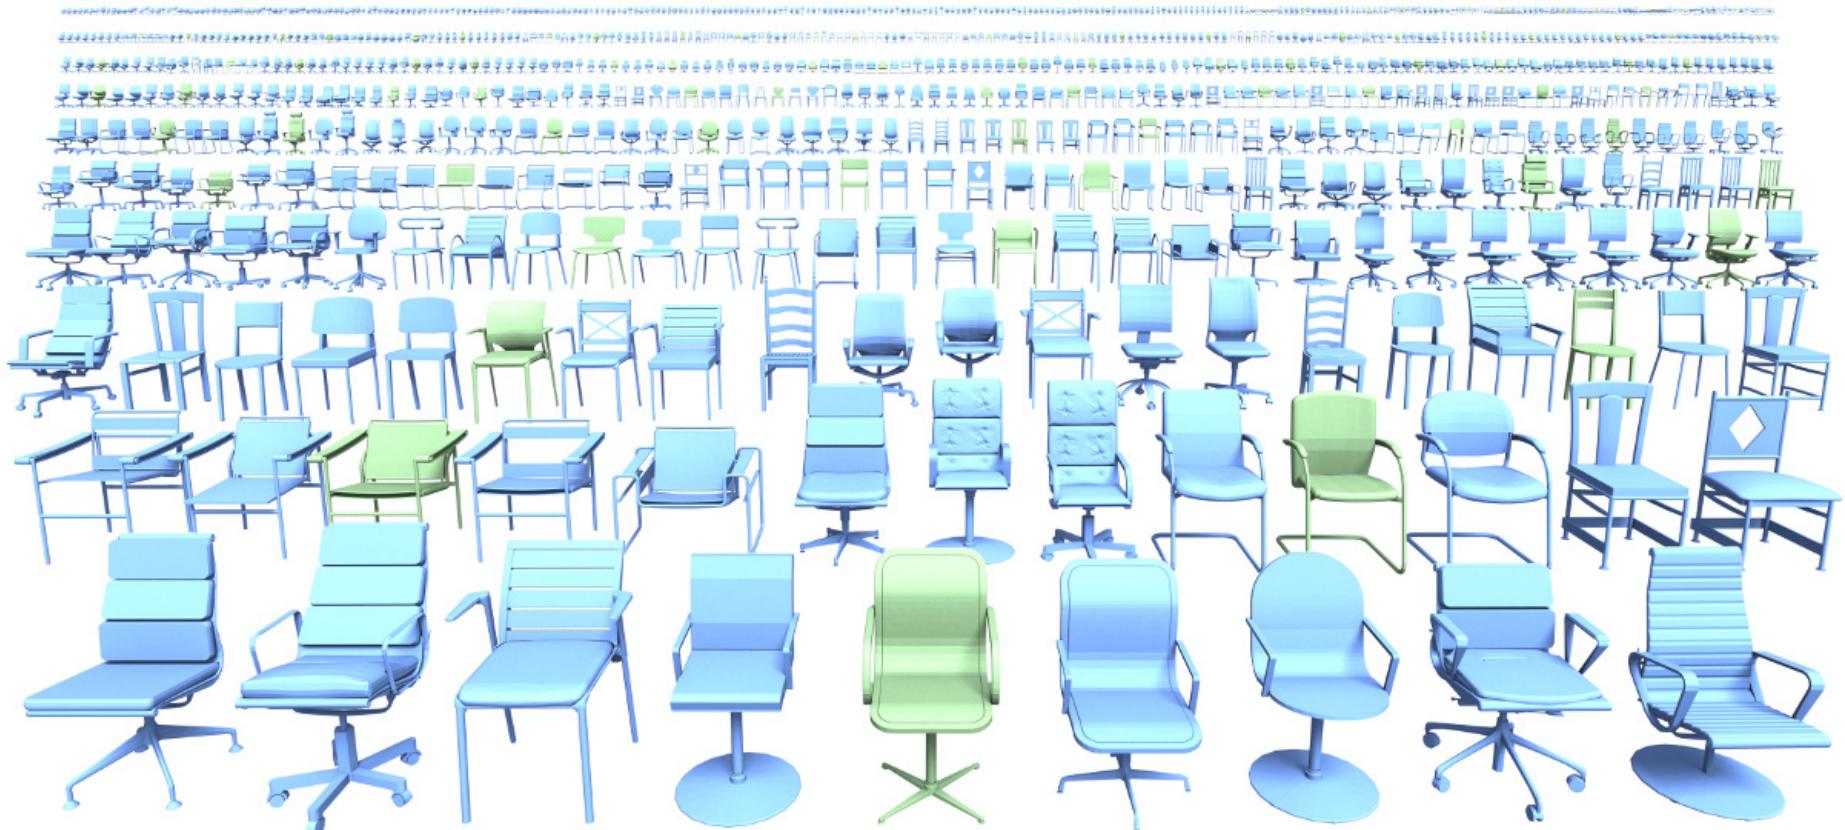
*Slides from Evangelos Kalogerakis, Siddhartha Chaudhuri, Daphne Koller, Vladlen Koltun
A Probabilistic Model for Component-Based Synthesis, 2012*

Shape Synthesis - Airplanes



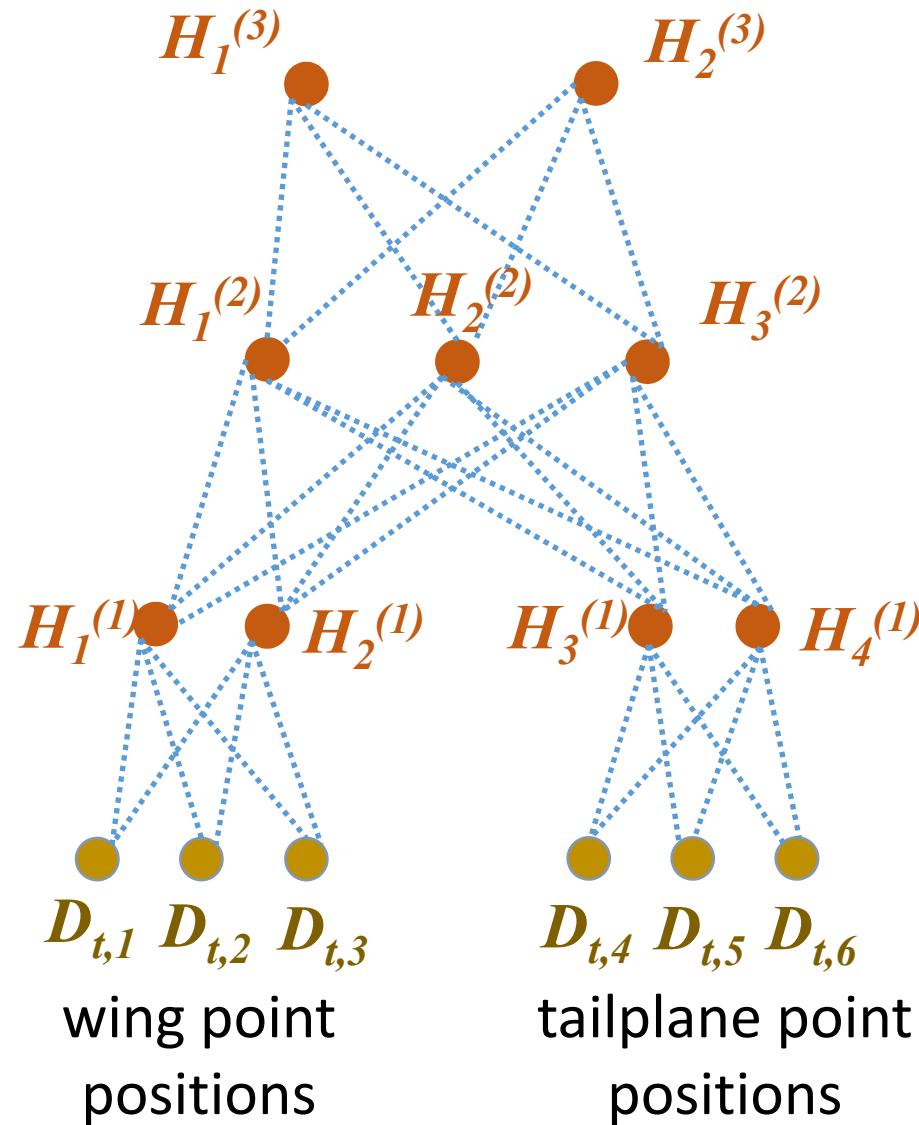
*Slides from Evangelos Kalogerakis, Siddhartha Chaudhuri, Daphne Koller, Vladlen Koltun
A Probabilistic Model for Component-Based Synthesis, 2012*

Shape Synthesis - Airplanes



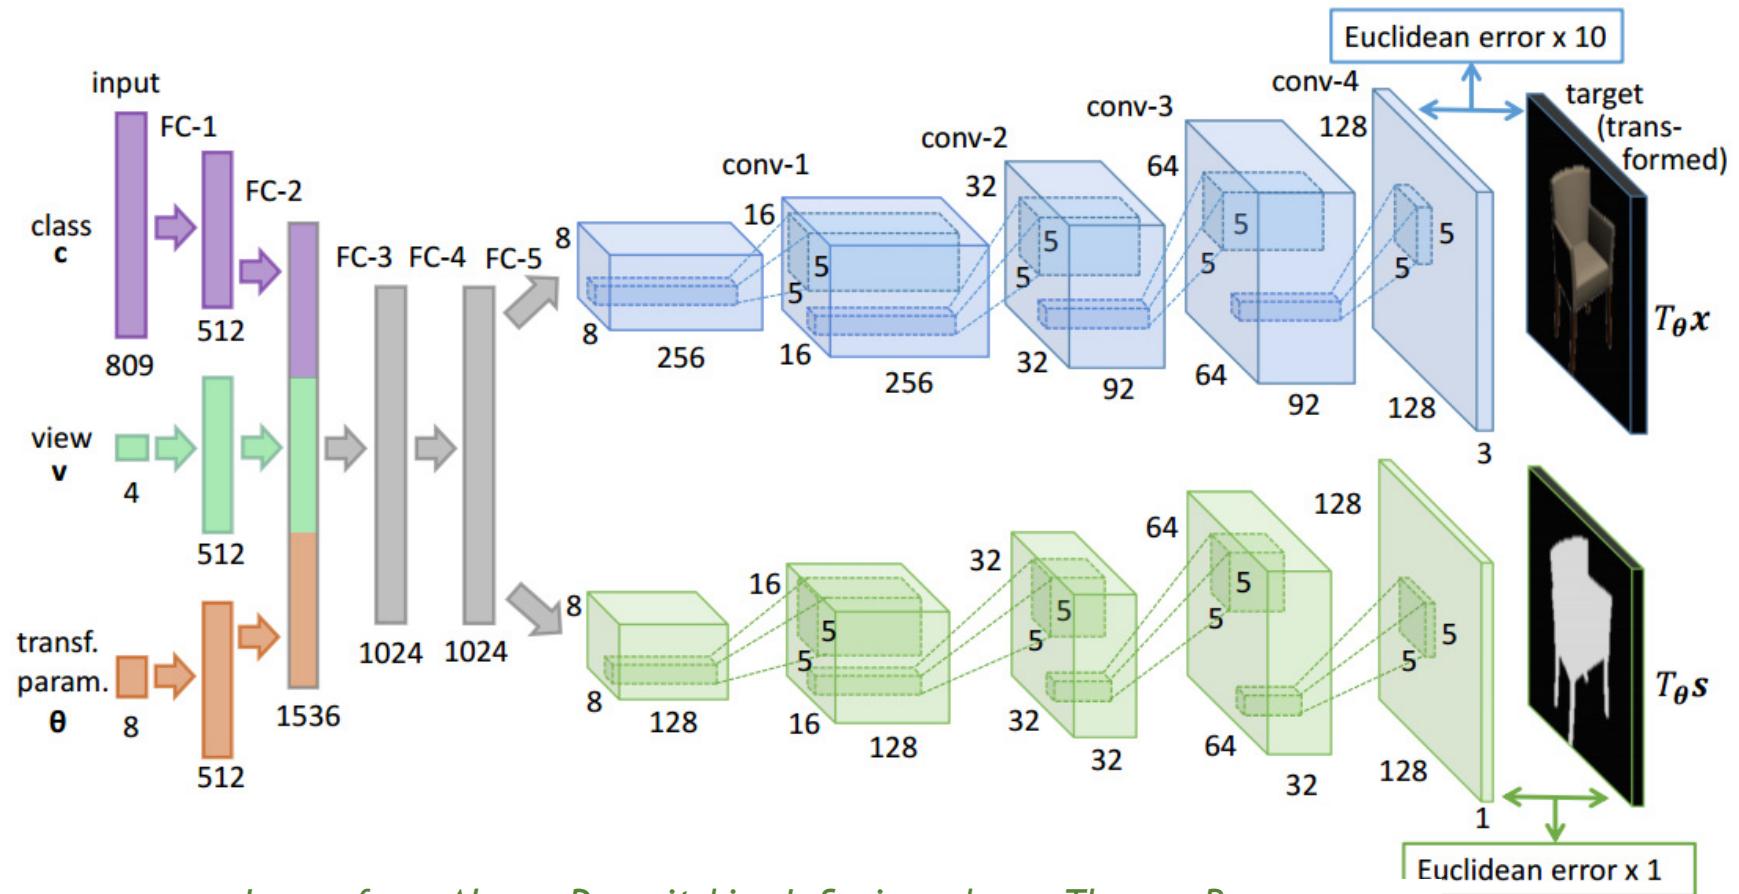
*Slides from Evangelos Kalogerakis, Siddhartha Chaudhuri, Daphne Koller, Vladlen Koltun
A Probabilistic Model for Component-Based Synthesis, 2012*

Generative models of surface geometry



Learning to Generate Chairs

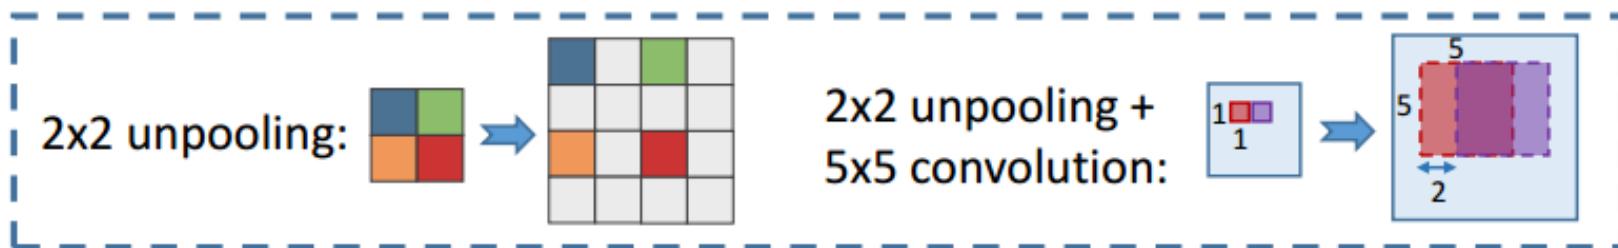
Inverting the CNN...



*Image from Alexey Dosovitskiy, J. Springenberg, Thomas Brox
Learning to Generate Chairs with Convolutional Neural Networks 2015*

Learning to Generate Chairs

Inverting the CNN...



*Image from Alexey Dosovitskiy, J. Springenberg, Thomas Brox
Learning to Generate Chairs with Convolutional Neural Networks 2015
to access video: <http://lmb.informatik.uni-freiburg.de/Publications/2015/DB15/>*

Summary

Welcome to the era where machines **learn to generate 3D visual content!**

Data-driven techniques with (deep) learning are highly promising directions

Summary

Welcome to the era where machines **learn to generate 3D visual content!**

Data-driven techniques with (deep) learning are highly promising directions

Some challenges:

- Generate **plausible, detailed, novel** 3D geometry from **high-level specifications, approximate** directions
- What **shape representation** should deep networks operate on?
- Integrate with approaches that optimize for **function, style** and **human-object interaction**